



Armadillo

C++ library for linear algebra & scientific computing

[About](#) [Documentation](#) [Questions](#) [Speed](#) [Contact](#) [Download](#)

API Documentation for Armadillo 12.6

Preamble

- For converting Matlab/Octave programs, see the [syntax conversion table](#)
- First time users: please see the short [example program](#)
- If you discover any bugs or regressions, please [report them](#)
- History of [API additions](#)

- Please cite the following papers if you use Armadillo in your research and/or software. Citations are useful for the continued development and maintenance of the library.

Conrad Sanderson and Ryan Curtin.

Armadillo: a template-based C++ library for linear algebra.

Journal of Open Source Software, Vol. 1, No. 2, pp. 26, 2016.

Conrad Sanderson and Ryan Curtin.

Practical sparse matrices in C++ with hybrid storage and template-based expression optimisation.

Mathematical and Computational Applications, Vol. 24, No. 3, 2019.

Overview

- matrix, vector, cube and field classes
- member functions & variables
- generated vectors / matrices / cubes
- functions of vectors / matrices / cubes
- decompositions, factorisations, inverses and equation solvers (dense matrices)
- decompositions, factorisations, and equation solvers (sparse matrices)
- signal & image processing
- statistics and clustering
- miscellaneous (constants, configuration options, ...)

Matrix, Vector, Cube and Field Classes

Mat<type>, mat, cx_mat	dense matrix class
Col<type>, colvec, vec	dense column vector class
Row<type>, rowvec	dense row vector class
Cube<type>, cube, cx_cube	dense cube class ("3D matrix")
field<object type>	class for storing arbitrary objects in matrix-like or cube-like layouts
SpMat<type>, sp_mat, sp_cx_mat	sparse matrix class
operators	+ - * % / == != <= >= < > &&

Member Functions & Variables

attributes	.n_rows, .n_cols, .n_elem, .n_slices, ...
element access	element/object access via (), [] and .at()
element initialisation	set elements via initialiser lists
.zeros	set all elements to zero
.ones	set all elements to one

<code>.eye</code> <code>.randu / .randn</code>	set elements along main diagonal to one and off-diagonal elements to zero set all elements to random values
<code>.fill</code> <code>.imbue</code>	set all elements to specified value imbue (fill) with values provided by functor or lambda function
<code>.clean</code> <code>.replace</code> <code>.clamp</code>	replace elements below a threshold with zeros replace specific elements with a new value clamp values to lower and upper limits
<code>.transform</code> <code>.for_each</code>	transform each element via functor or lambda function apply a functor or lambda function to each element
<code>.set_size</code> <code>.reshape</code> <code>.resize</code> <code>.copy_size</code> <code>.reset</code>	change size without keeping elements (fast) change size while keeping elements change size while keeping elements and preserving layout change size to be same as given object change size to empty
submatrix views subcube views subfield views	read/write access to contiguous and non-contiguous submatrices read/write access to contiguous and non-contiguous subcubes read/write access to contiguous subfields
<code>.diag</code> <code>.each_col / .each_row</code> <code>.each_slice</code>	read/write access to matrix diagonals vector operations applied to each column/row of matrix (aka "broadcasting") matrix operations applied to each slice of cube (aka "broadcasting")
<code>.set_imag / .set_real</code> <code>.insert_rows / cols / slices</code> <code>.shed_rows / cols / slices</code> <code>.swap_rows / cols</code> <code>.swap</code>	set imaginary/real part insert vector/matrix/cube at specified row/column/slice remove specified rows/columns/slices swap specified rows or columns swap contents with given object

<code>.memptr</code>	raw pointer to memory
<code>.colptr</code>	raw pointer to memory used by specified column
<code>iterators (matrices)</code>	iterators and associated member functions for dense matrices and vectors
<code>iterators (cubes)</code>	iterators and associated member functions for cubes
<code>iterators (sparse matrices)</code>	iterators and associated member functions for sparse matrices
<code>iterators (submatrices)</code>	iterators and associated member functions for submatrices & subcubes
<code>compat. container functions</code>	compatibility container functions
<code>.as_col / .as_row</code>	return flattened matrix as column or row vector
<code>.col_as_mat / .row_as_mat</code>	return matrix representation of cube column or cube row
<code>.t / .st</code>	return matrix transpose
<code>.i</code>	return inverse of square matrix
<code>.min / .max</code>	return extremum value
<code>.index_min / .index_max</code>	return index of extremum value
<code>.eval</code>	force evaluation of delayed expression
<code>.in_range</code>	check whether given location or span is valid
<code>.is_empty</code>	check whether object is empty
<code>.is_vec</code>	check whether matrix is a vector
<code>.is_sorted</code>	check whether vector or matrix is sorted
<code>.is_trimatu / .is_trimatl</code>	check whether matrix is upper/lower triangular
<code>.is_diagmat</code>	check whether matrix is diagonal
<code>.is_square</code>	check whether matrix is square sized
<code>.is_symmetric</code>	check whether matrix is symmetric
<code>.is_hermitian</code>	check whether matrix is hermitian
<code>.is_sympd</code>	check whether matrix is symmetric/hermitian positive definite
<code>.is_zero</code>	check whether all elements are zero
<code>.is_finite</code>	check whether all elements are finite

<code>.has_inf</code>	check whether any element is \pm infinity
<code>.has_nan</code>	check whether any element is NaN
<code>.print</code>	print object to <i>std::cout</i> or user specified stream
<code>.raw_print</code>	print object without formatting
<code>.brief_print</code>	print object in shortened/abridged form
<code>.save/.load (matrices & cubes)</code>	save/load matrices and cubes in files or streams
<code>.save/.load (fields)</code>	save/load fields in files or streams

Generated Vectors / Matrices / Cubes

<code>linspace</code>	generate vector with linearly spaced elements
<code>logspace</code>	generate vector with logarithmically spaced elements
<code>regspace</code>	generate vector with regularly spaced elements
<code>randperm</code>	generate vector with random permutation of a sequence of integers
<code>eye</code>	generate identity matrix
<code>ones</code>	generate object filled with ones
<code>zeros</code>	generate object filled with zeros
<code>randu</code>	generate object with random values (uniform distribution)
<code>randn</code>	generate object with random values (normal distribution)
<code>randg</code>	generate object with random values (gamma distribution)
<code>randi</code>	generate object with random integer values in specified interval
<code>speye</code>	generate sparse identity matrix
<code>spones</code>	generate sparse matrix with non-zero elements set to one
<code>sprandu / sprandn</code>	generate sparse matrix with non-zero elements set to random values
<code>toeplitz</code>	generate Toeplitz matrix

Functions of Vectors / Matrices / Cubes

<code>abs</code>	obtain magnitude of each element
<code>accu</code>	accumulate (sum) all elements

affmul	affine matrix multiplication
all	check whether all elements are non-zero, or satisfy a relational condition
any	check whether any element is non-zero, or satisfies a relational condition
approx_equal	approximate equality
arg	phase angle of each element
as_scalar	convert 1x1 matrix to pure scalar
clamp	obtain clamped elements according to given limits
cond	condition number of matrix
conj	obtain complex conjugate of each element
conv_to	convert/cast between matrix types
cross	cross product
cumsum	cumulative sum
cumprod	cumulative product
det	determinant
diagmat	generate diagonal matrix from given matrix or vector
diagvec	extract specified diagonal
diags / spdiags	generate band matrix from given set of vectors
diff	differences between adjacent elements
dot / cdot / norm_dot	dot product
eps	obtain distance of each element to next largest floating point representation
expmat	matrix exponential
expmat_sym	matrix exponential of symmetric matrix
find	find indices of non-zero elements, or elements satisfying a relational condition
find_finite	find indices of finite elements
find_nonfinite	find indices of non-finite elements
find_nan	find indices of NaN elements
find_unique	find indices of unique elements
fliplr / flipud	flip matrix left to right or upside down
imag / real	extract imaginary/real part
ind2sub	convert linear index to subscripts
index_min / index_max	indices of extremum values
inplace_trans	in-place transpose

intersect	find common elements in two vectors/matrices
join_rows / join_cols	concatenation of matrices
join_slices	concatenation of cubes
kron	Kronecker tensor product
log_det	log determinant
log_det_sympd	log determinant of symmetric positive definite matrix
logmat	matrix logarithm
logmat_sympd	matrix logarithm of symmetric matrix
min / max	return extremum values
nonzeros	return non-zero values
norm	various norms of vectors and matrices
norm2est	fast estimate of the matrix 2-norm
normalise	normalise vectors to unit p -norm
pow	element-wise power
powmat	matrix power
prod	product of elements
rank	rank of matrix
rcond	reciprocal condition number
repelem	replicate elements
repmat	replicate matrix in block-like fashion
reshape	change size while keeping elements
resize	change size while keeping elements and preserving layout
reverse	reverse order of elements
roots	roots of polynomial
shift	shift elements
shuffle	randomly shuffle elements
size	obtain dimensions of given object
sort	sort elements
sort_index	vector describing sorted order of elements
sqrmat	square root of matrix
sqrmat_sympd	square root of symmetric matrix
sum	sum of elements

sub2ind	convert subscripts to linear index
symmatu / symmatl	generate symmetric matrix from given matrix
trace	sum of diagonal elements
trans	transpose of matrix
trapz	trapezoidal numerical integration
trimatu / trimatl	copy upper/lower triangular part
trimatu_ind / trimatl_ind	obtain indices of upper/lower triangular part
unique	return unique elements
vecnorm	obtain vector norm of each row or column of a matrix
vectorise	flatten matrix into vector
misc functions	miscellaneous element-wise functions: exp, log, sqrt, round, sign, ...
trig functions	trigonometric element-wise functions: cos, sin, tan, ...

Decompositions, Factorisations, Inverses and Equation Solvers (Dense Matrices)

chol	Cholesky decomposition
eig_sym	eigen decomposition of dense symmetric/hermitian matrix
eig_gen	eigen decomposition of dense general square matrix
eig_pair	eigen decomposition for pair of general dense square matrices
hess	upper Hessenberg decomposition
inv	inverse of general square matrix
inv_sympd	inverse of symmetric positive definite matrix
lu	lower-upper decomposition
null	orthonormal basis of null space
orth	orthonormal basis of range space
pinv	pseudo-inverse / generalised inverse
qr	QR decomposition
qr_econ	economical QR decomposition
qz	generalised Schur decomposition
schur	Schur decomposition
solve	solve systems of linear equations
svd	singular value decomposition

svd_econ	economical singular value decomposition
syl	Sylvester equation solver

Decompositions, Factorisations and Equation Solvers (Sparse Matrices)

eigs_sym	limited number of eigenvalues & eigenvectors of sparse symmetric real matrix
eigs_gen	limited number of eigenvalues & eigenvectors of sparse general square matrix
svds	truncated svd: limited number of singular values & singular vectors of sparse matrix
spsolve	solve sparse systems of linear equations
spsolve_factoriser	factoriser for solving sparse systems of linear equations

Signal & Image Processing

conv	1D convolution
conv2	2D convolution
fft / ifft	1D fast Fourier transform and its inverse
fft2 / ifft2	2D fast Fourier transform and its inverse
interp1	1D interpolation
interp2	2D interpolation
polyfit	find polynomial coefficients for data fitting
polyval	evaluate polynomial

Statistics & Clustering

stats functions	mean, median, standard deviation, variance
cov	covariance
cor	correlation
hist	histogram of counts
histc	histogram of counts with user specified edges
quantile	quantiles of a dataset
princomp	principal component analysis (PCA)

normpdf	probability density function of normal distribution
log_normpdf	logarithm version of probability density function of normal distribution
normcdf	cumulative distribution function of normal distribution
mvnrnd	random vectors from multivariate normal distribution
chi2rnd	random numbers from chi-squared distribution
wishrnd	random matrix from Wishart distribution
iwishrnd	random matrix from inverse Wishart distribution
running_stat	running statistics of scalars (one dimensional process/signal)
running_stat_vec	running statistics of vectors (multi-dimensional process/signal)
kmeans	cluster data into disjoint sets
gmm_diag/gmm_full	model and evaluate data using Gaussian Mixture Models (GMMs)

Miscellaneous

constants	pi, inf, NaN, eps, speed of light, ...
wall_clock	timer for measuring number of elapsed seconds
RNG seed setting	functions for changing RNG seeds
output streams	streams for printing warnings and errors
uword / sword	shorthand for unsigned and signed integers
cx_double / cx_float	shorthand for std::complex<double> and std::complex<float>
Matlab/Armadillo syntax differences	examples of Matlab syntax and conceptually corresponding Armadillo syntax
example program	short example program
config.hpp	configuration options
API additions	API stability and list of API additions

Matrix, Vector, Cube and Field Classes

Mat<type>
mat
cx_mat

- Classes for dense matrices, with elements stored in **column-major ordering** (ie. column by column)
- The root matrix class is **Mat<type>**, where *type* is one of:
 - *float*, *double*, *std::complex<float>*, *std::complex<double>*, *short*, *int*, *long*, and unsigned versions of *short*, *int*, *long*

- For convenience the following typedefs have been defined:

```
mat = Mat<double>
dmat = Mat<double>
fmat = Mat<float>
cx_mat = Mat<cx_double>
cx_dmat = Mat<cx_double>
cx_fmat = Mat<cx_float>
umat = Mat<uword>
imat = Mat<sword>
```

- In this documentation the *mat* type is used for convenience; it is possible to use other matrix types instead, eg. *fmat*
- Matrix types with integer elements (such as *umat* and *imat*) cannot hold special values such as NaN and Inf
- Functions which use LAPACK (generally matrix decompositions) are only valid for the following matrix types: *mat*, *dmat*, *fmat*, *cx_mat*, *cx_dmat*, *cx_fmat*

- Constructors:

```
mat()
mat(n_rows, n_cols)
mat(n_rows, n_cols, fill_form) (elements are initialised according to fill_form)
mat(size(X))
mat(size(X), fill_form) (elements are initialised according to fill_form)
```

<code>mat(mat)</code>	
<code>mat(vec)</code>	
<code>mat(rowvec)</code>	
<code>mat(initializer_list)</code>	
<code>mat(string)</code>	
<code>mat(std::vector)</code>	(treated as a column vector)
<code>mat(sp_mat)</code>	(for converting a sparse matrix to a dense matrix)
<code>cx_mat(mat,mat)</code>	(for constructing a complex matrix out of two real matrices)

- The elements can be explicitly initialised during construction by specifying *fill_form*, which is one of:

<code>fill::zeros</code>	⇒ set all elements to 0
<code>fill::ones</code>	⇒ set all elements to 1
<code>fill::eye</code>	⇒ set the elements on the main diagonal to 1 and off-diagonal elements to 0
<code>fill::randu</code>	⇒ set all elements to random values from a uniform distribution in the [0,1] interval
<code>fill::randn</code>	⇒ set all elements to random values from a normal/Gaussian distribution with zero mean and unit variance
<code>fill::value(scalar)</code>	⇒ set all elements to specified scalar
<code>fill::none</code>	⇒ do not initialise the elements

- **Caveat:**

- since Armadillo 10.5, the elements are initialised to zero by default
 - in Armadillo 10.4 and older versions, the elements are not initialised unless *fill_form* is specified; ie. without specifying *fill_form*, the elements may contain garbage values, including NaN
- For the *mat(string)* constructor, the format is elements separated by spaces, and rows denoted by semicolons; for example, the 2x2 identity matrix can be created using "1 0; 0 1"; note that string based initialisation is slower than directly **setting the elements** or using **element initialisation**
 - Each instance of *mat* automatically allocates and releases internal memory. All internally allocated memory used by an instance of *mat* is automatically released as soon as the instance goes out of scope. For example, if an instance of *mat* is declared inside a function, it will be automatically destroyed at the end of the function. To forcefully release memory at any point, use **.reset()**; note that in normal use this is not required.

- Advanced constructors:

```
mat(ptr_aux_mem, n_rows, n_cols, copy_aux_mem = true, strict = false)
```

Create a matrix using data from writable auxiliary (external) memory, where *ptr_aux_mem* is a pointer to the memory. By default the matrix allocates its own memory and copies data from the auxiliary memory (for safety). However, if *copy_aux_mem* is set to *false*, the matrix will instead directly use the auxiliary memory (ie. no copying); this is faster, but can be dangerous unless you know what you are doing!

The *strict* parameter comes into effect only when *copy_aux_mem* is set to *false* (ie. the matrix is directly using auxiliary memory)

- when *strict* is set to *false*, the matrix will use the auxiliary memory until a size change or an aliasing event
- when *strict* is set to *true*, the matrix will be bound to the auxiliary memory for its lifetime; the number of elements in the matrix can't be changed

```
mat(const ptr_aux_mem, n_rows, n_cols)
```

Create a matrix by copying data from read-only auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

```
mat::fixed<n_rows, n_cols>
```

Create a fixed size matrix, with the size specified via template arguments. Memory for the matrix is reserved at compile time. This is generally faster than dynamic memory allocation, but the size of the matrix can't be changed afterwards (directly or indirectly).

For convenience, there are several pre-defined typedefs for each matrix type (where the types are: *umat*, *imat*, *fmat*, *mat*, *cx_fmat*, *cx_mat*). The typedefs specify a square matrix size, ranging from 2x2 to 9x9. The typedefs were defined by appending a two digit form of the size to the matrix type; examples: *mat33* is equivalent to *mat::fixed<3,3>*, while *cx_mat44* is equivalent to *cx_mat::fixed<4,4>*.

```
mat::fixed<n_rows, n_cols>(fill_form)
```

Create a fixed size matrix, with the elements explicitly initialised according to *fill_form*

```
mat::fixed<n_rows, n_cols>(const ptr_aux_mem)
```

Create a fixed size matrix, with the size specified via template arguments; data is copied from auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

- Examples:

```
mat A(5, 5, fill::randu);
double x = A(1,2);

mat B = A + A;
mat C = A * B;
mat D = A % B;

cx_mat X(A,B);

B.zeros();
B.set_size(10,10);
B.ones(5,6);

B.print("B:");

mat::fixed<5,6> F;

double aux_mem[24];
mat H(&aux_mem[0], 4, 6, false); // use auxiliary memory
```

- See also:

- [matrix attributes](#)
- [accessing elements](#)
- [initialising elements](#)
- [math & relational operators](#)
- [submatrix views](#)
- [saving & loading matrices](#)
- [printing matrices](#)
- [element iterators](#)
- [.eval\(\)](#)

- `conv_to()` (convert between matrix types)
- `Col` class
- `Row` class
- `Cube` class
- `SpMat` class (sparse matrix with compressed sparse column format)
- `config.hpp`
- explanation of *typedef* (cplusplus.com)
- `C` data types (Wikipedia)

Col<type>
vec
cx_vec

- Classes for column vectors (dense matrices with one column)
- The **Col<type>** class is derived from the **Mat<type>** class and inherits most of the member functions
- For convenience the following typedefs have been defined:

```
vec = colvec = Col<double>
dvec = dcolvec = Col<double>
fvec = fcolvec = Col<float>
cx_vec = cx_colvec = Col<cx_double>
cx_dvec = cx_dcolvec = Col<cx_double>
cx_fvec = cx_fcolvec = Col<cx_float>
uvec = ucolvec = Col<uword>
ivec = icolvec = Col<sword>
```

- The **vec** and **colvec** types have the **same meaning** and are used **interchangeably**
- In this documentation, the types *vec* or *colvec* are used for convenience; it is possible to use other column vector types instead, eg. *fvec*, *fcolvec*
- Functions which take *Mat* as input can generally also take *Col* as input; main exceptions are functions which require square matrices
- Constructors:

```
vec()
vec(n_elem)
vec(n_elem, fill_form) (elements are initialised according to fill_form)
vec(size(X))
vec(size(X), fill_form) (elements are initialised according to fill_form)
vec(vec)
```

<code>vec(mat)</code>	(<i>std::logic_error</i> exception is thrown if the given matrix has more than one column)
<code>vec(initializer_list)</code>	
<code>vec(string)</code>	(elements separated by spaces)
<code>vec(std::vector)</code>	
<code>cx_vec(vec, vec)</code>	(for constructing a complex vector out of two real vectors)

- **Caveat:**

- since Armadillo 10.5, the elements are initialised to zero by default
- in Armadillo 10.4 and older versions, the elements are not initialised unless *fill_form* is specified; ie. without specifying *fill_form*, the elements may contain garbage values, including NaN; see the [Mat class](#) for details on *fill_form*

- Advanced constructors:

```
vec(ptr_aux_mem, number_of_elements, copy_aux_mem = true, strict = false)
```

Create a column vector using data from writable auxiliary (external) memory, where *ptr_aux_mem* is a pointer to the memory. By default the vector allocates its own memory and copies data from the auxiliary memory (for safety). However, if *copy_aux_mem* is set to *false*, the vector will instead directly use the auxiliary memory (ie. no copying); this is faster, but can be dangerous unless you know what you are doing!

The *strict* parameter comes into effect only when *copy_aux_mem* is set to *false* (ie. the vector is directly using auxiliary memory)

- when *strict* is set to *false*, the vector will use the auxiliary memory until a size change or an aliasing event
- when *strict* is set to *true*, the vector will be bound to the auxiliary memory for its lifetime; the number of elements in the vector can't be changed

```
vec(const ptr_aux_mem, number_of_elements)
```

Create a column vector by copying data from read-only auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

```
vec::fixed<number_of_elements>
```

Create a fixed size column vector, with the size specified via the template argument. Memory for the vector is reserved at compile time. This is generally faster than dynamic memory allocation, but the size of the vector can't be changed afterwards (directly or indirectly).

For convenience, there are several pre-defined typedefs for each vector type (where the types are: *uvec*, *ivec*, *fvec*, *vec*, *cx_fvec*, *cx_vec* as well as the corresponding *colvec* versions). The pre-defined typedefs specify vector sizes ranging from 2 to 9. The typedefs were defined by appending a single digit form of the size to the vector type; examples: *vec3* is equivalent to *vec::fixed<3>*, while *cx_vec4* is equivalent to *cx_vec::fixed<4>*.

```
vec::fixed<number_of_elements>(fill_form)
```

Create a fixed size column vector, with the elements explicitly initialised according to *fill_form*

```
vec::fixed<number_of_elements>(const ptr_aux_mem)
```

Create a fixed size column vector, with the size specified via the template argument; data is copied from auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

- Examples:

```
vec x(10);  
vec y(10, fill::ones);  
  
mat A(10, 10, fill::randu);  
vec z = A.col(5); // extract a column vector
```

- See also:

- [element initialisation](#)
- [Mat class](#)
- [Row class](#)

Row<type>
rowvec
cx_rowvec

- Classes for row vectors (dense matrices with one row)
- The template **Row<type>** class is derived from the **Mat<type>** class and inherits most of the member functions
- For convenience the following typedefs have been defined:

```
rowvec = Row<double>
drowvec = Row<double>
frowvec = Row<float>
cx_rowvec = Row<cx_double>
cx_drowvec = Row<cx_double>
cx_frowvec = Row<cx_float>
urowvec = Row<uword>
irowvec = Row<sword>
```
- In this documentation, the *rowvec* type is used for convenience; it is possible to use other row vector types instead, eg. *frowvec*
- Functions which take *Mat* as input can generally also take *Row* as input; main exceptions are functions which require square matrices

- Constructors:

```
rowvec()
rowvec(n_elem)
rowvec(n_elem, fill_form)  (elements are initialised according to fill_form)
rowvec(size(X))
rowvec(size(X), fill_form)  (elements are initialised according to fill_form)
rowvec(rowvec)
```

<code>rowvec(mat)</code>	(<code>std::logic_error</code> exception is thrown if the given matrix has more than one row)
<code>rowvec(initializer_list)</code>	
<code>rowvec(string)</code>	(elements separated by spaces)
<code>rowvec(std::vector)</code>	
<code>cx_rowvec(rowvec, rowvec)</code>	(for constructing a complex row vector out of two real row vectors)

- **Caveat:**

- since Armadillo 10.5, the elements are initialised to zero by default
- in Armadillo 10.4 and older versions, the elements are not initialised unless *fill_form* is specified; ie. without specifying *fill_form*, the elements may contain garbage values, including NaN; see the [Mat class](#) for details on *fill_form*

- Advanced constructors:

```
rowvec(ptr_aux_mem, number_of_elements, copy_aux_mem = true, strict = false)
```

Create a row vector using data from writable auxiliary (external) memory, where *ptr_aux_mem* is a pointer to the memory. By default the vector allocates its own memory and copies data from the auxiliary memory (for safety). However, if *copy_aux_mem* is set to *false*, the vector will instead directly use the auxiliary memory (ie. no copying); this is faster, but can be dangerous unless you know what you are doing!

The *strict* parameter comes into effect only when *copy_aux_mem* is set to *false* (ie. the vector is directly using auxiliary memory)

- when *strict* is set to *false*, the vector will use the auxiliary memory until a size change or an aliasing event
- when *strict* is set to *true*, the vector will be bound to the auxiliary memory for its lifetime; the number of elements in the vector can't be changed

```
rowvec(const ptr_aux_mem, number_of_elements)
```

Create a row vector by copying data from read-only auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

```
rowvec::fixed<number_of_elements>
```

Create a fixed size row vector, with the size specified via the template argument. Memory for the vector is reserved at compile time. This is generally faster than dynamic memory allocation, but the size of the vector can't be changed afterwards (directly or indirectly).

For convenience, there are several pre-defined typedefs for each vector type (where the types are: *urowvec*, *irowvec*, *frowvec*, *rowvec*, *cx_frowvec*, *cx_rowvec*). The pre-defined typedefs specify vector sizes ranging from 2 to 9. The typedefs were defined by appending a single digit form of the size to the vector type; examples: *rowvec3* is equivalent to *rowvec::fixed<3>*, while *cx_rowvec4* is equivalent to *cx_rowvec::fixed<4>*.

```
rowvec::fixed<number_of_elements>(fill_form)
```

Create a fixed size row vector, with the elements explicitly initialised according to *fill_form*

```
rowvec::fixed<number_of_elements>(const ptr_aux_mem)
```

Create a fixed size row vector, with the size specified via the template argument; data is copied from auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

- Examples:

```
rowvec x(10);  
rowvec y(10, fill::ones);  
  
mat    A(10, 10, fill::randu);  
rowvec z = A.row(5); // extract a row vector
```

- See also:

- [element initialisation](#)
- [Mat class](#)
- [Col class](#)

Cube<type>
cube
cx_cube

- Classes for cubes (quasi 3rd order tensors), also known as "3D matrices"
- Data is stored as a set of slices (matrices) stored contiguously within memory; within each slice, elements are stored with column-major ordering (ie. column by column)
- The root cube class is **Cube<type>**, where *type* is one of:
 - *float, double, std::complex<float>, std::complex<double>, short, int, long* and unsigned versions of *short, int, long*

- For convenience the following typedefs have been defined:

```
cube = Cube<double>
dcube = Cube<double>
fcube = Cube<float>
cx_cube = Cube<cx_double>
cx_dcube = Cube<cx_double>
cx_fcube = Cube<cx_float>
ucube = Cube<uword>
icube = Cube<sword>
```

- In this documentation the *cube* type is used for convenience; it is possible to use other types instead, eg. *fcube*

- Constructors:

```
cube()
cube(n_rows, n_cols, n_slices)
cube(n_rows, n_cols, n_slices, fill_form) (elements are initialised according to fill_form)
cube(size(X))
cube(size(X), fill_form) (elements are initialised according to fill_form)
cube(cube)
cx_cube(cube, cube) (for constructing a complex cube out of two real cubes)
```

- The elements can be explicitly initialised during construction by specifying *fill_form*, which is one of:

<code>fill::zeros</code>	→ set all elements to 0
<code>fill::ones</code>	→ set all elements to 1
<code>fill::randu</code>	→ set all elements to random values from a uniform distribution in the [0,1] interval
<code>fill::randn</code>	→ set all elements to random values from a normal/Gaussian distribution with zero mean and unit variance
<code>fill::value(scalar)</code>	→ set all elements to specified scalar
<code>fill::none</code>	→ do not initialise the elements

- **Caveat:**

- since Armadillo 10.5, the elements are initialised to zero by default
- in Armadillo 10.4 and older versions, the elements are not initialised unless *fill_form* is specified; ie. without specifying *fill_form*, the elements may contain garbage values, including NaN

- Each instance of *cube* automatically allocates and releases internal memory. All internally allocated memory used by an instance of *cube* is automatically released as soon as the instance goes out of scope. For example, if an instance of *cube* is declared inside a function, it will be automatically destroyed at the end of the function. To forcefully release memory at any point, use `.reset()`; note that in normal use this is not required.

- Advanced constructors:

```
cube::fixed<n_rows, n_cols, n_slices>
```

Create a fixed size cube, with the size specified via template arguments. Memory for the cube is reserved at compile time. This is generally faster than dynamic memory allocation, but the size of the cube can't be changed afterwards (directly or indirectly).

```
cube(ptr_aux_mem, n_rows, n_cols, n_slices, copy_aux_mem = true, strict = false)
```

Create a cube using data from writable auxiliary (external) memory, where *ptr_aux_mem* is a pointer to the memory. By default the cube allocates its own memory and copies data from the auxiliary memory (for safety). However, if *copy_aux_mem* is set to *false*, the cube will instead directly use the auxiliary memory (ie. no copying); this is faster, but can be dangerous unless you

know what you are doing!

The *strict* parameter comes into effect only when *copy_aux_mem* is set to *false* (ie. the cube is directly using auxiliary memory)

- when *strict* is set to *false*, the cube will use the auxiliary memory until a size change or an aliasing event
- when *strict* is set to *true*, the cube will be bound to the auxiliary memory for its lifetime; the number of elements in the cube can't be changed

```
cube(const ptr_aux_mem, n_rows, n_cols, n_slices)
```

Create a cube by copying data from read-only auxiliary memory, where *ptr_aux_mem* is a pointer to the memory

- Examples:

```
cube x(1, 2, 3);
cube y(4, 5, 6, fill::randu);

mat A = y.slice(1); // extract a slice from the cube
                  // (each slice is a matrix)

mat B(4, 5, fill::randu);
y.slice(2) = B;     // set a slice in the cube

cube q = y + y;     // cube addition
cube r = y % y;     // element-wise cube multiplication

cube::fixed<4,5,6> f;
f.ones();
```

- Notes:

- Each cube slice can be interpreted as a matrix, hence functions which take *Mat* as input can generally also take cube slices as input
- The size of individual slices can't be changed. For example, the following **will not** work:

```
cube c(5,6,7);  
c.slice(0) = randu<mat>(10,20); // wrong size
```

- See also:
 - [cube attributes](#)
 - [accessing elements](#)
 - [math & relational operators](#)
 - [subcube views and slices](#)
 - [saving & loading cubes](#)
 - [element iterators](#)
 - [field class](#)
 - [Mat class](#)

field<object_type>

- Class for storing arbitrary objects in matrix-like or cube-like layouts
- Somewhat similar to a matrix or cube, but instead of each element being a scalar, each element can be a vector, or matrix, or cube
- Each element can have an arbitrary size (eg. in a field of matrices, each matrix can have a unique size)

- Constructors, where *object_type* is another class, eg. *vec*, *mat*, *std::string*, etc:

```
field<object_type>()
field<object_type>(n_elem)
field<object_type>(n_rows, n_cols)
field<object_type>(n_rows, n_cols, n_slices)
field<object_type>(size(X))
field<object_type>(field<object_type>)
```

- **Caveat:** to store a set of matrices of the same size, the **Cube** class is more efficient

- Examples:

```
mat A = randn(2,3);
mat B = randn(4,5);

field<mat> F(2,1);
F(0,0) = A;
F(1,0) = B;

F.print("F:");

F.save("mat_field");
```

- See also:
 - **field attributes**
 - **subfield views**
 - **saving / loading fields**
 - **Cube class**

SpMat<type>

sp_mat

sp_cx_mat

- Classes for sparse matrices; intended for storing large matrices, where most of the elements are zeros
- The root sparse matrix class is **SpMat<type>**, where *type* is one of:
 - *float, double, std::complex<float>, std::complex<double>, short, int, long* and unsigned versions of *short, int, long*

- For convenience the following typedefs have been defined:

```
sp_mat = SpMat<double>
sp_dmat = SpMat<double>
sp_fmat = SpMat<float>
sp_cx_mat = SpMat<cx_double>
sp_cx_dmat = SpMat<cx_double>
sp_cx_fmat = SpMat<cx_float>
sp_umat = SpMat<uword>
sp_imat = SpMat<sword>
```

- In this documentation the *sp_mat* type is used for convenience; it is possible to use other types instead, eg. *sp_fmat*

- Constructors:

```
sp_mat()
sp_mat(n_rows, n_cols)
sp_mat(size(X))
sp_mat(sp_mat)
sp_mat(mat) (for converting a dense matrix to a sparse matrix)
sp_cx_mat(sp_mat, sp_mat) (for constructing a complex matrix out of two real matrices)
```

- All elements are treated as zero by default (ie. the matrix is initialised to contain zeros)
- Non-zero elements are stored in **compressed sparse column (CSC) format** (ie. column-major ordering); zero-

valued elements are never stored

- This class behaves in a similar manner to the dense matrix **Mat** class; however, member functions which set all elements to non-zero values (and hence do not make sense for sparse matrices) have been deliberately omitted; examples of omitted functions: `.fill()`, `.ones()`, `+= scalar`, etc.
- **Caveat:** the sparse matrix class is not intended for small matrices (eg. with size $\leq 100 \times 100$), due to the overhead of the compressed storage format; for small matrices use the **Mat** class, even if most of the elements are zeros
- Batch insertion constructors:
 - form 1: `sp_mat(locations, values, sort_locations = true)`
 - form 2: `sp_mat(locations, values, n_rows, n_cols, sort_locations = true, check_for_zeros = true)`
 - form 3: `sp_mat(add_values, locations, values, n_rows, n_cols, sort_locations = true, check_for_zeros = true)`
 - form 4: `sp_mat(rowind, colptr, values, n_rows, n_cols, check_for_zeros = true)`
 - For forms 1, 2, 3, *locations* is a dense matrix of type *umat*, with a size of $2 \times N$, where N is the number of values to be inserted; the location of the i -th element is specified by the contents of the i -th column of the *locations* matrix, where the row is in *locations*(0, i), and the column is in *locations*(1, i)
 - For form 4, *rowind* is a dense column vector of type *uvec* containing the row indices of the values to be inserted, and *colptr* is a dense column vector of type *uvec* (with length $n_cols + 1$) containing indices of *values* corresponding to the start of new columns; the vectors correspond to the arrays used by the **compressed sparse column format**; this form is useful for copying data from other CSC sparse matrix containers
 - For all forms, *values* is a dense column vector containing the values to be inserted; it must have the same element type as the sparse matrix. For forms 1 and 2, the value in *values*[i] will be inserted at the location specified by the i -th column of the *locations* matrix.
 - For form 3, *add_values* is either *true* or *false*; when set to *true*, identical locations are allowed, and the values at identical locations are added
 - The size of the constructed matrix is either automatically determined from the maximal locations in the *locations* matrix (form 1), or manually specified via *n_rows* and *n_cols* (forms 2, 3, 4)

- If *sort_locations* is set to *false*, the *locations* matrix is assumed to contain locations that are already sorted according to column-major ordering; do not set this to *false* unless you know what you are doing!
- If *check_for_zeros* is set to *false*, the *values* vector is assumed to contain no zero values; do not set this to *false* unless you know what you are doing!
- The following subset of operations & functions is available for sparse matrices:
 - fundamental arithmetic **operations** (such as addition and multiplication)
 - **submatrix views**: most contiguous forms and the non-contiguous form of *X.cols(vector_of_column_indices)*
 - **diagonal views**
 - **saving and loading** (using *arma_binary*, *coord_ascii*, and *csv_ascii* formats)
 - element-wise functions: **abs()**, **ceil()**, **conj()**, **floor()**, **imag()**, **real()**, **round()**, **sign()**, **sqrt()**, **square()**, **trunc()**
 - scalar functions of matrices: **accu()**, **as_scalar()**, **dot()**, **norm()**, **norm2est()**, **trace()**
 - vector valued functions of matrices: **diagvec()**, **min()**, **max()**, **nonzeros()**, **sum()**, **mean()**, **var()**, **vecnorm()**, **vectorise()**
 - matrix valued functions of matrices: **clamp()**, **diagmat()**, **spdiags()**, **flipud()/fliplr()**, **join_rows()**, **join_cols()**, **kron()**, **normalise()**, **repelem()**, **repmat()**, **reshape()**, **resize()**, **reverse()**, **symmatu()/symmatl()**, **trimatu()/trimatl()**, **.t()**, **trans()**
 - generated matrices: **speye()**, **spones()**, **sprandu()**, **sprandn()**, **zeros()**
 - eigen decompositions and SVD: **eigs_sym()**, **eigs_gen()**, **svds()**
 - solution of sparse linear systems: **spsolve()**
 - miscellaneous: **approx_equal()**, **element access**, **element iterators**, **.as_col() / .as_row()**, **.for_each()**, **.print()**, **.clean()**, **.replace()**, **.transform()**, **.is_finite()**, **.is_symmetric()**, **.is_hermitian()**, **.is_trimatu()**, **.is_trimatl()**, **.is_diagmat()**

- Examples:

```

sp_mat A = sprandu(1000, 2000, 0.01);
sp_mat B = sprandu(2000, 1000, 0.01);

sp_mat C = 2*B;
sp_mat D = A*C;

sp_mat E(1000,1000);

```

```
E(1,2) = 123;
```

```
// batch insertion of 3 values at  
// locations (1, 2), (7, 8), (9, 9)
```

```
umat locations = { { 1, 7, 9 },  
                  { 2, 8, 9 } };
```

```
vec values = { 1.0, 2.0, 3.0 };
```

```
sp_mat X(locations, values);
```

- See also:
 - [element access](#)
 - [element iterators \(sparse matrices\)](#)
 - [printing matrices](#)
 - [Sparse Matrix in Wikipedia](#)
 - [Mat class](#) (dense matrix)

operators: + - * % / == != <= >= < > && ||

- Overloaded operators for *Mat*, *Col*, *Row* and *Cube* classes
- Operations:
 - + addition of two objects
 - subtraction of one object from another or negation of an object
 - * matrix multiplication of two objects; not applicable to the *Cube* class unless multiplying by a scalar
 - % element-wise multiplication of two objects (Schur product)
 - / element-wise division of an object by another object or a scalar
 - = element-wise equality evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - != element-wise non-equality evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - >= element-wise "greater than or equal to" evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - <= element-wise "less than or equal to" evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - > element-wise "greater than" evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - < element-wise "less than" evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - && element-wise logical AND evaluation of two objects; generates a matrix/cube of type *umat/ucube*
 - || element-wise logical OR evaluation of two objects; generates a matrix/cube of type *umat/ucube*
- For element-wise relational and logical operations (ie. ==, !=, >=, <=, >, <, &&, ||) each element in the generated object is either 0 or 1, depending on the result of the operation

- **Caveat:** operators involving equality comparison (ie. `==`, `!=`, `>=`, `<=`) are not recommended for matrices of type *mat* or *fmat*, due to the necessarily **limited precision of floating-point** element types; consider using **`approx_equal()`** instead
- If the `+`, `-` and `%` operators are chained, Armadillo aims to avoid the generation of temporaries; no temporaries are generated if all given objects are of the same type and size
- If the `*` operator is chained, Armadillo aims to find an efficient ordering of the matrix multiplications
- Broadcasting operations are available via **`.each_col()`**, **`.each_row()`**, **`.each_slice()`**
- If incompatible object sizes are used, a `std::logic_error` exception is thrown

- Examples:

```
mat A(5, 10, fill::randu);
mat B(5, 10, fill::randu);
mat C(10, 5, fill::randu);

mat P = A + B;
mat Q = A - B;
mat R = -B;
mat S = A / 123.0;
mat T = A % B;
mat U = A * C;

// V is constructed without temporaries
mat V = A + B + A + B;

imat AA = "1 2 3; 4 5 6; 7 8 9;";
imat BB = "3 2 1; 6 5 4; 9 8 7;";

// compare elements
umat ZZ = (AA >= BB);
```

- See also:
 - **`approx_equal()`**
 - **`pow()`**
 - **`any()`**

- `all()`
- `affmul()`
- `accu()`
- `as_scalar()`
- `find()`
- `.replace()`
- `.transform()`
- `.each_col()` & `.each_row()` (vector operations applied to each column or row)
- miscellaneous element-wise functions (exp, log, sqrt, square, round, ...)
- floating point arithmetic in Wikipedia
- floating point representation in MathWorld

Member Functions & Variables

attributes

<code>.n_rows</code>	number of rows; present in <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>field</i> and <i>SpMat</i>
<code>.n_cols</code>	number of columns; present in <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>field</i> and <i>SpMat</i>
<code>.n_elem</code>	total number of elements; present in <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>field</i> and <i>SpMat</i>
<code>.n_slices</code>	number of slices; present in <i>Cube</i> and <i>field</i>
<code>.n_nonzero</code>	number of non-zero elements; present in <i>SpMat</i>

- The variables are of type `uword`
- The variables are read-only; to change the size, use `.set_size()`, `.copy_size()`, `.zeros()`, `.ones()`, or `.reset()`
- For the *Col* and *Row* classes, `n_elem` also indicates vector length

- Examples:

```
mat X(4,5);  
cout << "X has " << X.n_cols << " columns" << endl;
```

- See also:

- `.set_size()`
- `.copy_size()`
- `.zeros()`
- `.ones()`
- `.reset()`
- `size()`

element/object access via (), [] and .at()

- Provide access to individual elements or objects stored in a container object (ie. *Mat*, *Col*, *Row*, *Cube*, *field*)

(i) For *vec* and *rowvec*, access the element stored at index *i*. For *mat*, *cube* and *field*, access the element/object stored at index *i* under the assumption of a flat layout, with column-major ordering of data (ie. column by column). An exception is thrown if the requested element is out of bounds.

.at(i) or [i] As for (i), but without a bounds check; not recommended; see the caveats below

(r,c) For *mat* and *2D field* classes, access the element/object stored at row *r* and column *c*. An exception is thrown if the requested element is out of bounds.

.at(r,c) As for (r,c), but without a bounds check; not recommended; see the caveats below

(r,c,s) For *cube* and *3D field* classes, access the element/object stored at row *r*, column *c*, and slice *s*. An exception is thrown if the requested element is out of bounds.

.at(r,c,s) As for (r,c,s), but without a bounds check; not recommended; see the caveats below

- The indices of elements are specified via the `uword` type, which is a typedef for an `unsigned integer type`. When using loops to access elements, it is best to use `uword` instead of `int`. For example:

```
for(uword i=0; i<X.n_elem; ++i) { X(i) = ... }
```

- **Caveats:**

- accessing elements without bounds checks is slightly faster, but is not recommended until your code has been thoroughly debugged first
- indexing in C++ starts at 0
- accessing elements via `[r,c]` and `[r,c,s]` does not work correctly in C++; instead use `(r,c)` and `(r,c,s)`

- **Examples:**

```
mat M(10, 10, fill::randu);
```

```
M(9,9) = 123.0;  
double x = M(1,2);
```

```
vec v(10, fill::randu);  
v(9) = 123.0;  
double y = v(0);
```

- See also:

- `.in_range()`
- element initialisation
- `ind2sub()`
- `sub2ind()`
- `.index_min()` / `.index_max()`
- submatrix views
- `.memptr()`
- `.transform()`
- `.for_each()`
- iterators (dense matrices)
- iterators (cubes)
- iterators (sparse matrices)
- `config.hpp`

element initialisation

- Set elements in *Mat*, *Col*, *Row* via braced initialiser lists

- Examples:

```
vec v = { 1, 2, 3 };
```

```
mat A = { {1, 3, 5},  
          {2, 4, 6} };
```

- See also:
 - [element access](#)
 - [.reshape\(\)](#)
 - [.print\(\)](#)
 - [saving & loading matrices](#)
 - [advanced constructors \(matrices\)](#)

.zeros()	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>SpMat</i> , <i>Cube</i>)
.zeros(n_elem)	(member function of <i>Col</i> and <i>Row</i>)
.zeros(n_rows, n_cols)	(member function of <i>Mat</i> and <i>SpMat</i>)
.zeros(n_rows, n_cols, n_slices)	(member function of <i>Cube</i>)
.zeros(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>SpMat</i>)

- Set the elements of an object to zero, optionally first changing the size to specified dimensions

- Examples:

```
mat A;
A.zeros(5, 10);    // or:  mat A(5, 10, fill::zeros);
```

```
mat B;
B.zeros( size(A) );
```

```
mat C(5, 10, fill::randu);
C.zeros();
```

- See also:

- `zeros()` (standalone function)
- `.ones()`
- `.clean()`
- `.is_zero()`
- `.randu()`
- `.fill()`
- `.imbue()`
- `.reset()`
- `.set_size()`
- `size()`

.ones()	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)
.ones(n_elem)	(member function of <i>Col</i> and <i>Row</i>)
.ones(n_rows, n_cols)	(member function of <i>Mat</i>)
.ones(n_rows, n_cols, n_slices)	(member function of <i>Cube</i>)
.ones(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)

- Set all the elements of an object to one, optionally first changing the size to specified dimensions

- Examples:

```
mat A;
A.ones(5, 10);    // or:  mat A(5, 10, fill::ones);
```

```
mat B;
B.ones( size(A) );
```

```
mat C(5, 10, fill::randu);
C.ones();
```

- See also:

- **ones()** (standalone function)
- **.eye()**
- **.zeros()**
- **.fill()**
- **.imbue()**
- **.randu()**
- **size()**

.eye()
.eye(n_rows, n_cols)
.eye(size(X))

- Member functions of *Mat* and *SpMat*
- Set the elements along the main diagonal to one and off-diagonal elements to zero, optionally first changing the size to specified dimensions
- An identity matrix is generated when $n_rows = n_cols$

- Examples:

```
mat A;  
A.eye(5, 5); // or: mat A(5, 5, fill::eye);
```

```
mat B;  
B.eye( size(A) );
```

```
mat C(5, 5, fill::randu);  
C.eye();
```

- See also:
 - `.ones()`
 - `.diag()`
 - `diagmat()`
 - `diagvec()`
 - `eye()` (standalone function)
 - `size()`

.randu()	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)
.randu(n_elem)	(member function of <i>Col</i> and <i>Row</i>)
.randu(n_rows, n_cols)	(member function of <i>Mat</i>)
.randu(n_rows, n_cols, n_slices)	(member function of <i>Cube</i>)
.randu(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)
.randn()	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)
.randn(n_elem)	(member function of <i>Col</i> and <i>Row</i>)
.randn(n_rows, n_cols)	(member function of <i>Mat</i>)
.randn(n_rows, n_cols, n_slices)	(member function of <i>Cube</i>)
.randn(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i>)

- Set all the elements to random values, optionally first changing the size to specified dimensions
- *.randu()* uses a uniform distribution in the [0,1] interval
- *.randn()* uses a normal/Gaussian distribution with zero mean and unit variance
- Examples:

```
mat A;
A.randu(5, 10);    // or:  mat A(5, 10, fill::randu);

mat B;
B.randu( size(A) );

mat C(5, 10, fill::zeros);
C.randu();
```

- See also:
 - *randu()* (standalone function with extended functionality)
 - *randn()* (standalone function with extended functionality)
 - *.fill()*
 - *.imbue()*
 - *.ones()*

- `.zeros()`
- `size()`
- RNG seed setting
- uniform distribution in Wikipedia
- normal distribution in Wikipedia

.fill(value)

- Member function of *Mat*, *Col*, *Row*, *Cube*, *field*
- Sets the elements to a specified value
- The type of value must match the type of elements used by the container object (eg. for *mat* the type is *double*)

- Examples:

```
mat A(5, 6);
```

```
A.fill(123.0);    // or:  mat A(5, 6, fill::value(123.0));
```

- **Note:** to set all elements to zero during matrix construction, use the following more compact form:

```
mat A(5, 6, fill::zeros);
```

- See also:

- `.imbue()`
- `.ones()`
- `.zeros()`
- `.randu()` & `.randn()`
- `.replace()`
- constants (`pi`, `nan`, `inf`, ...)

.imbue(functor)

.imbue(lambda_function)

- Member functions of *Mat*, *Col*, *Row* and *Cube*
- Imbue (fill) with values provided by a functor or lambda function
- For matrices, filling is done column-by-column (ie. column 0 is filled, then column 1, ...)
- For cubes, filling is done slice-by-slice, with each slice treated as a matrix
- Examples:

```
std::mt19937 engine; // Mersenne twister random number engine

std::uniform_real_distribution<double> distr(0.0, 1.0);

mat A(4, 5, fill::none);

A.imbue( [&]() { return distr(engine); } );
```

- See also:
 - [.fill\(\)](#)
 - [.transform\(\)](#)
 - [element access](#)
 - [function object](#) at Wikipedia
 - [C++11 lambda functions](#) at Wikipedia
 - [lambda function](#) at cprogramming.com

.clean(threshold)

- Member function of *Mat*, *Col*, *Row*, *Cube* and *SpMat*
- For objects with non-complex elements: each element with an absolute value $\leq threshold$ is replaced by zero
- For objects with complex elements: for each element, each component (real and imaginary) with an absolute value $\leq threshold$ is replaced by zero
- Can be used to sparsify a matrix, in the sense of zeroing values with small magnitudes
- **Caveat:** to explicitly convert from dense storage to sparse storage, use the **SpMat class**

- Examples:

```
sp_mat A;  
  
A.sprandu(1000, 1000, 0.01);  
  
A(12,34) = datum::eps;  
A(56,78) = -datum::eps;  
  
A.clean(datum::eps);
```

- See also:

- **.replace()**
- **.clamp()**
- **.transform()**
- **.is_zero()**
- **.zeros()**
- **nonzeros()**
- **datum::eps**

.replace(old_value, new_value)

- Member function of *Mat*, *Col*, *Row*, *Cube* and *SpMat*
- For all elements equal to *old_value*, set them to *new_value*
- The type of *old_value* and *new_value* must match the type of elements used by the container object (eg. for *mat* the type is *double*)
- **Caveats:**
 - floating point numbers (*float* and *double*) are approximations due to their necessarily limited precision
 - for sparse matrices (*SpMat*), replacement is not done when *old_value* = 0

- Examples:

```
mat A(5, 6, fill::randu);  
  
A.diag().fill(datum::nan);  
  
A.replace(datum::nan, 0); // replace each NaN with 0
```

- See also:
 - `.transform()`
 - `.for_each()`
 - `.clean()`
 - `.clamp()`
 - `.fill()`
 - `.has_nan()`
 - `.has_inf()`
 - `find()`
 - relational operators
 - constants (`pi`, `nan`, `inf`, ...)

.clamp(min_value, max_value)

- Member function of *Mat*, *Col*, *Row*, *Cube* and *SpMat*
- Clamp each element to the $[min_val, max_val]$ interval; any value lower than *min_val* will be set to *min_val*, and any value higher than *max_val* will be set to *max_val*
- For complex elements, the real and imaginary components are clamped separately
- For sparse matrices, clamping is applied only to the non-zero elements
- Examples:

```
mat A(5, 6, fill::randu);
```

```
A.clamp(0.2, 0.8);
```

- See also:
 - `.replace()`
 - `.clean()`
 - `.transform()`
 - `.for_each()`
 - `clamp()` (standalone function)
 - `relational operators`

.transform(functor)

.transform(lambda_function)

- Member functions of *Mat*, *Col*, *Row*, *Cube* and *SpMat*
- Transform each element using a functor or lambda function
- For dense matrices, transformation is done column-by-column for all elements
- For sparse matrices, transformation is done column-by-column for non-zero elements
- For cubes, transformation is done slice-by-slice, with each slice treated as a matrix
- Examples:

```
mat A(4, 5, fill::ones);  
  
// add 123 to every element  
A.transform( [](double val) { return (val + 123.0); } );
```

- See also:
 - [.for_each\(\)](#)
 - [.replace\(\)](#)
 - [.imbue\(\)](#)
 - [.clean\(\)](#)
 - [.clamp\(\)](#)
 - [element access](#)
 - [overloaded operators](#)
 - [miscellaneous element-wise functions](#) (exp, log, sqrt, square, round, ...)
 - [function object](#) at Wikipedia
 - [C++11 lambda functions](#) at Wikipedia
 - [lambda function](#) at cprogramming.com

.for_each(functor)
.for_each(lambda_function)

- Member functions of *Mat*, *Col*, *Row*, *Cube*, *SpMat* and *field*
- For each element, pass its reference to a functor or lambda function
- For dense matrices and fields, the processing is done column-by-column for all elements
- For sparse matrices, the processing is done column-by-column for non-zero elements
- For cubes, processing is done slice-by-slice, with each slice treated as a matrix
- Examples:

```
// add 123 to each element in a dense matrix
mat A(4, 5, fill::ones);

A.for_each( [](mat::elem_type& val) { val += 123.0; } ); // NOTE: the '&' is crucial!

// add 123 to each non-zero element in a sparse matrix
sp_mat S; S.sprandu(1000, 2000, 0.1);

S.for_each( [](sp_mat::elem_type& val) { val += 123.0; } ); // NOTE: the '&' is crucial!

// set the size of all matrices in field F
field<mat> F(2,3);

F.for_each( [](mat& X) { X.zeros(4,5); } ); // NOTE: the '&' is crucial!
```

- See also:
 - `.transform()`
 - `.replace()`
 - `.each_col()` & `.each_row()`
 - `.each_slice()`

- [element access](#)
- [miscellaneous element-wise functions](#) (exp, log, sqrt, square, round, ...)
- [function object](#) at Wikipedia
- [C++11 lambda functions](#) at Wikipedia
- [lambda function](#) at cprogramming.com

.set_size(n_elem)	(member function of <i>Col</i> , <i>Row</i> , <i>field</i>)
.set_size(n_rows, n_cols)	(member function of <i>Mat</i> , <i>SpMat</i> , <i>field</i>)
.set_size(n_rows, n_cols, n_slices)	(member function of <i>Cube</i> and <i>field</i>)
.set_size(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>SpMat</i> , <i>field</i>)

- Change the size of an object, without explicitly preserving data and without initialising the elements (ie. elements may contain garbage values, including NaN)
- To initialise the elements to zero while changing the size, use `.zeros()` instead
- To explicitly preserve data while changing the size, use `.reshape()` or `.resize()` instead;
NOTE: `.reshape()` and `.resize()` are considerably slower than `.set_size()`

- Examples:

```
mat A;
A.set_size(5, 10);      // or:  mat A(5, 10, fill::none);

mat B;
B.set_size( size(A) );  // or:  mat B(size(A), fill::none);

vec v;
v.set_size(100);        // or:  vec v(100, fill::none);
```

- See also:

- `.reset()`
- `.copy_size()`
- `.reshape()`
- `.resize()`
- `.zeros()`
- `size()`

.reshape(n_rows, n_cols) (member function of *Mat* and *SpMat*)
.reshape(n_rows, n_cols, n_slices) (member function of *Cube*)
.reshape(size(X)) (member function of *Mat*, *Cube*, *SpMat*)

- Recreate the object according to given size specifications, with the elements taken from the previous version of the object in a column-wise manner; the elements in the generated object are placed column-wise (ie. the first column is filled up before filling the second column)
- The layout of the elements in the recreated object will be different to the layout in the previous version of the object
- If the total number of elements in the previous version of the object is less than the specified size, the extra elements in the recreated object are set to zero
- If the total number of elements in the previous version of the object is greater than the specified size, only a subset of the elements is taken

- **Caveats:**

- to change the size without preserving data, use `.set_size()` instead, which is much faster
- to grow/shrink the object while preserving the elements **as well as** the layout of the elements, use `.resize()` instead
- to flatten a matrix into a vector, use `vectorise()` or `.as_col()` / `.as_row()` instead

- **Examples:**

```
mat A(4, 5, fill::randu);
```

```
A.reshape(5,4);
```

- **See also:**

- `.resize()`
- `.set_size()`
- `.copy_size()`
- `.zeros()`
- `.reset()`
- `.as_col()` / `.as_row()`

- `reshape()` (standalone function)
- `vectorise()`
- `size()`

.resize(n_elem)	(member function of <i>Col</i> , <i>Row</i>)
.resize(n_rows, n_cols)	(member function of <i>Mat</i> and <i>SpMat</i>)
.resize(n_rows, n_cols, n_slices)	(member function of <i>Cube</i>)
.resize(size(X))	(member function of <i>Mat</i> , <i>Col</i> , <i>Row</i> , <i>Cube</i> , <i>SpMat</i>)

- Recreate the object according to given size specifications, while preserving the elements as well as the layout of the elements
- Can be used for growing or shrinking an object (ie. adding/removing rows, and/or columns, and/or slices)
- **Caveat:** to change the size without preserving data, use `.set_size()` instead, which is much faster
- Examples:

```
mat A(4, 5, fill::randu);
```

```
A.resize(7,6);
```

- See also:
 - `.reshape()`
 - `.set_size()`
 - `.copy_size()`
 - `.zeros()`
 - `.reset()`
 - `.insert_rows / cols / slices`
 - `.shed_rows / cols / slices`
 - `resize()` (standalone function)
 - `vectorise()`
 - `size()`

.copy_size(A)

- Set the size to be the same as object A
- Object A must be of the same root type as the object being modified (eg. the size of a matrix can't be set by providing a cube)

- Examples:

```
mat A(5, 6, fill::randu);
```

```
mat B;  
B.copy_size(A);
```

```
cout << B.n_rows << endl;  
cout << B.n_cols << endl;
```

- See also:

- `.reset()`
- `.set_size()`
- `.reshape()`
- `.resize()`
- `.zeros()`
- `size()`

.reset()

- Reset the size to zero (the object will have no elements)

- Examples:

```
mat A(5, 5, fill::randu);  
A.reset();
```

- See also:

- [.set_size\(\)](#)
- [.is_empty\(\)](#)
- [.zeros\(\)](#)

submatrix views

- A collection of member functions of *Mat*, *Col* and *Row* classes that provide read/write access to submatrix views

- contiguous views for matrix X:

X.col(col_number)
X.row(row_number)

X.cols(first_col, last_col)
X.rows(first_row, last_row)

X.submat(first_row, first_col, last_row, last_col)

X(span(first_row, last_row), span(first_col, last_col))

X(first_row, first_col, size(n_rows, n_cols))
X(first_row, first_col, size(Y)) [*Y is a matrix*]

X(span(first_row, last_row), col_number)
X(row_number, span(first_col, last_col))

X.head_cols(number_of_cols)
X.head_rows(number_of_rows)

X.tail_cols(number_of_cols)
X.tail_rows(number_of_rows)

X.unsafe_col(col_number) [*use with caution*]

- contiguous views for vector V:

V(span(first_index, last_index))
V.subvec(first_index, last_index)

- non-contiguous views for matrix or vector X

X.elem(vector_of_indices)
X(vector_of_indices)

X.cols(vector_of_column_indices)
X.rows(vector_of_row_indices)

X.submat(vector_of_row_indices, vector_of_col_indices)
X(vector_of_row_indices, vector_of_col_indices)

- related matrix views (documented separately)

X.diag()
X.each_row()
X.each_col()

V.subvec(first_index, size(W)) [*W is a vector*]

V.head(number_of_elements)

V.tail(number_of_elements)

- Instances of *span(start,end)* can be replaced by *span::all* to indicate the entire range
- For functions requiring one or more vector of indices, eg.
X.submat(vector_of_row_indices, vector_of_column_indices), each vector of indices must be of type *uvec*
- In the function *X.elem(vector_of_indices)*, elements specified in *vector_of_indices* are accessed. *X* is interpreted as one long vector, with column-by-column ordering of the elements of *X*. The *vector_of_indices* must evaluate to a vector of type *uvec* (eg. generated by the *find()* function). The aggregate set of the specified elements is treated as a column vector (ie. the output of *X.elem()* is always a column vector).
- The function *.unsafe_col()* is provided for speed reasons and should be used only if you know what you are doing. It creates a seemingly independent *Col* vector object (eg. *vec*), but uses memory from the existing matrix object. As such, the created vector is not alias safe, and does not take into account that the underlying matrix memory could be freed (eg. due to any operation involving a size change of the matrix).
- Examples:

```
mat A(5, 10, fill::zeros);
```

```
A.submat( 0,1, 2,3 )      = randu<mat>(3,3);
```

```
A( span(0,2), span(1,3) ) = randu<mat>(3,3);
```

```
A( 0,1, size(3,3) )      = randu<mat>(3,3);
```

```
mat B = A.submat( 0,1, 2,3 );
```

```
mat C = A( span(0,2), span(1,3) );
```

```
mat D = A( 0,1, size(3,3) );
```

```
A.col(1)      = randu<mat>(5,1);
```

```
A(span::all, 1) = randu<mat>(5,1);
```

```
mat X(5, 5, fill::randu);
```

```
// get all elements of X that are greater than 0.5
```

```
vec q = X.elem( find(X > 0.5) );
```

```
// add 123 to all elements of X greater than 0.5
```

```
X.elem( find(X > 0.5) ) += 123.0;

// set four specific elements of X to 1
uvec indices = { 2, 3, 6, 8 };

X.elem(indices) = ones<vec>(4);

// add 123 to the last 5 elements of vector a
vec a(10, fill::randu);
a.tail(5) += 123.0;

// add 123 to the first 3 elements of column 2 of X
X.col(2).head(3) += 123;
```

- See also:

- diagonal views
- `.each_col()` & `.each_row()` (vector operations applied to each column or row)
- `.colptr()`
- `.in_range()`
- `find()`
- `join_rows / cols / slices`
- `.shed_rows / cols / slices`
- `.insert_rows / cols / slices`
- `size()`
- subcube views

subcube views and slices

- A collection of member functions of the *Cube* class that provide subcube views

- contiguous views for cube Q:

Q.slice(slice_number)
Q.slices(first_slice, last_slice)

Q.row(row_number)
Q.rows(first_row, last_row)

Q.col(col_number)
Q.cols(first_col, last_col)

Q.subcube(first_row, first_col, first_slice, last_row, last_col, last_slice)

Q(span(first_row, last_row), span(first_col, last_col), span(first_slice, last_slice))

Q(first_row, first_col, first_slice, size(n_rows, n_cols, n_slices))
Q(first_row, first_col, first_slice, size(R)) *[R is a cube]*

Q.head_slices(number_of_slices)
Q.tail_slices(number_of_slices)

Q.tube(row, col)
Q.tube(first_row, first_col, last_row, last_col)
Q.tube(span(first_row, last_row), span(first_col, last_col))
Q.tube(first_row, first_col, size(n_rows, n_cols))

- Instances of *span(a,b)* can be replaced by:

- non-contiguous views for cube Q:

Q.elem(vector_of_indices)
Q(vector_of_indices)

Q.slices(vector_of_slice_indices)

- related cube views (documented separately)

Q.each_slice()
Q.col_as_mat()
Q.row_as_mat()

- `span()` or `span::all`, to indicate the entire range
- `span(a)`, to indicate a particular row, column or slice
- An individual slice, accessed via `.slice()`, is an instance of the `Mat` class (a reference to a matrix is provided)
- All `.tube()` forms are variants of `.subcube()`, using `first_slice = 0` and `last_slice = Q.n_slices-1`
- The `.tube(row,col)` form uses `row = first_row = last_row`, and `col = first_col = last_col`
- In the function `Q.elem(vector_of_indices)`, elements specified in `vector_of_indices` are accessed. `Q` is interpreted as one long vector, with slice-by-slice and column-by-column ordering of the elements of `Q`. The `vector_of_indices` must evaluate to a vector of type `uvec` (eg. generated by the `find()` function). The aggregate set of the specified elements is treated as a column vector (ie. the output of `Q.elem()` is always a column vector).
- In the function `Q.slices(vector_of_slice_indices)`, slices specified in `vector_of_slice_indices` are accessed. The `vector_of_slice_indices` must evaluate to a vector of type `uvec`.

- Examples:

```
cube A(2, 3, 4, fill::randu);

mat B = A.slice(1); // each slice is a matrix

A.slice(0) = randu<mat>(2,3);
A.slice(0)(1,2) = 99.0;

A.subcube(0,0,1, 1,1,2) = randu<cube>(2,2,2);
A( span(0,1), span(0,1), span(1,2) ) = randu<cube>(2,2,2);
A( 0,0,1, size(2,2,2) ) = randu<cube>(2,2,2);

// add 123 to all elements of A greater than 0.5
A.elem( find(A > 0.5) ) += 123.0;

cube C = A.head_slices(2); // get first two slices

A.head_slices(2) += 123.0;
```

- See also:

- `.in_range()`

- `.each_slice()`
- `.col_as_mat()` / `.row_as_mat()`
- `join_slices()`
- `shed_slices()`
- `insert_slices()`
- `size()`
- submatrix views

subfield views

- A collection of member functions of the *field* class that provide subfield views
- For a 2D field F , the subfields are accessed as:

F.row(row_number)
F.col(col_number)

F.rows(first_row, last_row)
F.cols(first_col, last_col)

F.subfield(first_row, first_col, last_row, last_col)

F(span(first_row, last_row), span(first_col, last_col))

F(first_row, first_col, size(G)) [G is a 2D field]
F(first_row, first_col, size(n_rows, n_cols))

- For a 3D field F , the subfields are accessed as:

F.slice(slice_number)

F.slices(first_slice, last_slice)

F.subfield(first_row, first_col, first_slice, last_row, last_col, last_slice)

F(span(first_row, last_row), span(first_col, last_col), span(first_slice, last_slice))

F(first_row, first_col, first_slice, size(G)) [G is a 3D field]
F(first_row, first_col, first_slice, size(n_rows, n_cols, n_slices))

- Instances of $\text{span}(a,b)$ can be replaced by:
 - $\text{span}()$ or $\text{span}::\text{all}$, to indicate the entire range
 - $\text{span}(a)$, to indicate a particular row or column

- See also:
 - `.in_range()`
 - `size()`
 - submatrix views
 - subcube views

.diag() **.diag(k)**

- Member function of *Mat* and *SpMat*
- Read/write access to a diagonal in a matrix
- The argument k is optional; by default $k = 0$
- The argument k specifies the diagonal to use:
 - $k = 0$ indicates the main diagonal (default setting)
 - $k < 0$ indicates the k -th sub-diagonal (below main diagonal, towards bottom-left corner)
 - $k > 0$ indicates the k -th super-diagonal (above main diagonal, towards top-right corner)
- The diagonal is interpreted as a column vector within expressions
- **Note:** to calculate only the diagonal elements of a compound expression, use `diagvec()` or `diagmat()`
- Examples:

```
mat X(5, 5, fill::randu);

vec a = X.diag();
vec b = X.diag(1);
vec c = X.diag(-2);

X.diag() = randu<vec>(5);
X.diag() += 6;
X.diag().ones();

sp_mat S = sprandu<sp_mat>(10,10,0.1);

vec v(S.diag()); // copy sparse diagonal to dense vector
```

- See also:
 - `.eye()`
 - `diagvec()`
 - `diagmat()`
 - `diags()` / `spdiags()`

- submatrix views
- `.each_col()` & `.each_row()`
- `trace()`

.each_col()	.each_row()	(form 1)
.each_col(vector_of_indices)	.each_row(vector_of_indices)	(form 2)
.each_col(lambda_function)	.each_row(lambda_function)	(form 3)

- Member functions of *Mat*
- Apply a vector operation to each column or row of a matrix
- Similar to "broadcasting" in Matlab / Octave
- Supported operations for forms 1 and 2:

+	addition	+=	in-place addition
-	subtraction	-=	in-place subtraction
%	element-wise multiplication	%=	in-place element-wise multiplication
/	element-wise division	/=	in-place element-wise division
=	assignment (copy)		
- For form 2:
 - the argument *vector_of_indices* contains a list of indices of the columns/rows to be used; it must evaluate to a vector of type *uvec*
 - arithmetic operations as per form 1 are supported
- For form 3:
 - apply the given *lambda_function* to each slice
 - the function must accept a reference to a *Col* or *Row* object with the same element type as the underlying matrix

- Examples:

```
mat X(6, 5, fill::ones);
vec v = linspace<vec>(10,15,6);

X.each_col() += v;          // in-place addition of v to each column vector of X

mat Y = X.each_col() + v;   // generate Y by adding v to each column vector of X
```

```
// subtract v from columns 0 through to 3 in X
X.cols(0,3).each_col() -= v;

uvec indices(2);
indices(0) = 2;
indices(1) = 4;

X.each_col(indices) = v;    // copy v to columns 2 and 4 in X

X.each_col( [](vec& a){ a.print(); } );    // lambda function with non-const vector

const mat& XX = X;
XX.each_col( [](const vec& b){ b.print(); } );    // lambda function with const vector
```

- See also:
 - [math & relational operators](#)
 - [submatrix views](#)
 - [diagonal views](#)
 - [repmat\(\)](#)
 - [pow\(\)](#)
 - [.for_each\(\)](#)
 - [.each_slice\(\)](#)

.each_slice() (form 1)
.each_slice(vector_of_indices) (form 2)
.each_slice(lambda_function) (form 3)
.each_slice(lambda_function, use_mp) (form 4)

- Member function of *Cube*
- Apply a matrix operation to each slice of a cube, with each slice treated as a matrix
- Similar to "broadcasting" in Matlab / Octave
- Supported operations for form 1:

+ addition	+= in-place addition
- subtraction	-= in-place subtraction
% element-wise multiplication	%= in-place element-wise multiplication
/ element-wise division	/= in-place element-wise division
* matrix multiplication	*= in-place matrix multiplication
= assignment (copy)	
- For form 2:
 - the argument *vector_of_indices* contains a list of indices of the slices to be used; it must evaluate to a vector of type *uvec*
 - arithmetic operations as per form 1 are supported, except for * and *= (ie. matrix multiplication)
- For form 3:
 - apply the given *lambda_function* to each slice
 - the function must accept a reference to a *Mat* object with the same element type as the underlying cube
- For form 4:
 - apply the given *lambda_function* to each slice, as per form 3
 - the argument *use_mp* is a bool to enable the use of OpenMP for multi-threaded execution of *lambda_function* on multiple slices at the same time
 - the order of processing the slices is not deterministic (eg. slice 2 can be processed before slice 1)

- *lambda_function* must be thread-safe, ie. it must not write to variables outside of its scope

- Examples:

```
cube C(4, 5, 6, fill::randu);

mat M = repmat(linspace<vec>(1,4,4), 1, 5);

C.each_slice() += M;           // in-place addition of M to each slice of C

cube D = C.each_slice() + M;   // generate D by adding M to each slice of C


uvec indices(2);
indices(0) = 2;
indices(1) = 4;

C.each_slice(indices) = M;     // copy M to slices 2 and 4 in C


C.each_slice( [](mat& X){ X.print(); } );    // lambda function with non-const matrix

const cube& CC = C;
CC.each_slice( [](const mat& X){ X.print(); } ); // lambda function with const matrix
```

- See also:

- [math & relational operators](#)
- [subcube views](#)
- [pow\(\)](#)
- [.for_each\(\)](#)
- [.each_col\(\)](#) & [.each_row\(\)](#)
- [lambda function](#) at cprogramming.com

.set_imag(X)

.set_real(X)

- Set the imaginary/real part of an object
- X must have the same size as the recipient object
- Examples:

```
mat A(4, 5, fill::randu);  
mat B(4, 5, fill::randu);  
  
cx_mat C(4, 5, fill::zeros);  
  
C.set_real(A);  
C.set_imag(B);
```

- **Caveat:** to directly construct a complex matrix out of two real matrices, the following code is faster:

```
mat A(4, 5, fill::randu);  
mat B(4, 5, fill::randu);  
  
cx_mat C = cx_mat(A,B);
```

- See also:
 - [matrix constructors](#)
 - [cube constructors](#)
 - [imag\(\) / real\(\)](#)

.insert_rows(row_number, X) (member functions of *Mat*, *Col* and *Cube*)
.insert_rows(row_number, number_of_rows)

.insert_cols(col_number, X) (member functions of *Mat*, *Row* and *Cube*)
.insert_cols(col_number, number_of_cols)

.insert_slices(slice_number, X) (member functions of *Cube*)
.insert_slices(slice_number, number_of_slices)

- Functions with the *X* argument: insert a copy of *X* at the specified row/column/slice
 - if inserting rows, *X* must have the same number of columns (and slices) as the recipient object
 - if inserting columns, *X* must have the same number of rows (and slices) as the recipient object
 - if inserting slices, *X* must have the same number of rows and columns as the recipient object (ie. all slices must have the same size)
- Functions with the *number_of_...* argument:
 - expand the object by creating new rows/columns/slices
 - the elements in the new rows/columns/slices are set to zero
- Examples:

```
mat A(5, 10, fill::randu);  
mat B(5, 2, fill::ones );  
  
// at column 2, insert a copy of B;  
// A will now have 12 columns  
A.insert_cols(2, B);  
  
// at column 1, insert 5 zeroed columns;  
// B will now have 7 columns  
B.insert_cols(1, 5);
```

- See also:
 - [.shed_rows / cols / slices](#)
 - [join_rows / cols / slices](#)
 - [.resize\(\)](#)
 - [submatrix views](#)

- subcube views

.shed_row(row_number)	(member function of <i>Mat</i> , <i>Col</i> , <i>SpMat</i> , <i>Cube</i>)
.shed_rows(first_row, last_row)	(member function of <i>Mat</i> , <i>Col</i> , <i>SpMat</i> , <i>Cube</i>)
.shed_rows(vector_of_indices)	(member function of <i>Mat</i> , <i>Col</i>)

.shed_col(column_number)	(member function of <i>Mat</i> , <i>Row</i> , <i>SpMat</i> , <i>Cube</i>)
.shed_cols(first_column, last_column)	(member function of <i>Mat</i> , <i>Row</i> , <i>SpMat</i> , <i>Cube</i>)
.shed_cols(vector_of_indices)	(member function of <i>Mat</i> , <i>Row</i>)

.shed_slice(slice_number)	(member functions of <i>Cube</i>)
.shed_slices(first_slice, last_slice)	
.shed_slices(vector_of_indices)	

- Functions with single scalar argument: remove the specified row/column/slice
- Functions with two scalar arguments: remove the specified range of rows/columns/slices
- The *vector_of_indices* must evaluate to a vector of type *uvec*; it contains the indices of rows/columns/slices to remove
- Examples:

```
mat A(5, 10, fill::randu);
mat B(5, 10, fill::randu);
```

```
A.shed_row(2);
A.shed_cols(2,4);
```

```
uvec indices = {4, 6, 8};
B.shed_cols(indices);
```

- See also:
 - *.insert_rows / cols / slices*
 - *join_rows / cols / slices*
 - *.resize()*
 - *submatrix views*
 - *subcube views*

.swap_rows(row1, row2)

.swap_cols(col1, col2)

- Member functions of *Mat*, *Col*, *Row* and *SpMat*
- Swap the contents of specified rows or columns
- Examples:

```
mat X(5, 5, fill::randu);  
X.swap_rows(0,4);
```

- See also:
 - `reverse()`
 - `fliplr()` & `flipud()`
 - `.swap()`

.swap(X)

- Member function of *Mat*, *Col*, *Row* and *Cube*
- Swap contents with object *X*
- Examples:

```
mat A(4, 5, fill::zeros);  
mat B(6, 7, fill::ones );
```

```
A.swap(B);
```

- See also:
 - `.swap_rows()` & `.swap_cols()`

.memptr()

- Member function of *Mat*, *Col*, *Row* and *Cube*
- Obtain a raw pointer to the memory used for storing elements
- Data for matrices is stored in a column-by-column order
- Data for cubes is stored in a slice-by-slice (matrix-by-matrix) order
- **Caveat:** the pointer becomes invalid after any operation involving a size change or aliasing
- **Caveat:** this function is not recommended for use unless you know what you are doing!
- Examples:

```
mat A(5, 5, fill::randu);  
const mat B(5, 5, fill::randu);  
  
double* A_mem = A.memptr();  
const double* B_mem = B.memptr();
```

- See also:
 - [.colptr\(\)](#)
 - [submatrix views](#)
 - [element access](#)
 - [iterators \(dense matrices\)](#)
 - [iterators \(cubes\)](#)
 - [advanced constructors \(matrices\)](#)
 - [advanced constructors \(cubes\)](#)

.colptr(col_number)

- Member function of *Mat*
- Obtain a raw pointer to the memory used by elements in the specified column
- **Caveat:** the pointer becomes invalid after any operation involving a size change or aliasing
- **Caveat:** this function is not recommended for use unless you know what you are doing -- it is safer to use **submatrix views** instead
- Examples:

```
mat A(5, 5, fill::randu);  
  
double* mem = A.colptr(2);
```

- See also:
 - **.memptr()**
 - **submatrix views**
 - **element access**
 - **iterators (dense matrices)**
 - **advanced constructors (matrices)**

iterators (dense matrices & vectors)

- Iterators and associated member functions of *Mat*, *Col*, *Row*
- Iterators for dense matrices and vectors traverse over all elements within the specified range
- Member functions:

.begin() iterator referring to the first element

.end() iterator referring to the *past-the-end* element

.begin_col(col_number) iterator referring to the first element of the specified column

.end_col(col_number) iterator referring to the *past-the-end* element of the specified column

.begin_row(row_number) iterator referring to the first element of the specified row

.end_row(row_number) iterator referring to the *past-the-end* element of the specified row

- Iterator types:

mat::iterator random access iterators, for read/write access to elements (which are
vec::iterator stored column by column)
rowvec::iterator

mat::const_iterator random access iterators, for read-only access to elements (which are
vec::const_iterator stored column by column)
rowvec::const_iterator

mat::col_iterator random access iterators, for read/write access to the elements of
vec::col_iterator specified columns
rowvec::col_iterator

mat::const_col_iterator random access iterators, for read-only access to the elements of
vec::const_col_iterator specified columns

rowvec::const_col_iterator

mat::row_iterator bidirectional iterator, for read/write access to the elements of specified rows

mat::const_row_iterator bidirectional iterator, for read-only access to the elements of specified rows

vec::row_iterator random access iterators, for read/write access to the elements of specified rows
rowvec::row_iterator

vec::const_row_iterator random access iterators, for read-only access to the elements of specified rows
rowvec::const_row_iterator

- Examples:

```
mat X(5, 6, fill::randu);

mat::iterator it      = X.begin();
mat::iterator it_end = X.end();

for(; it != it_end; ++it)
{
    cout << (*it) << endl;
}

mat::col_iterator col_it      = X.begin_col(1); // start of column 1
mat::col_iterator col_it_end = X.end_col(3);    // end of column 3

for(; col_it != col_it_end; ++col_it)
{
    cout << (*col_it) << endl;
    (*col_it) = 123.0;
}
```

- See also:

- [Mat class](#)
- [element access](#)

- `.for_each()`
- `.memptr()`
- `.colptr()`
- submatrix views
- iterators (submatrices)
- iterators (cubes)
- iterators (sparse matrices)
- iterator at cplusplus.com

iterators (cubes)

- Iterators and associated member functions of *Cube*
- Iterators for cubes traverse over all elements within the specified range
- Member functions:

.begin() iterator referring to the first element

.end() iterator referring to the *past-the-end* element

.begin_slice(slice_number) iterator referring to the first element of the specified slice

.end_slice(slice_number) iterator referring to the *past-the-end* element of the specified slice

- Iterator types:

cube::iterator random access iterator, for read/write access to elements; the elements are ordered slice by slice; the elements within each slice are ordered column by column

cube::const_iterator random access iterator, for read-only access to elements

cube::slice_iterator random access iterator, for read/write access to the elements of a particular slice; the elements are ordered column by column

cube::const_slice_iterator random access iterator, for read-only access to the elements of a particular slice

- Examples:

```
cube X(2, 3, 4, fill::randu);
```

```
cube::iterator it      = X.begin();
```

```
cube::iterator it_end = X.end();
```

```
for(; it != it_end; ++it)
{
    cout << (*it) << endl;
}

cube::slice_iterator s_it      = X.begin_slice(1); // start of slice 1
cube::slice_iterator s_it_end = X.end_slice(2);    // end of slice 2

for(; s_it != s_it_end; ++s_it)
{
    cout << (*s_it) << endl;
    (*s_it) = 123.0;
}
```

- See also:
 - [Cube class](#)
 - [element access](#)
 - [.for_each\(\)](#)
 - [.memptr\(\)](#)
 - [subcube views](#)
 - [iterators \(subcubes\)](#)
 - [iterators \(dense matrices\)](#)
 - [iterator at cplusplus.com](#)

iterators (sparse matrices)

- Iterators and associated member functions of *SpMat*
- Iterators for sparse matrices traverse over non-zero elements within the specified range
- **Caveats:**
 - writing a zero value into a sparse matrix through an iterator will invalidate all current iterators associated with the sparse matrix
 - to modify the non-zero elements in a safer manner, use `.transform()` or `.for_each()` instead of iterators
- Member functions:

.begin() iterator referring to the first element

.end() iterator referring to the *past-the-end* element

.begin_col(col_number) iterator referring to the first element of the specified column

.end_col(col_number) iterator referring to the *past-the-end* element of the specified column

.begin_row(row_number) iterator referring to the first element of the specified row

.end_row(row_number) iterator referring to the *past-the-end* element of the specified row

- Iterator types:

sp_mat::iterator bidirectional iterator, for read/write access to elements (which are stored column by column)

sp_mat::const_iterator bidirectional iterator, for read-only access to elements (which are stored column by column)

sp_mat::col_iterator bidirectional iterator, for read/write access to the elements of a specific column

sp_mat::const_col_iterator bidirectional iterator, for read-only access to the elements of a specific column

sp_mat::row_iterator bidirectional iterator, for read/write access to the elements of a specific row

sp_mat::const_row_iterator bidirectional iterator, for read-only access to the elements of a specific row

- The iterators have `.row()` and `.col()` functions which return the row and column of the current element; the returned values are of type `uword`
- Examples:

```
sp_mat X = sprandu<sp_mat>(1000, 2000, 0.1);

sp_mat::const_iterator it = X.begin();
sp_mat::const_iterator it_end = X.end();

for(; it != it_end; ++it)
{
    cout << "val: " << (*it) << endl;
    cout << "row: " << it.row() << endl;
    cout << "col: " << it.col() << endl;
}
```

- See also:
 - `SpMat` class
 - element access
 - `.transform()`
 - `.for_each()`
 - `.replace()`
 - submatrix views
 - iterators (dense matrices)
 - iterator at cplusplus.com

iterators (dense submatrices & subcubes)

- iterators for dense [submatrix](#) and [subcube](#) views, allowing range-based for loops
- **Caveat:** These iterators are intended **only** to be used with range-based for loops. Any other use is not supported. For example, the direct use of the *begin()* and *end()* functions, as well as the underlying iterators types is not supported. The implementation of submatrices and subcubes uses short-lived temporary objects that are subject to automatic deletion, and as such are error-prone to handle manually.

- Examples:

```
mat X(100, 200, fill::randu);

for( double& val : X(span(40,60), span(50,100)) )
{
    cout << val << endl;
    val = 123.0;
}
```

- See also:
 - [submatrix views](#)
 - [.for_each\(\)](#)
 - [iterators \(dense matrices\)](#)
 - [iterators \(cubes\)](#)
 - [range-based for](#) (cppreference.com)

compatibility container functions

- Member functions to mimic the functionality of containers in the C++ standard library:

.front() access the first element in a vector

.back() access the last element in a vector

.clear() causes an object to have no elements

.empty() returns *true* if the object has no elements; returns *false* if the object has one or more elements

.size() returns the total number of elements

- Examples:

```
mat A(5, 5, fill::randu);  
cout << A.size() << endl;
```

```
A.clear();  
cout << A.empty() << endl;
```

- See also:
 - [iterators \(dense matrices\)](#)
 - [iterators \(cubes\)](#)
 - [iterators \(sparse matrices\)](#)
 - [matrix and vector attributes](#)
 - [.is_empty\(\)](#)
 - [.reset\(\)](#)

.as_col()
.as_row()

- Member functions of any matrix expression
- `.as_col()`: return a flattened version of the matrix as a column vector; flattening is done by concatenating all columns
- `.as_row()`: return a flattened version of the matrix as a row vector; flattening is done by concatenating all rows
- **Caveat:** concatenating columns is faster than concatenating rows
- Examples:

```
mat X(4, 5, fill::randu);  
vec v = X.as_col();
```

- See also:
 - `.reshape()`
 - `.t()` / `.st()`
 - `as_scalar()`
 - `vectorise()`

.col_as_mat(col_number)
.row_as_mat(row_number)

- Member functions of any cube expression
- *.col_as_mat(col_number)*:
 - return a matrix representation of the specified cube column
 - the number of rows is preserved
 - given a cube with size $R \times C \times S$, the resultant matrix size is $R \times S$
- *.row_as_mat(row_number)*:
 - return a matrix representation of the specified cube row
 - the number of columns is preserved
 - given a cube with size $R \times C \times S$, the resultant matrix size is $S \times C$
- Examples:

```
cube Q(5, 4, 3, fill::randu);  
  
mat A = Q.col_as_mat(2);  // size of A: 5x3  
  
mat B = Q.row_as_mat(2);  // size of B: 3x4
```
- See also:
 - *.slice()*
 - *vectorise()*

.t() **.st()**

- Member functions of any matrix or vector expression
- For real (non-complex) matrix:
 - `.t()` provides a transposed copy of the matrix
 - `.st()` is not applicable
- For complex matrix:
 - `.t()` provides a Hermitian (conjugate) transposed copy (ie. signs of imaginary components are flipped)
 - `.st()` provides a simple transposed copy (ie. signs of imaginary components are not flipped)

- Examples:

```
mat A(4, 5, fill::randu);  
mat B = A.t();
```

- See also:
 - `trans()`
 - `reverse()`
 - `.as_col()` / `.as_row()`
 - [transpose in Wikipedia](#)
 - [transpose in MathWorld](#)
 - [conjugate transpose in Wikipedia](#)
 - [conjugate transpose in MathWorld](#)

.i()

- Member function of any matrix expression
- Provides an inverse of the matrix expression
- If the matrix expression is not square sized, a *std::logic_error* exception is thrown
- If the matrix expression appears to be singular, the output matrix is reset and a *std::runtime_error* exception is thrown
- **Caveats:**
 - if matrix *A* is known to be symmetric positive definite, it is faster to use `inv_sympd()` instead
 - to solve a system of linear equations, such as $Z = \text{inv}(X)*Y$, using `solve()` can be faster and/or more accurate
- Examples:

```
mat A(4, 4, fill::randu);  
  
mat X = A.i();  
  
mat Y = (A+A).i();
```
- See also:
 - `inv()`
 - `rcond()`
 - `pinv()`
 - `solve()`

.min() **.max()**

- Return the extremum value of any matrix or cube expression
- For objects with complex numbers, absolute values are used for comparison
- Examples:

```
mat A(5, 5, fill::randu);
```

```
double max_val = A.max();
```

- See also:
 - `.index_min()` & `.index_max()`
 - `min()` & `max()` (standalone functions with extended functionality)
 - `clamp()`
 - `running_stat`
 - `running_stat_vec`

.index_min() **.index_max()**

- Return the linear index of the extremum value of any matrix or cube expression
- For objects with complex numbers, absolute values are used for comparison
- The returned index is of type **uword**

- Examples:

```
mat A(5, 5, fill::randu);
```

```
uword i = A.index_max();
```

```
double max_val = A(i);
```

- See also:
 - **.min()** & **.max()**
 - **index_min()** & **index_max()** (standalone functions with extended functionality)
 - **ind2sub()**
 - **sort_index()**
 - **find()**
 - **element access**

.eval()

- Member function of any matrix or vector expression
- Explicitly forces the evaluation of a delayed expression and outputs a matrix
- This function should be used sparingly and only in cases where it is absolutely necessary; indiscriminate use can degrade performance

- Examples:

```
cx_mat A( randu<mat>(4,4), randu<mat>(4,4) );
```

```
real(A).eval().save("A_real.dat", raw_ascii);
```

```
imag(A).eval().save("A_imag.dat", raw_ascii);
```

- See also:
 - [as_scalar\(\)](#)
 - [Mat class](#)

.in_range(i)	(member of <i>Mat, Col, Row, Cube, SpMat, field</i>)
.in_range(span(start, end))	(member of <i>Mat, Col, Row, Cube, SpMat, field</i>)
.in_range(row, col)	(member of <i>Mat, Col, Row, SpMat, field</i>)
.in_range(span(start_row, end_row), span(start_col, end_col))	(member of <i>Mat, Col, Row, SpMat, field</i>)
.in_range(row, col, slice)	(member of <i>Cube</i> and <i>field</i>)
.in_range(span(start_row, end_row), span(start_col, end_col), span(start_slice, end_slice))	(member of <i>Cube</i> and <i>field</i>)
.in_range(first_row, first_col, size(X)) (<i>X is a matrix or field</i>)	(member of <i>Mat, Col, Row, SpMat, field</i>)
.in_range(first_row, first_col, size(n_rows, n_cols))	(member of <i>Mat, Col, Row, SpMat, field</i>)
.in_range(first_row, first_col, first_slice, size(Q)) (<i>Q is a cube or field</i>)	(member of <i>Cube</i> and <i>field</i>)
.in_range(first_row, first_col, first_slice, size(n_rows, n_cols, n_slices))	(member of <i>Cube</i> and <i>field</i>)

- Returns *true* if the given location or span is currently valid
- Returns *false* if the object is empty, the location is out of bounds, or the span is out of bounds
- Instances of *span(a,b)* can be replaced by:
 - *span()* or *span::all*, to indicate the entire range
 - *span(a)*, to indicate a particular row, column or slice

- Examples:

```
mat A(4, 5, fill::randu);

cout << A.in_range(0,0) << endl; // true
cout << A.in_range(3,4) << endl; // true
cout << A.in_range(4,5) << endl; // false
```

- See also:
 - [element access](#)
 - [submatrix views](#)
 - [subcube views](#)
 - [subfield views](#)
 - [.set_size\(\)](#)

.is_empty()

- Returns *true* if the object has no elements
- Returns *false* if the object has one or more elements

- Examples:

```
mat A(5, 5, fill::randu);  
cout << A.is_empty() << endl;
```

```
A.reset();  
cout << A.is_empty() << endl;
```

- See also:

- `.is_square()`
- `.is_vec()`
- `.is_finite()`
- `.reset()`

.is_vec()
.is_colvec()
.is_rowvec()

- Member functions of *Mat* and *SpMat*
- `.is_vec()`:
 - returns *true* if the matrix can be interpreted as a vector (either column or row vector)
 - returns *false* if the matrix does not have exactly one column or one row
- `.is_colvec()`:
 - returns *true* if the matrix can be interpreted as a column vector
 - returns *false* if the matrix does not have exactly one column
- `.is_rowvec()`:
 - returns *true* if the matrix can be interpreted as a row vector
 - returns *false* if the matrix does not have exactly one row
- **Caveat:** do not assume that the vector has elements if these functions return *true*; it is possible to have an empty vector (eg. `0x1`)
- Examples:

```
mat A(1, 5, fill::randu);
mat B(5, 1, fill::randu);
mat C(5, 5, fill::randu);

cout << A.is_vec() << endl;
cout << B.is_vec() << endl;
cout << C.is_vec() << endl;
```
- See also:
 - `.is_empty()`
 - `.is_square()`
 - `.is_finite()`

.is_sorted()

.is_sorted(sort_direction)

.is_sorted(sort_direction, dim)

- Member function of *Mat*, *Row* and *Col*
- If the object is a vector, return a *bool* indicating whether the elements are sorted
- If the object is a matrix, return a *bool* indicating whether the elements are sorted in each column (*dim* = 0), or each row (*dim* = 1)
- The *sort_direction* argument is optional; *sort_direction* is one of:
 - "ascend" ↪ elements are ascending; consecutive elements can be equal; this is the **default operation**
 - "descend" ↪ elements are descending; consecutive elements can be equal
 - "strictascend" ↪ elements are strictly ascending; consecutive elements cannot be equal
 - "strictdescend" ↪ elements are strictly descending; consecutive elements cannot be equal
- The *dim* argument is optional; by default *dim* = 0 is used
- For matrices and vectors with complex numbers, order is checked via absolute values

- Examples:

```
vec a(10, fill::randu);  
vec b = sort(a);
```

```
bool check1 = a.is_sorted();  
bool check2 = b.is_sorted();
```

```
mat A(10, 10, fill::randu);
```

```
// check whether each column is sorted in descending manner  
cout << A.is_sorted("descend") << endl;
```

```
// check whether each row is sorted in ascending manner  
cout << A.is_sorted("ascend", 1) << endl;
```

- See also:
 - `sort()`
 - `sort_index()`

.is_trimatu()

.is_trimatl()

- Member functions of *Mat* and *SpMat*
- *.is_trimatu()*:
 - return *true* if the matrix is upper triangular, ie. the matrix is square sized and all elements below the main diagonal are zero; return *false* otherwise
 - **caveat:** if this function returns *true*, do not assume that the matrix contains non-zero elements on or above the main diagonal
- *.is_trimatl()*:
 - return *true* if the matrix is lower triangular, ie. the matrix is square sized and all elements above the main diagonal are zero; return *false* otherwise
 - **caveat:** if this function returns *true*, do not assume that the matrix contains non-zero elements on or below the main diagonal

- Examples:

```
mat A(5, 5, fill::randu);  
mat B = trimatl(A);  
  
cout << A.is_trimatu() << endl;  
cout << B.is_trimatl() << endl;
```

- See also:
 - [trimatu\(\) / trimatl\(\)](#)
 - [.is_symmetric\(\)](#)
 - [.is_diagmat\(\)](#)
 - [Triangular matrix in MathWorld](#)
 - [Triangular matrix in Wikipedia](#)

.is_diagmat()

- Member function of *Mat* and *SpMat*
- Return *true* if the matrix is diagonal, ie. all elements outside of the main diagonal are zero
- Return *false* otherwise
- **Caveat:** if this function returns *true*, do not assume that the matrix contains non-zero elements on the main diagonal
- Examples:

```
mat A(5, 5, fill::randu);  
mat B = diagmat(A);  
  
cout << A.is_diagmat() << endl;  
cout << B.is_diagmat() << endl;
```

- See also:
 - [diagmat\(\)](#)
 - [.is_trimatu\(\)](#) / [.is_trimatl\(\)](#)
 - [.is_symmetric\(\)](#)
 - [Diagonal matrix in MathWorld](#)
 - [Diagonal matrix in Wikipedia](#)

.is_square()

- Member function of *Mat* and *SpMat*
- Returns *true* if the matrix is square, ie. number of rows is equal to the number of columns
- Returns *false* if the matrix is not square
- Examples:

```
mat A(5, 5, fill::randu);  
mat B(6, 7, fill::randu);
```

```
cout << A.is_square() << endl;  
cout << B.is_square() << endl;
```

- See also:
 - [.is_symmetric\(\)](#)
 - [.is_hermitian\(\)](#)
 - [.is_empty\(\)](#)
 - [.is_vec\(\)](#)
 - [.is_finite\(\)](#)

.is_symmetric()

.is_symmetric(tol)

- Member function of *Mat* and *SpMat*
- Returns *true* if the matrix is symmetric
- Returns *false* if the matrix is not symmetric
- The *tol* argument is optional; if *tol* is specified, the given matrix *X* is considered symmetric if $\text{norm}(X - X.st(), "inf") / \text{norm}(X, "inf") \leq \text{tol}$

- Examples:

```
mat A(5, 5, fill::randu);  
mat B = A.t() * A;  
  
cout << A.is_symmetric() << endl;  
cout << B.is_symmetric() << endl;
```

- See also:

- [.is_hermitian\(\)](#)
- [.is_sympd\(\)](#)
- [.is_square\(\)](#)
- [.is_trimatu\(\)](#) / [.is_trimatl\(\)](#)
- [Symmetric matrix in Wikipedia](#)
- [Symmetric matrix in MathWorld](#)

.is_hermitian()

.is_hermitian(tol)

- Member function of *Mat* and *SpMat*
- Returns *true* if the matrix is hermitian (self-adjoint)
- Returns *false* if the matrix is not hermitian
- The *tol* argument is optional; if *tol* is specified, the given matrix *X* is considered hermitian if $\text{norm}(X - X.t(), "inf") / \text{norm}(X, "inf") \leq \text{tol}$
- Examples:

```
cx_mat A(5, 5, fill::randu);  
cx_mat B = A.t() * A;  
  
cout << A.is_hermitian() << endl;  
cout << B.is_hermitian() << endl;
```

- See also:
 - [.is_symmetric\(\)](#)
 - [.is_sympd\(\)](#)
 - [.is_square\(\)](#)
 - [Hermitian matrix in Wikipedia](#)
 - [Hermitian matrix in MathWorld](#)

.is_sympd() **.is_sympd(tol)**

- Member function of *Mat* and any dense matrix expression
- Returns *true* if the matrix is symmetric/hermitian positive definite within the tolerance given by *tol*
- Returns *false* otherwise
- The *tol* argument is optional; if *tol* is not specified, by default $tol = 100 * datum::eps * norm(X, "fro")$
- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = A.t() * A;  
  
cout << A.is_sympd() << endl;  
cout << B.is_sympd() << endl;
```

- See also:
 - [.is_symmetric\(\)](#)
 - [.is_hermitian\(\)](#)
 - [datum::eps](#)

.is_zero()

.is_zero(tolerance)

- For objects with non-complex elements: return *true* if each element has an absolute value $\leq tolerance$; return *false* otherwise
- For objects with complex elements: return *true* if for each element, each component (real and imaginary) has an absolute value $\leq tolerance$; return *false* otherwise
- The argument *tolerance* is optional; by default *tolerance* = 0
- Examples:

```
mat A(5, 5, fill::zeros);
```

```
A(0,0) = datum::eps;
```

```
cout << A.is_zero() << endl;
```

```
cout << A.is_zero(datum::eps) << endl;
```

- See also:
 - `.clean()`
 - `all()`
 - `datum::eps`
 - `approx_equal()`

.is_finite()

- Member function of *Mat*, *Col*, *Row*, *Cube*, *SpMat*
- Returns *true* if all elements of the object are finite
- Returns *false* if at least one of the elements of the object is non-finite (\pm infinity or NaN)
- Examples:

```
mat A(5, 5, fill::randu);  
mat B(5, 5, fill::randu);  
  
B(1,1) = datum::inf;  
  
cout << A.is_finite() << endl;  
cout << B.is_finite() << endl;
```

- See also:
 - `.has_inf()`
 - `.has_nan()`
 - `find_finite()`
 - `find_nonfinite()`
 - constants (`pi`, `nan`, `inf`, ...)

.has_inf()

- Member function of *Mat*, *Col*, *Row*, *Cube*, *SpMat*
- Returns *true* if at least one of the elements of the object is \pm infinity
- Returns *false* otherwise
- Examples:

```
mat A(5, 5, fill::randu);  
mat B(5, 5, fill::randu);  
  
B(1,1) = datum::inf;  
  
cout << A.has_inf() << endl;  
cout << B.has_inf() << endl;
```

- See also:
 - [.has_nan\(\)](#)
 - [.replace\(\)](#)
 - [.is_finite\(\)](#)
 - [find_nonfinite\(\)](#)
 - [constants \(pi, nan, inf, ...\)](#)

.has_nan()

- Member function of *Mat*, *Col*, *Row*, *Cube*, *SpMat*
- Returns *true* if at least one of the elements of the object is NaN (not-a-number)
- Returns *false* otherwise
- **Caveat:** NaN is not equal to anything, even itself

- Examples:

```
mat A(5, 5, fill::randu);
mat B(5, 5, fill::randu);

B(1,1) = datum::nan;

cout << A.has_nan() << endl;
cout << B.has_nan() << endl;
```

- See also:
 - `.has_inf()`
 - `.replace()`
 - `.is_finite()`
 - `find_nonfinite()`
 - `constants (pi, nan, inf, ...)`

.print()
.print(header)

.print(stream)
.print(stream, header)

- Member functions of *Mat*, *Col*, *Row*, *SpMat*, *Cube* and *field*
- Print the contents of an object to the *std::cout* stream (default), or a user specified stream, with an optional header string
- Objects can also be printed using the << stream operator
- Elements of a field can only be printed if there is an associated *operator<<* function defined

- Examples:

```
mat A(5, 5, fill::randu);
mat B(6, 6, fill::randu);

A.print();

// print a transposed version of A
A.t().print();

// "B:" is the optional header line
B.print("B:");

cout << A << endl;

cout << "B:" << endl;
cout << B << endl;
```

- See also:
 - [.raw_print\(\)](#)
 - [.brief_print\(\)](#)
 - [saving & loading matrices](#)
 - [initialising elements](#)
 - [output streams](#)

.raw_print()
.raw_print(header)

.raw_print(stream)
.raw_print(stream, header)

- Member functions of *Mat*, *Col*, *Row*, *SpMat* and *Cube*
- Similar to the `.print()` member function, with the difference that no formatting of the output is done; the stream's parameters such as precision, cell width, etc. can be set manually
- If the cell width is set to zero, a space is printed between the elements
- Examples:

```
mat A(5, 5, fill::randu);  
  
cout.precision(11);  
cout.setf(ios::fixed);  
  
A.raw_print(cout, "A:");
```

- See also:
 - `.print()`
 - `.brief_print()`
 - `std::ios_base::fmtflags` (cppreference.com)
 - `std::ios_base::fmtflags` (cplusplus.com)

.brief_print()
.brief_print(header)

.brief_print(stream)
.brief_print(stream, header)

- Member functions of *Mat*, *Col*, *Row*, *SpMat* and *Cube*
- Print a shortened/abridged version of an object and its size to the *std::cout* stream (default), or a user specified stream, with an optional header string
- Can be useful for gaining a glimpse of large matrices
- Examples:

```
mat A(123, 456, fill::randu);

A.brief_print("A:");

// possible output:
//
// A:
// [matrix size: 123x456]
//   0.8402   0.7605   0.6218   ...   0.9744
//   0.3944   0.9848   0.0409   ...   0.7799
//   0.7831   0.9350   0.4140   ...   0.8835
//       :       :       :       :       :
//   0.4954   0.1826   0.9848   ...   0.1918
```

- See also:
 - [.print\(\)](#)
 - [.raw_print\(\)](#)

saving / loading matrices & cubes

.save(filename)

.save(filename, file_type)

.save(stream)

.save(stream, file_type)

.save(hdf5_name(filename, dataset))

.save(hdf5_name(filename, dataset, settings))

.save(csv_name(filename, header))

.save(csv_name(filename, header, settings))

.load(filename)

.load(filename, file_type)

.load(stream)

.load(stream, file_type)

.load(hdf5_name(filename, dataset))

.load(hdf5_name(filename, dataset, settings))

.load(csv_name(filename, header))

.load(csv_name(filename, header, settings))

- Member functions of *Mat*, *Col*, *Row*, *Cube* and *SpMat*
- Store/retrieve data in a file or stream (**caveat:** the stream must be opened in binary mode)
- On success, *.save()* and *.load()* return a *bool* set to *true*
- On failure, *.save()* and *.load()* return a *bool* set to *false*; additionally, *.load()* resets the object so that it has no elements
- *file_type* can be one of the following:

auto_detect Used only by *.load()* only: attempt to automatically detect the file type as one of the formats described below;
[default operation for *.load()*]

arma_binary Numerical data stored in machine dependent binary format, with a simple header to speed up loading. The header indicates the type and size of matrix/cube.
[default operation for *.save()*]

- arma_ascii** Numerical data stored in human readable text format, with a simple header to speed up loading. The header indicates the type and size of matrix/cube.
- raw_binary** Numerical data stored in machine dependent raw binary format, without a header. Matrices are loaded to have one column, while cubes are loaded to have one slice with one column. The `.reshape()` function can be used to alter the size of the loaded matrix/cube without losing data.
- raw_ascii** Numerical data stored in raw ASCII format, without a header. The numbers are separated by whitespace. The number of columns must be the same in each row. Cubes are loaded as one slice. Data which was saved in Matlab/Octave using the `-ascii` option can be read in Armadillo, except for complex numbers. Complex numbers are stored in standard C++ notation, which is a tuple surrounded by brackets: eg. `(1.23,4.56)` indicates $1.24 + 4.56i$.
- csv_ascii** Numerical data stored in comma separated value (CSV) text format, **without** a header. To save/load **with** a header, use the **csv_name(filename,header)** specification instead (more details below). Handles complex numbers stored in the compound form of $1.24+4.56i$. Applicable to *Mat* and *SpMat*.
- coord_ascii** Numerical data stored as a text file in coordinate list format, without a header. Only non-zero values are stored.
For real matrices, each line contains information in the following format: row column value
For complex matrices, each line contains information in the following format: row column real_value imag_value
The rows and columns start at zero.
Armadillo ≥ 10.3 : applicable to *Mat* and *SpMat*; Armadillo ≤ 10.2 : applicable to *SpMat* only.
Caveat: not supported by *auto_detect*.

pgm_binary Image data stored in Portable Gray Map (PGM) format. Applicable to *Mat* only. Saving *int*, *float* or *double* matrices is a lossy operation, as each element is copied and converted to an 8 bit representation. As such the matrix should have values in the [0,255] interval, otherwise the resulting image may not display correctly.

ppm_binary Image data stored in Portable Pixel Map (PPM) format. Applicable to *Cube* only. Saving *int*, *float* or *double* matrices is a lossy operation, as each element is copied and converted to an 8 bit representation. As such the cube/field should have values in the [0,255] interval, otherwise the resulting image may not display correctly.

hdf5_binary Numerical data stored in portable **HDF5** binary format.

- for saving, the default dataset name within the HDF5 file is "dataset"
- for loading, the order of operations is: (1) try loading a dataset named "dataset", (2) try loading a dataset named "value", (3) try loading the first available dataset
- to explicitly control the dataset name, specify it via the **hdf5_name()** argument (more details below)

- **Caveat:** for saving / loading HDF5 files, support for HDF5 must be enabled within Armadillo's **configuration**; the *hdf5.h* header file must be available on your system and you will need to link with the HDF5 library (eg. *-lhdf5*). HDF5 support can be enabled by defining **ARMA_USE_HDF5** before including the armadillo header:

```
#define ARMA_USE_HDF5
#include <armadillo>
```

- By providing either **hdf5_name(filename, dataset)** or **hdf5_name(filename, dataset, settings)**, the *file_type* type is assumed to be *hdf5_binary*
 - the *dataset* argument specifies an HDF5 dataset name (eg. "my_dataset") that can include a full path (eg. "/group_name/my_dataset"); if a blank dataset name is specified (ie. ""), it is assumed to be "dataset"
 - the *settings* argument is optional; it is one of the following, or a combination thereof:

<code>hdf5_opts::trans</code>	save/load the data with columns transposed to rows (and vice versa)
<code>hdf5_opts::append</code>	instead of overwriting the file, append the specified dataset to the file; the specified dataset must not already exist in the file

`hdf5_opts::replace` instead of overwriting the file, replace the specified dataset in the file
caveat: HDF5 may not automatically reclaim deleted space; use `h5repack` to clean HDF5 files

the above settings can be combined using the + operator; for example:

```
hdf5_opts::trans + hdf5_opts::append
```

- By providing either **`csv_name(filename, header)`** or **`csv_name(filename, header, settings)`**, the file is assumed to have data in comma separated value (CSV) text format
 - the *header* argument specifies the object which stores the separate elements of the header line; it must have the type `field<std::string>`
 - the optional *settings* argument is one of the following, or a combination thereof:

<code>csv_opts::trans</code>	save/load the data with columns transposed to rows (and vice versa)
<code>csv_opts::no_header</code>	assume there is no header line; the <i>header</i> argument is not referenced
<code>csv_opts::semicolon</code>	use semicolon (;) instead of comma (,) as the separator character
<code>csv_opts::strict</code>	interpret missing values as NaN (not applicable to sparse matrices)

the above settings can be combined using the + operator; for example:

```
csv_opts::trans + csv_opts::no_header
```

- Examples:

```
mat A(5, 6, fill::randu);

// default save format is arma_binary
A.save("A.bin");

// save in raw_ascii format
A.save("A.txt", raw_ascii);

// save in CSV format without a header
A.save("A.csv", csv_ascii);

// save in CSV format with a header
field<std::string> header(A.n_cols);
header(0) = "foo";
```

```

header(1) = "bar"; // etc
A.save( csv_name("A.csv", header) );

// save in HDF5 format with internal dataset named as "my_data"
A.save(hdf5_name("A.h5", "my_data"));

// automatically detect format type while loading
mat B;
B.load("A.bin");

// force loading in arma_ascii format
mat C;
C.load("A.txt", arma_ascii);

// example of testing for success
mat D;
bool ok = D.load("A.bin");

if(ok == false)
{
    cout << "problem with loading" << endl;
}

```

- See also:
 - [HDF](#) in Wikipedia
 - [CSV](#) in Wikipedia
 - [saving / loading fields](#)

saving / loading fields

<code>.save(name)</code>	<code>.load(name)</code>
<code>.save(name, file_type)</code>	<code>.load(name, file_type)</code>
<code>.save(stream)</code>	<code>.load(stream)</code>
<code>.save(stream, file_type)</code>	<code>.load(stream, file_type)</code>

- Store/retrieve data in a file or stream (**caveat:** the stream must be opened in binary mode)
- On success, `.save()` and `.load()` return a *bool* set to *true*
- On failure, `.save()` and `.load()` return a *bool* set to *false*; additionally, `.load()` resets the object so that it has no elements
- Fields with objects of type *std::string* are saved and loaded as raw text files. The text files do not have a header. Each string is separated by a whitespace. `load()` will only accept text files that have the same number of strings on each line. The strings can have variable lengths.
- Other than storing string fields as text files, the following file formats are supported:

auto_detect

- `.load()`: attempt to automatically detect the field format type as one of the formats described below; this is the default operation

arma_binary

- objects are stored in machine dependent binary format
- default type for fields of type *Mat*, *Col*, *Row* or *Cube*
- only applicable to fields of type *Mat*, *Col*, *Row* or *Cube*

ppm_binary

- image data stored in Portable Pixmap Map (PPM) format
- only applicable to fields of type *Mat*, *Col* or *Row*
- `.load()`: loads the specified image and stores the red, green and blue components as

three separate matrices; the resulting field is comprised of the three matrices, with the red, green and blue components in the first, second and third matrix, respectively

- `.save()`: saves a field with exactly three matrices of equal size as an image; it is assumed that the red, green and blue components are stored in the first, second and third matrix, respectively; saving *int*, *float* or *double* matrices is a lossy operation, as each matrix element is copied and converted to an 8 bit representation

- See also:

- [saving/loading matrices and cubes](#)

Generated Vectors / Matrices / Cubes

linspace(start, end)
linspace(start, end, N)

- Generate a vector with N elements; the values of the elements are linearly spaced from *start* to (and including) *end*
- The argument N is optional; by default $N = 100$
- Usage:
 - `vec v = linspace(start, end, N)`
 - `vector_type v = linspace<vector_type>(start, end, N)`
- **Caveat:** for $N = 1$, the generated vector will have a single element equal to *end*
- Examples:

```
vec a = linspace(0, 5, 6);  
rowvec b = linspace<rowvec>(5, 0, 6);
```
- See also:
 - `regspace()`
 - `logspace()`
 - `randperm()`
 - `ones()`
 - `interp1()`

logspace(A, B)
logspace(A, B, N)

- Generate a vector with N elements; the values of the elements are logarithmically spaced from 10^A to (and including) 10^B
- The argument N is optional; by default $N = 50$
- Usage:
 - `vec v = logspace(A, B, N)`
 - `vector_type v = logspace<vector_type>(A, B, N)`
- Examples:

```
vec a = logspace(0, 5, 6);  
rowvec b = logspace<rowvec>(5, 0, 6);
```
- See also:
 - `linspace()`
 - `regspace()`

regspace(start, end)

regspace(start, delta, end)

- Generate a vector with regularly spaced elements:
 $[(start + 0*delta), (start + 1*delta), (start + 2*delta), \dots, (start + M*delta)]$
where $M = \text{floor}((end-start) / delta)$, so that $(start + M*delta) \leq end$
- Similar in operation to the Matlab/Octave colon operator, ie. $start:end$ and $start:delta:end$
- If $delta$ is not specified:
 - $delta = +1$, if $start \leq end$
 - $delta = -1$, if $start > end$ (**caveat:** this is different to Matlab/Octave)
- An empty vector is generated when one of the following conditions is true:
 - $start < end$, and $delta < 0$
 - $start > end$, and $delta > 0$
 - $delta = 0$
- Usage:
 - `vec v = regspace(start, end)`
 - `vec v = regspace(start, delta, end)`
 - `vector_type v = regspace<vector_type>(start, end)`
 - `vector_type v = regspace<vector_type>(start, delta, end)`

- Examples:

```
vec a = regspace(0, 9);           // 0, 1, ..., 9
```

```
uvec b = regspace<uvec>(2, 2, 10); // 2, 4, ..., 10
```

```
ivec c = regspace<ivec>(0, -1, -10); // 0, -1, ..., -10
```

- **Caveat:** do not use `regspace()` to specify ranges for contiguous submatrix views; use `span()` instead
- See also:
 - `linspace()`
 - `logspace()`

- `randperm()`

randperm(N)
randperm(N, M)

- Generate a vector with a random permutation of integers from 0 to $N-1$
- The optional argument M indicates the number of elements to return, sampled without replacement from 0 to $N-1$

- Examples:

```
uvec X = randperm(10);  
uvec Y = randperm(10,2);
```

- See also:

- [randi\(\)](#)
- [shuffle\(\)](#)
- [linspace\(\)](#)
- [regspace\(\)](#)
- [RNG seed setting](#)

eye(n_rows, n_cols)
eye(size(X))

- Generate a matrix with the elements along the main diagonal set to one and off-diagonal elements set to zero
- An identity matrix is generated when $n_rows = n_cols$
- Usage:
 - `mat X = eye(n_rows, n_cols)`
 - `matrix_type X = eye<matrix_type>(n_rows, n_cols)`
 - `matrix_type Y = eye<matrix_type>(size(X))`

- Examples:

```
mat A = eye(5,5); // or: mat A(5,5,fill::eye);
```

```
fmat B = 123.0 * eye<fmat>(5,5);
```

```
cx_mat C = eye<cx_mat>( size(B) );
```

- See also:
 - `.eye()` (member function of Mat)
 - `.diag()`
 - `ones()`
 - `diagmat()`
 - `diagvec()`
 - `speye()`
 - `size()`

ones(n_elem)
ones(n_rows, n_cols)
ones(n_rows, n_cols, n_slices)
ones(size(X))

- Generate a vector, matrix or cube with all elements set to one
- Usage:
 - `vector_type v = ones<vector_type>(n_elem)`
 - `matrix_type X = ones<matrix_type>(n_rows, n_cols)`
 - `matrix_type Y = ones<matrix_type>(size(X))`
 - `cube_type Q = ones<cube_type>(n_rows, n_cols, n_slices)`
 - `cube_type R = ones<cube_type>(size(Q))`
- **Caveat:** specifying `fill::ones` during object construction is more compact, eg. `mat A(5, 6, fill::ones)`

- Examples:

```
vec v = ones(10);      // or: vec v(10, fill::ones);
uvec u = ones<uvec>(10);
rowvec r = ones<rowvec>(10);

mat A = ones(5,6);      // or: mat A(5, 6, fill::ones);
fmat B = ones<fmat>(5,6);
umat C = ones<umat>(5,6);

cube Q = ones(5,6,7);   // or: cube Q(5, 6, 7, fill::ones);
fcube R = ones<fcube>(5,6,7);
```

- See also:
 - `.ones()` (member function of *Mat*, *Col*, *Row* and *Cube*)
 - `.fill()`
 - `eye()`
 - `linspace()`
 - `regspace()`
 - `zeros()`
 - `randu()`
 - `spones()`

- `size()`

zeros(n_elem)
zeros(n_rows, n_cols)
zeros(n_rows, n_cols, n_slices)
zeros(size(X))

- Generate a vector, matrix or cube with the elements set to zero
- Usage:
 - `vector_type v = zeros<vector_type>(n_elem)`
 - `matrix_type X = zeros<matrix_type>(n_rows, n_cols)`
 - `matrix_type Y = zeros<matrix_type>(size(X))`
 - `cube_type Q = zeros<cube_type>(n_rows, n_cols, n_slices)`
 - `cube_type R = zeros<cube_type>(size(Q))`
- **Caveat:** specifying `fill::zeros` during object construction is more compact, eg. `mat A(5, 6, fill::zeros)`

- Examples:

```
vec v = zeros(10);      // or: vec v(10, fill::zeros);
uvec u = zeros<uvec>(10);
rowvec r = zeros<rowvec>(10);

mat A = zeros(5,6);      // or: mat A(5, 6, fill::zeros);
fmat B = zeros<fmat>(5,6);
umat C = zeros<umat>(5,6);

cube Q = zeros(5,6,7);   // or: cube Q(5, 6, 7, fill::zeros);
fcube R = zeros<fcube>(5,6,7);
```

- See also:
 - `.zeros()` (member function of *Mat*, *Col*, *Row*, *SpMat* and *Cube*)
 - `.fill()`
 - `ones()`
 - `randu()`
 - `size()`

randu()
randu(distr_param(a,b))

randu(n_elem)
randu(n_elem, distr_param(a,b))

randu(n_rows, n_cols)
randu(n_rows, n_cols, distr_param(a,b))

randu(n_rows, n_cols, n_slices)
randu(n_rows, n_cols, n_slices, distr_param(a,b))

randu(size(X))
randu(size(X), distr_param(a,b))

- Generate a scalar, vector, matrix or cube with the elements set to random **floating point** values uniformly distributed in the $[a,b]$ interval
- The default distribution parameters are $a = 0$ and $b = 1$
- Usage:
 - *scalar_type* s = randu<scalar_type>(), where *scalar_type* \in { float, double, **cx_float**, **cx_double** }
 - *scalar_type* s = randu<scalar_type>(distr_param(a,b)), where *scalar_type* \in { float, double, **cx_float**, **cx_double** }
 - *vector_type* v = randu<vector_type>(n_elem)
 - *vector_type* v = randu<vector_type>(n_elem, distr_param(a,b))
 - *matrix_type* X = randu<matrix_type>(n_rows, n_cols)
 - *matrix_type* X = randu<matrix_type>(n_rows, n_cols, distr_param(a,b))
 - *cube_type* Q = randu<cube_type>(n_rows, n_cols, n_slices)
 - *cube_type* Q = randu<cube_type>(n_rows, n_cols, n_slices, distr_param(a,b))
- **Caveat:** to generate a matrix with random integer values instead of floating point values, use **randi()** instead

- Examples:

```
double a = randu();  
double b = randu(distr_param(10,20));  
  
vec v1 = randu(5);    // or: vec v1(5, fill::randu);  
vec v2 = randu(5, distr_param(10,20));  
  
rowvec r1 = randu<rowvec>(5);  
rowvec r2 = randu<rowvec>(5, distr_param(10,20));  
  
mat A1 = randu(5, 6); // or: mat A1(5, 6, fill::randu);  
mat A2 = randu(5, 6, distr_param(10,20));  
  
fmat B1 = randu<fmat>(5, 6);  
fmat B2 = randu<fmat>(5, 6, distr_param(10,20));
```

- See also:

- `.randu()` (member function)
- `randn()`
- `randg()`
- `randi()`
- `.imbue()`
- `ones()`
- `zeros()`
- `shuffle()`
- `sprandu()`
- `size()`
- RNG seed setting
- uniform distribution in Wikipedia

randn()
randn(distr_param(mu,sd))

randn(n_elem)
randn(n_elem, distr_param(mu,sd))

randn(n_rows, n_cols)
randn(n_rows, n_cols, distr_param(mu,sd))

randn(n_rows, n_cols, n_slices)
randn(n_rows, n_cols, n_slices, distr_param(mu,sd))

randn(size(X))
randn(size(X), distr_param(mu,sd))

- Generate a scalar, vector, matrix or cube with the elements set to random values with normal / Gaussian distribution, parameterised by mean μ and standard deviation σ
- The default distribution parameters are $\mu = 0$ and $\sigma = 1$
- Usage:
 - $\text{scalar_type } s = \text{randn}<\text{scalar_type}>()$, where $\text{scalar_type} \in \{ \text{float}, \text{double}, \text{cx_float}, \text{cx_double} \}$
 - $\text{scalar_type } s = \text{randn}<\text{scalar_type}>(\text{distr_param}(\mu, \sigma))$, where $\text{scalar_type} \in \{ \text{float}, \text{double}, \text{cx_float}, \text{cx_double} \}$
 - $\text{vector_type } v = \text{randn}<\text{vector_type}>(n_elem)$
 - $\text{vector_type } v = \text{randn}<\text{vector_type}>(n_elem, \text{distr_param}(\mu, \sigma))$
 - $\text{matrix_type } X = \text{randn}<\text{matrix_type}>(n_rows, n_cols)$
 - $\text{matrix_type } X = \text{randn}<\text{matrix_type}>(n_rows, n_cols, \text{distr_param}(\mu, \sigma))$
 - $\text{cube_type } Q = \text{randn}<\text{cube_type}>(n_rows, n_cols, n_slices)$
 - $\text{cube_type } Q = \text{randn}<\text{cube_type}>(n_rows, n_cols, n_slices, \text{distr_param}(\mu, \sigma))$
- Examples:

```
double a = randn();  
double b = randn(distr_param(10,5));  
  
vec v1 = randn(5);    // or: vec v1(5, fill::randn);  
vec v2 = randn(5, distr_param(10,5));  
  
rowvec r1 = randn<rowvec>(5);  
rowvec r2 = randn<rowvec>(5, distr_param(10,5));  
  
mat A1 = randn(5, 6); // or: mat A1(5, 6, fill::randn);  
mat A2 = randn(5, 6, distr_param(10,5));  
  
fmat B1 = randn<fmat>(5, 6);  
fmat B2 = randn<fmat>(5, 6, distr_param(10,5));
```

- See also:
 - `.randn()` (member function)
 - `randu()`
 - `randg()`
 - `randi()`
 - `mvnrnd()`
 - `normpdf()`
 - `.imbue()`
 - `sprandn()`
 - `size()`
 - RNG seed setting
 - normal distribution in Wikipedia

randg()
randg(distr_param(a,b))

randg(n_elem)
randg(n_elem, distr_param(a,b))

randg(n_rows, n_cols)
randg(n_rows, n_cols, distr_param(a,b))

randg(n_rows, n_cols, n_slices)
randg(n_rows, n_cols, n_slices, distr_param(a,b))

randg(size(X))
randg(size(X), distr_param(a,b))

- Generate a scalar, vector, matrix or cube with the elements set to random values from a gamma distribution:

$$p(x | a, b) = \frac{x^{a-1} \exp(-x/b)}{b^a \Gamma(a)}$$

where a is the shape parameter and b is the scale parameter, with constraints $a > 0$ and $b > 0$

- The default distribution parameters are $a = 1$ and $b = 1$
- Usage:
 - *scalar_type* $s = \text{randg}<\text{scalar_type}>()$, where *scalar_type* is either *float* or *double*
 - *scalar_type* $s = \text{randg}<\text{scalar_type}>(\text{distr_param}(a,b))$, where *scalar_type* is either *float* or *double*
 - *vector_type* $v = \text{randg}<\text{vector_type}>(n_elem)$
 - *vector_type* $v = \text{randg}<\text{vector_type}>(n_elem, \text{distr_param}(a,b))$
 - *matrix_type* $X = \text{randg}<\text{matrix_type}>(n_rows, n_cols)$
 - *matrix_type* $X = \text{randg}<\text{matrix_type}>(n_rows, n_cols, \text{distr_param}(a,b))$
 - *cube_type* $Q = \text{randg}<\text{cube_type}>(n_rows, n_cols, n_slices)$

- *cube_type* Q = randg<*cube_type*>(n_rows, n_cols, n_slices, distr_param(a,b))

- Examples:

```
vec v1 = randg(100);  
vec v2 = randg(100, distr_param(2,1));  
  
rowvec r1 = randg<rowvec>(100);  
rowvec r2 = randg<rowvec>(100, distr_param(2,1));  
  
mat A1 = randg(10, 10);  
mat A2 = randg(10, 10, distr_param(2,1));  
  
fmat B1 = randg<fmat>(10, 10);  
fmat B2 = randg<fmat>(10, 10, distr_param(2,1));
```

- See also:

- [randu\(\)](#)
- [randn\(\)](#)
- [randi\(\)](#)
- [chi2rnd\(\)](#)
- [.imbue\(\)](#)
- [size\(\)](#)
- [RNG seed setting](#)
- [gamma distribution in Wikipedia](#)

randi()
randi(distr_param(a,b))

randi(n_elem)
randi(n_elem, distr_param(a,b))

randi(n_rows, n_cols)
randi(n_rows, n_cols, distr_param(a,b))

randi(n_rows, n_cols, n_slices)
randi(n_rows, n_cols, n_slices, distr_param(a,b))

randi(size(X))
randi(size(X), distr_param(a,b))

- Generate a scalar, vector, matrix or cube with the elements set to random **integer** values uniformly distributed in the [a,b] interval
- The default distribution parameters are $a = 0$ and $b = \text{maximum_int}$
- Usage:
 - *scalar_type* v = randi<scalar_type>()
 - *scalar_type* v = randi<scalar_type>(distr_param(a,b))
 - *vector_type* v = randi<vector_type>(n_elem)
 - *vector_type* v = randi<vector_type>(n_elem, distr_param(a,b))
 - *matrix_type* X = randi<matrix_type>(n_rows, n_cols)
 - *matrix_type* X = randi<matrix_type>(n_rows, n_cols, distr_param(a,b))
 - *cube_type* Q = randi<cube_type>(n_rows, n_cols, n_slices)
 - *cube_type* Q = randi<cube_type>(n_rows, n_cols, n_slices, distr_param(a,b))
- **Caveat:** to generate a matrix with random floating point values (ie. *float* or *double*) instead of integers, use **randu()** instead

- Examples:

```
int a = randi();  
int b = randi(distr_param(-10, +20));  
  
imat A1 = randi(5, 6);  
imat A2 = randi(5, 6, distr_param(-10, +20));  
  
mat B1 = randi<mat>(5, 6);  
mat B2 = randi<mat>(5, 6, distr_param(-10, +20));
```

- See also:

- `randu()`
- `randperm()`
- `.imbue()`
- `ones()`
- `zeros()`
- `shuffle()`
- `size()`
- RNG seed setting

speye(n_rows, n_cols)
speye(size(X))

- Generate a sparse matrix with the elements along the main diagonal set to one and off-diagonal elements set to zero
- An identity matrix is generated when $n_rows = n_cols$
- Usage:
 - `sparse_matrix_type X = speye<sparse_matrix_type>(n_rows, n_cols)`
 - `sparse_matrix_type Y = speye<sparse_matrix_type>(size(X))`

- Examples:

```
sp_mat A = speye<sp_mat>(5,5);
```

- See also:
 - `spones()`
 - `sprandu()` & `sprandn()`
 - `eye()`
 - `size()`

spones(A)

- Generate a sparse matrix with the same structure as sparse matrix *A*, but with the non-zero elements set to one

- Examples:

```
sp_mat A = sprandu<sp_mat>(100, 200, 0.1);
```

```
sp_mat B = spones(A);
```

- See also:

- `speye()`
- `sprandu()` & `sprandn()`
- `ones()`

sprandu(n_rows, n_cols, density)
sprandn(n_rows, n_cols, density)

sprandu(size(X), density)
sprandn(size(X), density)

- Generate a sparse matrix with the non-zero elements set to random values
- The *density* argument specifies the percentage of non-zero elements; it must be in the [0,1] interval
- *sprandu()* uses a uniform distribution in the [0,1] interval
- *sprandn()* uses a normal/Gaussian distribution with zero mean and unit variance
- Usage:
 - *sparse_matrix_type* X = sprandu<*sparse_matrix_type*>(n_rows, n_cols, density)
 - *sparse_matrix_type* Y = sprandu<*sparse_matrix_type*>(size(X), density)

- Examples:

```
sp_mat A = sprandu<sp_mat>(100, 200, 0.1);
```

- See also:
 - [speye\(\)](#)
 - [spones\(\)](#)
 - [randu\(\)](#)
 - [randn\(\)](#)
 - [size\(\)](#)
 - [RNG seed setting](#)
 - [uniform distribution in Wikipedia](#)
 - [normal distribution in Wikipedia](#)

toeplitz(A)

toeplitz(A, B)

circ_toeplitz(A)

- `toeplitz()`: generate a Toeplitz matrix, with the first column specified by *A*, and (optionally) the first row specified by *B*
- `circ_toeplitz()`: generate a circulant Toeplitz matrix
- *A* and *B* must be vectors

- Examples:

```
vec A(5, fill::randu);  
mat X = toeplitz(A);  
mat Y = circ_toeplitz(A);
```

- See also:
 - [Toeplitz matrix in MathWorld](#)
 - [Toeplitz matrix in Wikipedia](#)
 - [Circulant matrix in Wikipedia](#)

Functions of Vectors / Matrices / Cubes

abs(X)

- Obtain the magnitude of each element
- Usage for non-complex X:
 - `Y = abs(X)`
 - X and Y must have the same matrix type or cube type, such as *mat* or *cube*
- Usage for complex X:
 - *real_object_type* `Y = abs(X)`
 - The type of X must be a complex matrix or complex cube, such as *cx_mat* or *cx_cube*
 - The type of Y must be the real counterpart to the type of X; if X has the type *cx_mat*, then the type of Y must be *mat*

- Examples:

```
mat A(5, 5, fill::randu);  
mat B = abs(A);
```

```
cx_mat X(5, 5, fill::randu);  
mat Y = abs(X);
```

- See also:
 - `arg()`
 - `conj()`
 - `imag() / real()`
 - `pow()`
 - miscellaneous element-wise functions

accu(X)

- Accumulate (sum) all elements of a vector, matrix or cube

- Examples:

```
mat A(5, 6, fill::randu);  
mat B(5, 6, fill::randu);
```

```
double x = accu(A);
```

```
double y = accu(A % B);
```

```
// accu(A % B) is a "multiply-and-accumulate" operation  
// as operator % performs element-wise multiplication
```

- See also:

- `sum()`
- `cumsum()`
- `trace()`
- `mean()`
- `dot()`
- `as_scalar()`

affmul(A, B)

- Multiply matrix A by an automatically extended form of B
- A is typically an affine transformation matrix
- B can be a vector or matrix, and is treated as having an additional row of ones
- The number of columns in A must be equal to number of rows in the extended form of B (ie. $A.n_cols = B.n_rows+1$)
- If A is 3×3 and B is 2×1 , the equivalent matrix multiplication is:

$$\begin{bmatrix} C_0 \\ C_1 \\ C_2 \end{bmatrix} = \begin{bmatrix} A_{00} & A_{01} & A_{02} \\ A_{10} & A_{11} & A_{12} \\ A_{20} & A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_0 \\ B_1 \\ 1 \end{bmatrix}$$

- If A is 2×3 and B is 2×1 , the equivalent matrix multiplication is:

$$\begin{bmatrix} C_0 \\ C_1 \end{bmatrix} = \begin{bmatrix} A_{00} & A_{01} & A_{02} \\ A_{10} & A_{11} & A_{12} \end{bmatrix} \times \begin{bmatrix} B_0 \\ B_1 \\ 1 \end{bmatrix}$$

- Examples:

```
mat A(2, 3, fill::randu);  
vec B(2, fill::randu);  
  
vec C = affmul(A,B);
```

- See also:
 - [math operators](#)
 - [Affine transformation in Wikipedia](#)
 - [Transformation matrix in Wikipedia](#)

all(V)
all(X)
all(X, dim)

- For vector *V*, return *true* if all elements of the vector are non-zero or satisfy a relational condition
- For matrix *X* and
 - *dim* = 0, return a row vector (of type *urowvec* or *umat*), with each element (0 or 1) indicating whether the corresponding column of *X* has all non-zero elements
 - *dim* = 1, return a column vector (of type *ucolvec* or *umat*), with each element (0 or 1) indicating whether the corresponding row of *X* has all non-zero elements
- The *dim* argument is optional; by default *dim* = 0 is used
- Relational operators can be used instead of *V* or *X*, eg. *A* > 0.5
- Examples:

```
vec V(10, fill::randu);  
mat X(5, 5, fill::randu);
```

```
// status1 will be set to true if vector V has all non-zero elements  
bool status1 = all(V);
```

```
// status2 will be set to true if vector V has all elements greater than 0.5  
bool status2 = all(V > 0.5);
```

```
// status3 will be set to true if matrix X has all elements greater than 0.6;  
// note the use of vectorise()  
bool status3 = all(vectorise(X) > 0.6);
```

```
// generate a row vector indicating which columns of X have all elements greater than 0.7  
umat A = all(X > 0.7);
```

- See also:
 - *any()*
 - *approx_equal()*
 - *find()*

- `.is_zero()`
- `conv_to()` (convert between matrix/vector types)
- `vectorise()`

any(V)
any(X)
any(X, dim)

- For vector *V*, return *true* if any element of the vector is non-zero or satisfies a relational condition
- For matrix *X* and
 - *dim* = 0, return a row vector (of type *urowvec* or *umat*), with each element (0 or 1) indicating whether the corresponding column of *X* has any non-zero elements
 - *dim* = 1, return a column vector (of type *ucolvec* or *umat*), with each element (0 or 1) indicating whether the corresponding row of *X* has any non-zero elements
- The *dim* argument is optional; by default *dim* = 0 is used
- Relational operators can be used instead of *V* or *X*, eg. *A* > 0.9
- Examples:

```
vec V(10, fill::randu);  
mat X(5, 5, fill::randu);
```

```
// status1 will be set to true if vector V has any non-zero elements  
bool status1 = any(V);
```

```
// status2 will be set to true if vector V has any elements greater than 0.5  
bool status2 = any(V > 0.5);
```

```
// status3 will be set to true if matrix X has any elements greater than 0.6;  
// note the use of vectorise()  
bool status3 = any(vectorise(X) > 0.6);
```

```
// generate a row vector indicating which columns of X have elements greater than 0.7  
umat A = any(X > 0.7);
```

- See also:
 - *all()*
 - *approx_equal()*
 - *find()*

- `conv_to()` (convert between matrix/vector types)
- `vectorise()`

approx_equal(A, B, method, tol)
approx_equal(A, B, method, abs_tol, rel_tol)

- Return *true* if all corresponding elements in *A* and *B* are approximately equal
- Return *false* if any of the corresponding elements in *A* and *B* are not approximately equal, or if *A* and *B* have different dimensions
- The argument *method* controls how the approximate equality is determined; it is one of:
 - "absdiff" \rightarrow scalars *x* and *y* are considered equal if $|x - y| \leq tol$
 - "reldiff" \rightarrow scalars *x* and *y* are considered equal if $|x - y| / \max(|x|, |y|) \leq tol$
 - "both" \rightarrow scalars *x* and *y* are considered equal if $|x - y| \leq abs_tol$ **or** $|x - y| / \max(|x|, |y|) \leq rel_tol$

- Examples:

```
mat A(5, 5, fill::randu);  
mat B = A + 0.001;  
  
bool same1 = approx_equal(A, B, "absdiff", 0.002);  
  
mat C = 1000 * randu<mat>(5,5);  
mat D = C + 1;  
  
bool same2 = approx_equal(C, D, "reldiff", 0.1);  
  
bool same3 = approx_equal(C, D, "both", 2, 0.1);
```

- See also:
 - `all()`
 - `any()`
 - `find()`
 - `.is_zero()`
 - relational operators

arg(X)

- Obtain the phase angle (in radians) of each element
- Usage for non-complex X:
 - $Y = \arg(X)$
 - X and Y must have the same matrix type or cube type, such as *mat* or *cube*
 - non-complex elements are treated as complex elements with zero imaginary component
- Usage for complex X:
 - *real_object_type* $Y = \arg(X)$
 - The type of X must be a complex matrix or complex cube, such as *cx_mat* or *cx_cube*
 - The type of Y must be the real counterpart to the type of X; if X has the type *cx_mat*, then the type of Y must be *mat*
- Examples:

```
cx_mat A(5, 5, fill::randu);  
mat B = arg(A);
```
- See also:
 - [abs\(\)](#)
 - [atan2\(\)](#)
 - [Argument \(complex analysis\) in Wikipedia](#)
 - [Complex Argument in MathWorld](#)

as_scalar(expression)

- Evaluate an expression that results in a 1x1 matrix, followed by converting the 1x1 matrix to a pure scalar
- Optimised expression evaluations are automatically used when a binary or trinary expression is given (ie. 2 or 3 terms)
- Examples:

```
rowvec r(5, fill::randu);
colvec q(5, fill::randu);

mat X(5, 5, fill::randu);

// examples of expressions which have optimised implementations

double a = as_scalar(r*q);
double b = as_scalar(r*X*q);
double c = as_scalar(r*diagmat(X)*q);
double d = as_scalar(r*inv(diagmat(X))*q);
```

- See also:
 - `.as_col()` / `.as_row()`
 - `vectorise()`
 - `accu()`
 - `trace()`
 - `dot()`
 - `norm()`
 - `conv_to()`

clamp(X, min_val, max_val)

- Create a copy of X with each element clamped to the $[min_val, max_val]$ interval; any value lower than min_val will be set to min_val , and any value higher than max_val will be set to max_val
- For objects with complex elements, the real and imaginary components are clamped separately
- If X is a sparse matrix, clamping is applied only to the non-zero elements

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = clamp(A, 0.2, 0.8);  
  
mat C = clamp(A, A.min(), 0.8);  
  
mat D = clamp(A, 0.2, A.max());
```

- See also:
 - `.clamp()` (member function)
 - `.min()` & `.max()`
 - `.clean()`
 - `.replace()`
 - `find()`

cond(A)

- Return the condition number of matrix A (the ratio of the largest singular value to the smallest)
- Matrix A is well-conditioned when the condition number is close to 1
- Matrix A is ill-conditioned (poorly conditioned) when the condition number is large
- The computation is based on singular value decomposition
- **Caveat:** calculating the approximate reciprocal condition number via `rcond()` is considerably more efficient
- Examples:

```
mat A(5, 5, fill::randu);
```

```
double c = cond(A);
```

- See also:
 - `rcond()`
 - `rank()`
 - `inv()`
 - `svd()`
 - [condition number in MathWorld](#)
 - [condition number in Wikipedia](#)

conj(X)

- Obtain the complex conjugate of each element in a complex matrix or cube

- Examples:

```
cx_mat X(5, 5, fill::randu);  
cx_mat Y = conj(X);
```

- See also:

- `abs()`
- `imag() / real()`
- `trans()`

conv_to< type >::from(X)

- Convert (cast) from one matrix type to another (eg. *mat* to *imat*), or one cube type to another (eg. *cube* to *icube*)
- Conversion between *std::vector* and Armadillo matrices/vectors is also possible
- Conversion of a *mat* object into *colvec*, *rowvec* or *std::vector* is possible if the object can be interpreted as a vector

- Examples:

```
mat A(5, 5, fill::randu);  
fmat B = conv_to<fmat>::from(A);  
  
typedef std::vector<double> stdvec;  
  
stdvec x(3);  
x[0] = 0.0; x[1] = 1.0; x[2] = 2.0;  
  
colvec y = conv_to< colvec >::from(x);  
stdvec z = conv_to< stdvec >::from(y);
```

- See also:
 - [as_scalar\(\)](#)
 - [abs\(\)](#)
 - [imag\(\) / real\(\)](#)
 - [advanced constructors \(matrices\)](#)
 - [advanced constructors \(cubes\)](#)

cross(A, B)

- Calculate the cross product between A and B, under the assumption that A and B are 3 dimensional vectors

- Examples:

```
vec a(3, fill::randu);  
vec b(3, fill::randu);  
  
vec c = cross(a,b);
```

- See also:

- [dot\(\)](#)
- [Cross product in Wikipedia](#)
- [Cross product in MathWorld](#)

cumsum(V)
cumsum(X)
cumsum(X, dim)

- For vector V , return a vector of the same orientation, containing the cumulative sum of elements
- For matrix X , return a matrix containing the cumulative sum of elements in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used
- Examples:

```
mat A(5, 5, fill::randu);  
mat B = cumsum(A);  
mat C = cumsum(A, 1);  
  
vec x(10, fill::randu);  
vec y = cumsum(x);
```

- See also:
 - `cumprod()`
 - `accu()`
 - `sum()`
 - `diff()`

cumprod(V)
cumprod(X)
cumprod(X, dim)

- For vector V , return a vector of the same orientation, containing the cumulative product of elements
- For matrix X , return a matrix containing the cumulative product of elements in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used
- Examples:

```
mat A(5, 5, fill::randu);  
mat B = cumprod(A);  
mat C = cumprod(A, 1);  
  
vec x(10, fill::randu);  
vec y = cumprod(x);
```

- See also:
 - `cumsum()`
 - `prod()`

val = det(A) (form 1)

det(val, A) (form 2)

- Calculate the determinant of square matrix A , based on LU decomposition
- form 1: return the determinant
- form 2: store the calculated determinant in val and return a bool indicating success
- If A is not square sized, a *std::logic_error* exception is thrown
- If the calculation fails:
 - $val = \det(A)$ throws a *std::runtime_error* exception
 - $\det(val, A)$ returns a bool set to *false* (exception is not thrown)
- **Caveat:** `log_det()` is preferred, as it's less likely to suffer from numerical underflows/overflows

- Examples:

```
mat A(5, 5, fill::randu);  
  
double val1 = det(A);           // form 1  
  
double val2;  
bool success = det(val2, A);    // form 2
```

- See also:
 - `log_det()`
 - `rcond()`
 - [determinant in MathWorld](#)
 - [determinant in Wikipedia](#)

diagmat(V)
diagmat(V, k)

diagmat(X)
diagmat(X, k)

- Generate a diagonal matrix from vector V or matrix X
- Given vector V :
 - generate a square matrix with the k -th diagonal containing a copy of the vector; all other elements are set to zero
- Given matrix X :
 - generate a matrix with the k -th diagonal containing a copy of the k -th diagonal of X ; all other elements are set to zero
 - if X is an expression, the evaluation of the expression aims to calculate only the diagonal elements
- The argument k is optional; by default $k = 0$
- The argument k specifies the diagonal to use:
 - $k = 0$ indicates the main diagonal (default setting)
 - $k < 0$ indicates the k -th sub-diagonal (below main diagonal, towards bottom-left corner)
 - $k > 0$ indicates the k -th super-diagonal (above main diagonal, towards top-right corner)

- Examples:

```
mat A(5, 5, fill::randu);  
mat B = diagmat(A);  
mat C = diagmat(A,1);  
  
vec v(5, fill::randu);  
mat D = diagmat(v);  
mat E = diagmat(v,1);
```

- See also:
 - `.diag()`
 - `.is_diagmat()`

- `diagvec()`
- `diags()` / `spdiags()`
- `trace()`
- `trimatu()` / `trimatl()`
- `symmatu()` / `symmatl()`
- `reshape()`
- `vectorise()`
- diagonal matrix in Wikipedia

diagvec(X)

diagvec(X, k)

- Extract the k -th diagonal from matrix X
- The argument k is optional; by default $k = 0$
- The argument k specifies the diagonal to extract:
 - $k = 0$ indicates the main diagonal (default setting)
 - $k < 0$ indicates the k -th sub-diagonal (below main diagonal, towards bottom-left corner)
 - $k > 0$ indicates the k -th super-diagonal (above main diagonal, towards top-right corner)
- The extracted diagonal is interpreted as a column vector
- If X is an expression, the evaluation of the expression aims to calculate only the diagonal elements

- Examples:

```
mat A(5, 5, fill::randu);
```

```
vec d = diagvec(A);
```

- See also:
 - `.diag()`
 - `diagmat()`
 - `trace()`
 - `vectorise()`

diags(V, D, n_rows, n_cols)
spdiags(V, D, n_rows, n_cols)

- Generate a matrix with diagonals specified by vector D copied from corresponding column vectors in matrix V
- Function *diags()* generates a dense matrix, while function *spdiags()* generates a sparse matrix; in both cases the generated matrix has the same element type as matrix V
- Given vector D must be a vector of type *ivec* or *irowvec*
- The number of columns in given matrix V must match the number of elements in vector D
- Each element in vector D specifies diagonal k , where:
 - $k = 0$ indicates the main diagonal
 - $k < 0$ indicates the k -th sub-diagonal (below main diagonal, towards bottom-left corner)
 - $k > 0$ indicates the k -th super-diagonal (above main diagonal, towards top-right corner)
- For $k < 0$, vectors from matrix V are truncated at the end
- For $k > 0$, vectors from matrix V are truncated from the start
- Examples:

```
mat V(10, 3, fill::randu);  
ivec D = {-1, 0, +1};  
mat X = diags(V, D, 10, 10);  
sp_mat Y = spdiags(V, D, 10, 10);
```
- See also:
 - *.diag()*
 - *diagmat()*
 - [band matrix in Wikipedia](#)

diff(V)
diff(V, k)

diff(X)
diff(X, k)
diff(X, k, dim)

- For vector V , return a vector of the same orientation, containing the differences between consecutive elements
- For matrix X , return a matrix containing the differences between consecutive elements in each column ($dim = 0$), or each row ($dim = 1$)
- The optional argument k indicates that the differences are calculated recursively k times; by default $k = 1$ is used
- The resulting number of differences is $n - k$, where n is the number of elements; if $n \leq k$, the number of differences is zero (ie. an empty vector/matrix is returned)
- The argument dim is optional; by default $dim = 0$

- Examples:

```
vec a = linspace<vec>(1,10,10);
```

```
vec b = diff(a);
```

- See also:
 - [trapz\(\)](#)
 - [numerical differentiation in Wikipedia](#)
 - [numerical differentiation in MathWorld](#)

dot(A, B)
cdot(A, B)
norm_dot(A, B)

- *dot(A,B)*: dot product of *A* and *B*, treating *A* and *B* as vectors
- *cdot(A,B)*: as per *dot(A,B)*, but the complex conjugate of *A* is used
- *norm_dot(A,B)*: normalised dot product; equivalent to $\text{dot}(A,B) / (\|A\| \cdot \|B\|)$
- **Caveat:** *norm()* is more robust for calculating the norm, as it handles underflows and overflows
- Examples:

```
vec a(10, fill::randu);  
vec b(10, fill::randu);  
  
double x = dot(a,b);
```

- See also:
 - *norm()*
 - *as_scalar()*
 - *cross()*
 - *conj()*

eps(X)

- Obtain the positive distance of the absolute value of each element of X to the next largest representable floating point number
- X can be a scalar (eg. *double*), vector or matrix
- Examples:

```
mat A(4, 5, fill::randu);  
mat B = eps(A);
```

- See also:
 - [datum::eps](#)
 - [Floating-Point Arithmetic in MathWorld](#)
 - [IEEE Standard for Floating-Point Arithmetic in Wikipedia](#)

B = expmat(A)
expmat(B, A)

- Matrix exponential of general square matrix A
- If A is not square sized, a *std::logic_error* exception is thrown
- If the matrix exponential cannot be found:
 - $B = \text{expmat}(A)$ resets B and throws a *std::runtime_error* exception
 - $\text{expmat}(B,A)$ resets B and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - the matrix exponential operation is generally **not** the same as applying the `exp()` function to each element
 - if matrix A is symmetric, using `expmat_sym()` is faster

- Examples:

```
mat A(5, 5, fill::randu);
```

```
mat B = expmat(A);
```

- See also:
 - `expmat_sym()`
 - `logmat()`
 - `sqrtmat()`
 - miscellaneous element-wise functions
 - matrix exponential in Wikipedia
 - matrix exponential in MathWorld
 - [Nineteen Dubious Ways to Compute the Exponential of a Matrix, Twenty-Five Years Later](#)

B = expmat_sym(A)
expmat_sym(B, A)

- Matrix exponential of symmetric/hermitian matrix A
- The computation is based on eigen decomposition
- If A is not square sized, a *std::logic_error* exception is thrown
- If the matrix exponential cannot be found:
 - $B = \text{expmat_sym}(A)$ resets B and throws a *std::runtime_error* exception
 - $\text{expmat_sym}(B,A)$ resets B and returns a bool set to *false* (exception is not thrown)
- **Caveat:** the matrix exponential operation is generally **not** the same as applying the `exp()` function to each element

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = A*A.t();    // make symmetric matrix  
  
mat C = expmat_sym(B);
```

- See also:
 - `expmat()`
 - `logmat_sympd()`
 - `sqrtmat_sympd()`
 - `.is_symmetric()`
 - `.is_hermitian()`
 - miscellaneous element-wise functions
 - matrix exponential in Wikipedia
 - matrix exponential in MathWorld
 - [Nineteen Dubious Ways to Compute the Exponential of a Matrix, Twenty-Five Years Later](#)

find(X)
find(X, k)
find(X, k, s)

- Return a column vector containing the indices of elements of X that are non-zero or satisfy a relational condition
- The output vector must have the type `uvec` (ie. the indices are stored as unsigned integers of type `uword`)
- X is interpreted as a vector, with column-by-column ordering of the elements of X
- Relational operators can be used instead of X , eg. $A > 0.5$
- If $k = 0$ (default), return the indices of all non-zero elements, otherwise return at most k of their indices
- If $s = \text{"first"}$ (default), return at most the first k indices of the non-zero elements
- If $s = \text{"last"}$, return at most the last k indices of the non-zero elements
- **Caveats:**
 - to clamp values to an interval, `clamp()` is more efficient
 - to replace a specific value, `.replace()` is more efficient

- Examples:

```
mat A(5, 5, fill::randu);
mat B(5, 5, fill::randu);

uvec q1 = find(A > B);
uvec q2 = find(A > 0.5);
uvec q3 = find(A > 0.5, 3, "last");

// change elements of A greater than 0.5 to 1
A.elem( find(A > 0.5) ).ones();
```

- See also:
 - `all()`
 - `any()`

- `clamp()`
- `.transform()`
- `find_finite()`
- `find_nonfinite()`
- `find_nan()`
- `find_unique()`
- `nonzeros()`
- `unique()`
- `sort_index()`
- `ind2sub()`
- `.index_min()` & `.index_max()`
- `trimatu_ind()` / `trimatl_ind()`
- submatrix views
- subcube views

find_finite(X)

- Return a column vector containing the indices of elements of X that are finite (ie. not $\pm\text{Inf}$ and not NaN)
- The output vector must have the type **uvec** (ie. the indices are stored as unsigned integers of type **uword**)
- X is interpreted as a vector, with column-by-column ordering of the elements of X
- Examples:

```
mat A(5, 5, fill::randu);  
  
A(1,1) = datum::inf;  
  
// accumulate only finite elements  
double val = accu( A.elem( find_finite(A) ) );
```

- See also:
 - **find()**
 - **find_nonfinite()**
 - **find_nan()**
 - **.is_finite()**
 - **.replace()**
 - **.has_inf()**
 - **.has_nan()**
 - **submatrix views**
 - **constants (pi, nan, inf, ...)**

find_nonfinite(X)

- Return a column vector containing the indices of elements of X that are non-finite (ie. $\pm\text{Inf}$ or NaN)
- The output vector must have the type **uvec** (ie. the indices are stored as unsigned integers of type **uword**)
- X is interpreted as a vector, with column-by-column ordering of the elements of X
- Examples:

```
mat A(5, 5, fill::randu);  
  
A(1,1) = datum::inf;  
A(2,2) = datum::nan;  
  
// change non-finite elements to zero  
A.elem( find_nonfinite(A) ).zeros();
```

- **Caveat:** to replace instances of a specific non-finite value (eg. NaN or Inf), it is more efficient to use **.replace()**
- See also:
 - **find()**
 - **find_finite()**
 - **find_nan()**
 - **.is_finite()**
 - **.replace()**
 - **.has_inf()**
 - **.has_nan()**
 - **submatrix views**
 - **constants (pi, nan, inf, ...)**

find_nan(X)

- Return a column vector containing the indices of elements of X that are NaN (not-a-number)
- The output vector must have the type **uvec** (ie. the indices are stored as unsigned integers of type **uword**)
- X is interpreted as a vector, with column-by-column ordering of the elements of X

- Examples:

```
mat A(5, 5, fill::randu);  
  
A(2,3) = datum::nan;  
  
uvec indices = find_nan(A);
```

- **Caveat:** to replace instances of NaN values, it is more efficient to use **.replace()**
- See also:
 - **find()**
 - **find_finite()**
 - **find_nonfinite()**
 - **.replace()**
 - **.has_nan()**
 - **constants** (pi, nan, inf, ...)
 - **NaN** in Wikipedia

find_unique(X)

find_unique(X, ascending_indices)

- Return a column vector containing the indices of unique elements of X
- The output vector must have the type `uvec` (ie. the indices are stored as unsigned integers of type `uword`)
- X is interpreted as a vector, with column-by-column ordering of the elements of X
- The *ascending_indices* argument is optional; it is one of:
 - `true` = the returned indices are sorted to be ascending (**default setting**)
 - `false` = the returned indices are in arbitrary order (faster operation)

- Examples:

```
mat A = { { 2, 2, 4 },  
          { 4, 6, 6 } };
```

```
uvec indices = find_unique(A);
```

- See also:
 - `find()`
 - `unique()`
 - `intersect()`
 - `submatrix views`

fliplr(X)
flipud(X)

- *fliplr()*: generate a copy of matrix *X*, with the order of the columns reversed
- *flipud()*: generate a copy of matrix *X*, with the order of the rows reversed
- Examples:

```
mat A(5, 5, fill::randu);
```

```
mat B = fliplr(A);
```

```
mat C = flipud(A);
```

- See also:
 - *reverse()*
 - *shift()*
 - *.swap_rows()* & *.swap_cols()*
 - *.t()*

imag(X)
real(X)

- Extract the imaginary/real part of a complex matrix or cube

- Examples:

```
cx_mat C(5, 5, fill::randu);
```

```
mat    A = imag(C);  
mat    B = real(C);
```

- See also:

- `.set_imag()` / `.set_real()`
- `abs()`
- `conj()`
- `conv_to()`

`uvec sub = ind2sub(size(X), index)` (form 1)

`umat sub = ind2sub(size(X), vector_of_indices)` (form 2)

- Convert a linear index, or a vector of indices, to subscript notation
- The argument **size(X)** can be replaced with **size(n_rows, n_cols)** or **size(n_rows, n_cols, n_slices)**
- If an index is out of range, a *std::logic_error* exception is thrown
- When only one index is given (form 1), the subscripts are returned in a vector of type **uvec**
- When a vector of indices (of type **uvec**) is given (form 2), the corresponding subscripts are returned in each column of an $m \times n$ matrix of type **umat**; $m=2$ for matrix subscripts, while $m=3$ for cube subscripts

- Examples:

```
mat M(4, 5, fill::randu);

uvec s = ind2sub( size(M), 6 );

cout << "row: " << s(0) << endl;
cout << "col: " << s(1) << endl;

uvec indices = find(M > 0.5);
umat t       = ind2sub( size(M), indices );

cube Q(2,3,4);

uvec u = ind2sub( size(Q), 8 );

cout << "row:   " << u(0) << endl;
cout << "col:   " << u(1) << endl;
cout << "slice: " << u(2) << endl;
```

- See also:
 - **size()**
 - **sub2ind()**
 - **element access**

- `find()`

index_min(V)	index_max(V)
index_min(M)	index_max(M)
index_min(M, dim)	index_max(M, dim)
index_min(Q)	index_max(Q)
index_min(Q, dim)	index_max(Q, dim)

- For vector V , return the linear index of the extremum value; the returned index is of type **uword**
- For matrix M and:
 - $dim = 0$, return a row vector (of type **urowvec** or **umat**), with each column containing the index of the extremum value in the corresponding column of M
 - $dim = 1$, return a column vector (of type **uvec** or **umat**), with each row containing the index of the extremum value in the corresponding row of M
- For cube Q , return a cube (of type **ucube**) containing the indices of extremum values of elements along dimension dim , where $dim \in \{ 0, 1, 2 \}$
- For each column, row, or slice, the index starts at zero
- The dim argument is optional; by default $dim = 0$ is used
- For objects with complex numbers, absolute values are used for comparison
- Examples:

```
vec v(10, fill::randu);
```

```
uword i = index_max(v);
double max_val_in_v = v(i);
```

```
mat M(5, 6, fill::randu);
```

```
urowvec ii = index_max(M);
ucolvec jj = index_max(M,1);
```

```
double max_val_in_col_2 = M( ii(2), 2 );
```

```
double max_val_in_row_4 = M( 4, jj(4) );
```

- See also:
 - `min()` & `max()`
 - `.index_min()` & `.index_max()` (member functions)
 - `sort_index()`
 - `find()`

inplace_trans(X)
inplace_trans(X, method)

inplace_strans(X)
inplace_strans(X, method)

- In-place / in-situ transpose of matrix *X*
- For real (non-complex) matrix:
 - *inplace_trans()* performs a normal transpose
 - *inplace_strans()* not applicable
- For complex matrix:
 - *inplace_trans()* performs a Hermitian transpose (ie. the conjugate of the elements is taken during the transpose)
 - *inplace_strans()* provides a transposed copy without taking the conjugate of the elements
- The argument *method* is optional
- By default, a greedy transposition algorithm is used; a low-memory algorithm can be used instead by explicitly setting *method* to "lowmem"
- The low-memory algorithm is considerably slower than the greedy algorithm; using the low-memory algorithm is only recommended for cases where *X* takes up more than half of available memory (ie. very large *X*)

- Examples:

```
mat X(4,      5,      fill::randu);
mat Y(20000, 30000, fill::randu);

inplace_trans(X);           // use greedy algorithm by default

inplace_trans(Y, "lowmem"); // use low-memory (and slow) algorithm
```

- See also:
 - `.t()`

- `trans()`
- inplace matrix transpose in Wikipedia

C = intersect(A, B) (form 1)

intersect(C, iA, iB, A, B) (form 2)

- For form 1:
 - return the unique elements common to both *A* and *B*, sorted in ascending order
- For form 2:
 - store in *C* the unique elements common to both *A* and *B*, sorted in ascending order
 - store in *iA* and *iB* the indices of the unique elements, such that $C = A.\text{elem}(iA)$ and $C = B.\text{elem}(iB)$
 - *iA* and *iB* must have the type **uvec** (ie. the indices are stored as unsigned integers of type **uword**)
- *C* is a column vector if either *A* or *B* is a matrix or column vector; *C* is a row vector if both *A* and *B* are row vectors
- For matrices and vectors with complex numbers, ordering is via absolute values

- Examples:

```
ivec A = regspace<ivec>(4, 1); // 4, 3, 2, 1
ivec B = regspace<ivec>(3, 6); // 3, 4, 5, 6
```

```
ivec C = intersect(A,B);      // 3, 4
```

```
ivec CC;
uvec iA;
uvec iB;
```

```
intersect(CC, iA, iB, A, B);
```

- See also:
 - **unique()**
 - **find_unique()**
 - **submatrix views**

join_rows(A, B) **join_horiz(A, B)**
join_rows(A, B, C) **join_horiz(A, B, C)**
join_rows(A, B, C, D) **join_horiz(A, B, C, D)**

join_cols(A, B) **join_vert(A, B)**
join_cols(A, B, C) **join_vert(A, B, C)**
join_cols(A, B, C, D) **join_vert(A, B, C, D)**

- *join_rows()* and *join_horiz()*: horizontal concatenation; join the corresponding rows of the given matrices; the given matrices must have the same number of rows
- *join_cols()* and *join_vert()*: vertical concatenation; join the corresponding columns of the given matrices; the given matrices must have the same number of columns

- Examples:

```
mat A(4, 5, fill::randu);  
mat B(4, 6, fill::randu);  
mat C(6, 5, fill::randu);
```

```
mat AB = join_rows(A,B);  
mat AC = join_cols(A,C);
```

- See also:
 - [.shed_rows / cols / slices](#)
 - [.insert_rows / cols / slices](#)
 - [join_slices\(\)](#)
 - [submatrix views](#)

join_slices(cube C, cube D)
join_slices(mat M, mat N)

join_slices(mat M, cube C)
join_slices(cube C, mat M)

- for two cubes C and D : join the slices of C with the slices of D ; cubes C and D must have the same number of rows and columns (ie. all slices must have the same size)
- for two matrices M and N : treat M and N as cube slices and join them to form a cube with 2 slices; matrices M and N must have the same number of rows and columns
- for matrix M and cube C : treat M as a cube slice and join it with the slices of C ; matrix M and cube C must have the same number of rows and columns

- Examples:

```
cube C(5, 10, 3, fill::randu);  
cube D(5, 10, 4, fill::randu);
```

```
cube E = join_slices(C,D);
```

```
mat M(10, 20, fill::randu);  
mat N(10, 20, fill::randu);
```

```
cube Q = join_slices(M,N);
```

```
cube R = join_slices(Q,M);
```

```
cube S = join_slices(M,Q);
```

- See also:
 - [.shed_rows / cols / slices](#)
 - [.insert_rows / cols / slices](#)
 - [join_rows / cols](#)
 - [subcube views](#)

kron(A, B)

- Kronecker tensor product
- Given matrix A (with n rows and p columns) and matrix B (with m rows and q columns), generate a matrix (with nm rows and pq columns) that denotes the tensor product of A and B

- Examples:

```
mat A(4, 5, fill::randu);  
mat B(5, 4, fill::randu);  
  
mat K = kron(A,B);
```

- See also:
 - [repmat\(\)](#)
 - [repelem\(\)](#)
 - [Kronecker product in MathWorld](#)
 - [Kronecker product in Wikipedia](#)

complex result = log_det(A) (form 1)

log_det(val, sign, A) (form 2)

- Log determinant of square matrix A , based on LU decomposition
- form 1: return the complex log determinant
 - if matrix A is real and the determinant is positive:
 - the real part is the log determinant
 - the imaginary part is zero
 - if matrix A is real and the determinant is negative:
 - the real part is $\log \text{abs}(\text{determinant})$
 - the imaginary part is equal to *datum::pi*
- form 2: store the calculated log determinant in val and $sign$, and return a bool indicating success; the determinant is equal to $\exp(val) \cdot sign$
- If A is not square sized, a *std::logic_error* exception is thrown
- If the log determinant cannot be found:
 - *result = log_det(A)* throws a *std::runtime_error* exception
 - *log_det(val, sign, A)* returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 5, fill::randu);

cx_double result = log_det(A);    // form 1

double val;
double sign;

bool ok = log_det(val, sign, A);  // form 2
```

- See also:
 - *log_det_sympd()*
 - *det()*
 - *rcond()*

- `cx_double`
- determinant in MathWorld
- determinant in Wikipedia

result = log_det_sympd(A) (form 1)

log_det_sympd(result, A) (form 2)

- Log determinant of symmetric positive definite matrix A
- form 1: return the log determinant
- form 2: store the calculated log determinant in *result* and return a bool indicating success
- If A is not square sized, a *std::logic_error* exception is thrown
- If the log determinant cannot be found:
 - *result = log_det_sympd(A)* throws a *std::runtime_error* exception
 - *log_det_sympd(result, A)* returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = A.t() * A; // make symmetric matrix  
  
double result1 = log_det_sympd(B);           // form 1  
  
double result2;  
bool success = log_det_sympd(result2, B);    // form 2
```

- See also:
 - [log_det\(\)](#)
 - [rcond\(\)](#)
 - [determinant in MathWorld](#)
 - [determinant in Wikipedia](#)

B = logmat(A)
logmat(B, A)

- Complex matrix logarithm of general square matrix A
- If A is not square sized, a *std::logic_error* exception is thrown
- If the matrix logarithm cannot be found:
 - $B = \text{logmat}(A)$ resets B and throws a *std::runtime_error* exception
 - $\text{logmat}(B,A)$ resets B and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - the matrix logarithm operation is generally **not** the same as applying the `log()` function to each element
 - if matrix A is symmetric positive definite, using `logmat_sympd()` is faster

- Examples:

```
mat A(5, 5, fill::randu);  
cx_mat B = logmat(A);
```

- See also:
 - `logmat_sympd()`
 - `expmat()`
 - `sqrtmat()`
 - `real()`
 - miscellaneous element-wise functions
 - matrix logarithm in Wikipedia

B = logmat_sympd(A)
logmat_sympd(B, A)

- Matrix logarithm of symmetric/hermitian positive definite matrix A
- If A is not square sized, a `std::logic_error` exception is thrown
- If the matrix logarithm cannot be found:
 - `B = logmat_sympd(A)` resets B and throws a `std::runtime_error` exception
 - `logmat_sympd(B,A)` resets B and returns a bool set to `false` (exception is not thrown)
- **Caveat:** the matrix logarithm operation is generally **not** the same as applying the `log()` function to each element

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = A*A.t();    // make symmetric matrix  
  
mat C = logmat_sympd(B);
```

- See also:
 - `logmat()`
 - `expmat_sym()`
 - `sqrtmat_sympd()`
 - `.is_sympd()`
 - miscellaneous element-wise functions
 - matrix logarithm in Wikipedia
 - positive definite matrix in Wikipedia
 - positive definite matrix in MathWorld

min(V)	max(V)
min(M)	max(M)
min(M, dim)	max(M, dim)
min(Q)	max(Q)
min(Q, dim)	max(Q, dim)
min(A, B)	max(A, B)

- For vector V , return the extremum value
- For matrix M , return the extremum value for each column ($dim = 0$), or each row ($dim = 1$)
- For cube Q , return the extremum values of elements along dimension dim , where $dim \in \{ 0, 1, 2 \}$
- The dim argument is optional; by default $dim = 0$ is used
- For two matrices/cubes A and B , return a matrix/cube containing element-wise extremum values
- For objects with complex numbers, absolute values are used for comparison

- Examples:

```
colvec v(10, fill::randu);
double x = max(v);

mat M(10, 10, fill::randu);

rowvec a = max(M);
rowvec b = max(M,0);
colvec c = max(M,1);

// element-wise maximum
mat X(5, 6, fill::randu);
mat Y(5, 6, fill::randu);
mat Z = arma::max(X,Y); // use arma:: prefix to distinguish from std::max()
```

- See also:
 - `.min()` & `.max()` (member functions)
 - `index_min()` & `index_max()`
 - `clamp()`

- `statistics` functions
- `running_stat` - class for running statistics of scalars
- `running_stat_vec` - class for running statistics of vectors

nonzeros(X)

- Return a column vector containing the non-zero **values** of *X*
- *X* can be a sparse or dense matrix
- **Caveats:**
 - for dense matrices/vectors, to obtain the **number** of non-zero elements, the expression `accu(X != 0)` is more efficient
 - for sparse matrices, to obtain the **number** of non-zero elements, the `.n_nonzero` attribute is more efficient, eg. `X.n_nonzero`

- **Examples:**

```
sp_mat A = sprandu<sp_mat>(100, 100, 0.1);  
vec a = nonzeros(A);
```

```
mat B(100, 100, fill::eye);  
vec b = nonzeros(B);
```

- See also:
 - `find()`
 - `unique()`
 - `vectorise()`
 - `.clean()`
 - `.for_each()`

norm(X)

norm(X, p)

- Compute the p -norm of X , where X is a vector or matrix
- For vectors, p is an integer ≥ 1 , or one of: `"-inf"`, `"inf"`, `"fro"`
- For matrices, p is one of: 1, 2, `"inf"`, `"fro"`
- `"-inf"` is the minimum quasi-norm, `"inf"` is the maximum norm, `"fro"` is the Frobenius norm
- The argument p is optional; by default $p = 2$ is used
- For vector norm with $p = 2$ and matrix norm with $p = \text{"fro"}$, a robust algorithm is used to reduce the likelihood of underflows and overflows
- **Caveats:**
 - to obtain the zero/Hamming pseudo-norm (the number of non-zero elements), use this expression:
`accu(X != 0)`
 - to obtain the vector norm of each row or column of a matrix, use `vecnorm()`
 - matrix 2-norm (spectral norm) is based on SVD, which is computationally intensive for large matrices; a possible alternative is `norm2est()`
- Examples:

```
vec q(5, fill::randu);  
  
double x = norm(q, 2);  
double y = norm(q, "inf");
```
- See also:
 - `vecnorm()`
 - `norm2est()`
 - `normalise()`
 - `dot()`
 - [vector norm in Wikipedia](#)
 - [vector norm in MathWorld](#)

- [matrix norm in Wikipedia](#)
- [matrix norm in MathWorld](#)

norm2est(X)

norm2est(X, tol)

norm2est(X, tol, max_iter)

- Fast estimate of the 2-norm (spectral norm) of X , where X is a dense or sparse matrix
- The estimate is found via an iterative algorithm which finishes when one of the following conditions is met:
 - the relative difference between two consecutive estimates is less than the specified tolerance, ie. $|est_1 - est_2| / \max(est_1, est_2) < tol$
 - the number of iterations is equal to *max_iter*
- The optional argument *tol* specifies the tolerance for the relative difference; by default *tol* = $1e-6$ is used
- The optional argument *max_iter* specifies the maximum number of iterations; by default *max_iter* = 100 is used

- Examples:

```
mat X(2000, 3000, fill::randu);  
  
double x = norm2est(X);  
double y = norm2est(X, 1e-5);  
double z = norm2est(X, 1e-4, 10);
```

- See also:
 - [norm\(\)](#)
 - [matrix norm in Wikipedia](#)
 - [matrix norm in MathWorld](#)

normalise(V)
normalise(V, p)

normalise(X)
normalise(X, p)
normalise(X, p, dim)

- For vector V , return its normalised version (ie. having unit p -norm)
- For matrix X , return its normalised version, where each column ($dim = 0$) or row ($dim = 1$) has been normalised to have unit p -norm
- The p argument is optional; by default $p = 2$ is used
- The dim argument is optional; by default $dim = 0$ is used

- Examples:

```
vec A(10, fill::randu);  
vec B = normalise(A);  
vec C = normalise(A, 1);  
  
mat X(5, 6, fill::randu);  
mat Y = normalise(X);  
mat Z = normalise(X, 2, 1);
```

- See also:
 - [norm\(\)](#)
 - [vecnorm\(\)](#)
 - [norm_dot\(\)](#)
 - [Normalised vector in MathWorld](#)
 - [Unit vector in Wikipedia](#)

pow(A, scalar) (form 1)
pow(A, B) (form 2)
pow(M.each_col(), C) (form 3)
pow(M.each_row(), R) (form 4)
pow(Q.each_slice(), M) (form 5)

- Element-wise power operations
- form 1: raise all elements in A to the power denoted by the given scalar
- form 2: raise each element in A to the power denoted by the corresponding element in B ;
the sizes of A and B must be the same
- form 3: for each column vector of matrix M , raise each element to the power denoted by the corresponding element in column vector C ;
the number of rows in M and C must be the same
- form 4: for each row vector of matrix M , raise each element to the power denoted by the corresponding element in row vector R ;
the number of columns in M and R must be the same
- form 5: for each slice of cube Q , raise each element to the power denoted by the corresponding element in matrix M ;
the number of rows and columns in Q and M must be the same
- **Caveats:**
 - to raise all elements to the power 2, use `square()` instead
 - for the matrix power operation, which takes into account matrix structure, use `powmat()`

- **Examples:**

```
mat A(5, 6, fill::randu);  
mat B(5, 6, fill::randu);  
  
mat X = pow(A, 3.45);  
mat Y = pow(A, B);
```

```
vec C(5, fill::randu);  
rowvec R(6, fill::randu);  
  
mat Z1 = pow(A.each_col(), C);  
mat Z2 = pow(A.each_row(), R);
```

- See also:
 - `powmat()`
 - `abs()`
 - miscellaneous element-wise functions
 - `.each_col()` & `.each_row()`
 - `.each_slice()`

B = powmat(A, n)
powmat(B, A, n)

- Matrix power operation: raise square matrix *A* to the power of *n*, where *n* has the type *int* or *double*
- If *n* has the type *double*, the resultant matrix *B* always has complex elements
- For *n* = 0, an identity matrix is generated
- If *A* is not square sized, a *std::logic_error* exception is thrown
- If the matrix power cannot be found:
 - *B = powmat(A)* resets *B* and throws a *std::runtime_error* exception
 - *powmat(B,A)* resets *B* and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - the matrix power operation is generally **not** the same as applying the *pow()* function to each element
 - to find the inverse of a matrix, use *inv()* instead
 - to solve a system of linear equations, use *solve()* instead
 - to find the matrix square root, use *sqrtmat()* instead

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = powmat(A, 4);      //      integer exponent  
  
cx_mat C = powmat(A, 4.56); // non-integer exponent
```

- See also:
 - *pow()*
 - *sqrtmat()*
 - *inv()*
 - *eye()*
 - *operators*
 - *miscellaneous element-wise functions*

prod(V)
prod(M)
prod(M, dim)

- For vector V , return the product of all elements
- For matrix M , return the product of elements in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used

- Examples:

```
colvec v(10, fill::randu);  
double x = prod(v);  
  
mat M(10, 10, fill::randu);  
  
rowvec a = prod(M);  
rowvec b = prod(M,0);  
colvec c = prod(M,1);
```

- See also:
 - [cumprod\(\)](#)
 - [Schur product in operators](#)

r = rank(X) (form 1)

r = rank(X, tolerance)

rank(r, X) (form 2)

rank(r, X, tolerance)

- Calculate the rank of matrix X , based on singular value decomposition
- form 1: return the rank
- form 2: store the calculated rank in r and return a bool indicating success
- Any singular values less than *tolerance* are treated as zero
- The *tolerance* argument is optional; by default $tolerance = max_rc \cdot max_sv \cdot epsilon$, where:
 - $max_rc = \max(X.n_rows, X.n_cols)$
 - max_sv = maximum singular value of X (ie. spectral norm)
 - $epsilon$ = difference between 1 and the least value greater than 1 that is representable
- If the calculation fails:
 - $r = rank(X)$ throws a `std::runtime_error` exception
 - $rank(r,X)$ returns a bool set to *false* (exception is not thrown)
- **Caveat:** to distinguish from `std::rank`, use the `arma::` prefix, ie. `arma::rank(X)`

- Examples:

```
mat A(4, 5, fill::randu);

uword r1 = rank(A);           // form 1

uword r2;
bool success = rank(r2, A);   // form 2
```

- See also:
 - `rcond()`
 - `svd()`

- `orth()`
- `datum::eps`
- Rank in MathWorld
- Rank in Wikipedia

rcond(A)

- Return the 1-norm estimate of the reciprocal condition number of square matrix A
- Values close to 1 suggest that A is well-conditioned
- Values close to 0 suggest that A is badly conditioned
- If A is not square sized, a *std::logic_error* exception is thrown
- **Caveat:** to efficiently calculate rcond and matrix inverse at the same time, use `inv()`
- Examples:

```
mat A(5, 5, fill::randu);
```

```
double r = rcond(A);
```

- See also:
 - `cond()`
 - `rank()`
 - `det()`
 - `inv()`
 - `solve()`

repelem(A, num_copies_per_row, num_copies_per_col)

- Generate a matrix by replicating each element of matrix A
- The generated matrix has the following size:
 $n_rows = num_copies_per_row * A.n_rows$
 $n_cols = num_copies_per_col * A.n_cols$

- Examples:

```
mat A(2, 3, fill::randu);
```

```
mat B = repelem(A, 4, 5);
```

- See also:
 - `repmat()`
 - `kron()`
 - `reshape()`
 - `resize()`

repmat(A, num_copies_per_row, num_copies_per_col)

- Generate a matrix by replicating matrix *A* in a block-like fashion
- The generated matrix has the following size:
 $n_rows = num_copies_per_row \times A.n_rows$
 $n_cols = num_copies_per_col \times A.n_cols$
- **Caveat:** to apply a vector operation on each row or column of a matrix, it is generally more efficient to use `.each_row()` or `.each_col()`
- Examples:

```
mat A(2, 3, fill::randu);  
mat B = repmat(A, 4, 5);
```
- See also:
 - `.each_col()` & `.each_row()` (vector operations applied to each column or row)
 - `repelem()`
 - `kron()`
 - `reshape()`
 - `resize()`

reshape(X, n_rows, n_cols) (X is a vector or matrix)
reshape(X, size(Y))

reshape(Q, n_rows, n_cols, n_slices) (Q is a cube)
reshape(Q, size(R))

- Generate a vector/matrix/cube with given size specifications, whose elements are taken from the given object in a column-wise manner; the elements in the generated object are placed column-wise (ie. the first column is filled up before filling the second column)
- The layout of the elements in the generated object will be different to the layout in the given object
- If the total number of elements in the given object is less than the specified size, the remaining elements in the generated object are set to zero
- If the total number of elements in the given object is greater than the specified size, only a subset of elements is taken from the given object

- **Caveats:**

- to change the size without preserving data, use `.set_size()` instead, which is much faster
- to grow/shrink a matrix while preserving the elements **as well as** the layout of the elements, use `resize()` instead
- to flatten a matrix into a vector, use `vectorise()` or `.as_col()` / `.as_row()` instead

- **Examples:**

```
mat A(10, 5, fill::randu);
```

```
mat B = reshape(A, 5, 10);
```

- **See also:**

- `.reshape()` (member function)
- `.set_size()`
- `resize()`
- `vectorise()`
- `as_scalar()`

- `conv_to()`
- `diagmat()`
- `repmat()`
- `repelem()`
- `size()`
- `interp2()`

resize(X, n_rows, n_cols) (X is a vector or matrix)
resize(X, size(Y))

resize(Q, n_rows, n_cols, n_slices) (Q is a cube)
resize(Q, size(R))

- Generate a vector/matrix/cube with given size specifications, whose elements as well as the layout of the elements are taken from the given object
- **Caveat:** to change the size without preserving data, use `.set_size()` instead, which is much faster
- Examples:

```
mat A(4, 5, fill::randu);
```

```
mat B = resize(A, 7, 6);
```

- See also:
 - `.resize()` (member function of Mat and Cube)
 - `.set_size()` (member function of Mat and Cube)
 - `reshape()`
 - `vectorise()`
 - `as_scalar()`
 - `conv_to()`
 - `repmat()`
 - `repelem()`
 - `size()`

reverse(V)
reverse(X)
reverse(X, dim)

- For vector V , generate a copy of the vector with the order of elements reversed
- For matrix X , generate a copy of the matrix with the order of elements reversed in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used
- Examples:

```
vec v(123, fill::randu);  
vec y = reverse(v);  
  
mat A(4, 5, fill::randu);  
mat B = reverse(A);  
mat C = reverse(A,1);
```

- See also:
 - `fliplr()` & `flipud()`
 - `shift()`
 - `sort()`
 - `trans()`
 - `.t()`
 - `.swap_rows()` & `.swap_cols()`

R = roots(P)
roots(R, P)

- Find the complex roots of a polynomial function represented via vector P and store them in column vector R

- The polynomial function is modelled as:

$$y = p_0 x^N + p_1 x^{N-1} + p_2 x^{N-2} + \dots + p_{N-1} x^1 + p_N$$

where p_i is the i -th polynomial coefficient in vector P

- The computation is based on eigen decomposition; if the decomposition fails:
 - $R = \text{roots}(P)$ resets R and throws a `std::runtime_error` exception
 - $\text{roots}(R,P)$ resets R and returns a bool set to *false* (exception is not thrown)

- Examples:

```
vec P(5, fill::randu);
```

```
cx_vec R = roots(P);
```

- See also:
 - [polyval\(\)](#)
 - [polyfit\(\)](#)
 - [real\(\)](#)
 - [solve\(\)](#)
 - [zero of a function in Wikipedia](#)

shift(V, N)
shift(X, N)
shift(X, N, dim)

- For vector V , generate a copy of the vector with the elements shifted by N positions in a circular manner
- For matrix X , generate a copy of the matrix with the elements shifted by N positions in each column ($dim = 0$), or each row ($dim = 1$)
- N can be positive or negative
- The dim argument is optional; by default $dim = 0$ is used
- Examples:

```
mat A(4, 5, fill::randu);  
mat B = shift(A, -1);  
mat C = shift(A, +1);
```

- See also:
 - `shuffle()`
 - `fliplr()` & `flipud()`
 - `reverse()`

shuffle(V)
shuffle(X)
shuffle(X, dim)

- For vector V , generate a copy of the vector with the elements shuffled
- For matrix X , generate a copy of the matrix with the elements shuffled in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used
- Examples:

```
mat A(4, 5, fill::randu);  
mat B = shuffle(A);
```

- See also:
 - `shift()`
 - `sort()`
 - `unique()`
 - `randi()`
 - `randperm()`
 - RNG seed setting

size(X)
size(n_rows, n_cols)
size(n_rows, n_cols, n_slices)

- Obtain the dimensions of object *X*, or explicitly specify the dimensions
- Dimensions provided via *size(...)* can be used in conjunction with **matrix constructors**, **submatrix views**, **random matrix generation**, **ind2sub()**, **sub2ind()**, etc.
- *size(...)* objects support simple arithmetic operations such as addition and multiplication
- Two *size(...)* objects can be compared for equality/inequality
- *size(...)* objects can be passed to an `std::ostream` (eg. `cout`) via `<<`
- **Caveat:** to prevent interference from ***std::size()*** in C++17, preface Armadillo's *size(...)* with the *arma* namespace qualification, eg. *arma::size(...)*
- Examples:

```
mat A(5,6);

mat B(size(A), fill::zeros);

mat C; C.randu(size(A));

mat D = ones<mat>(size(A));

mat E(10,20, fill::ones);
E(3,4,size(C)) = C;    // access submatrix of E

mat F( size(A) + size(E) );

mat G( size(A) * 2 );

cout << "size of A: " << size(A) << endl;

bool is_same_size = (size(A) == size(E));
```

- See also:

- attributes

sort(V)
sort(V, sort_direction)

sort(X)
sort(X, sort_direction)
sort(X, sort_direction, dim)

- For vector V , return a vector which is a sorted version of the input vector
- For matrix X , return a matrix with the elements of the input matrix sorted in each column ($dim = 0$), or each row ($dim = 1$)
- The dim argument is optional; by default $dim = 0$ is used
- The $sort_direction$ argument is optional; $sort_direction$ is either "ascend" or "descend"; by default "ascend" is used
- For matrices and vectors with complex numbers, sorting is via absolute values
- Examples:

```
mat A(10, 10, fill::randu);  
mat B = sort(A);
```

- See also:
 - `sort_index()`
 - `.is_sorted()`
 - `shuffle()`
 - `unique()`
 - `reverse()`
 - `randi()`

sort_index(X)

sort_index(X, sort_direction)

stable_sort_index(X)

stable_sort_index(X, sort_direction)

- Return a vector which describes the sorted order of the elements of X (ie. it contains the indices of the elements of X)
- The output vector must have the type `uvec` (ie. the indices are stored as unsigned integers of type `uword`)
- X is interpreted as a vector, with column-by-column ordering of the elements of X
- The *sort_direction* argument is optional; *sort_direction* is either "ascend" or "descend"; by default "ascend" is used
- The *stable_sort_index()* variant preserves the relative order of elements with equivalent values
- For matrices and vectors with complex numbers, sorting is via absolute values
- Examples:

```
vec q(10, fill::randu);
```

```
uvec indices = sort_index(q);
```

- See also:
 - `sort()`
 - `find()`
 - `.is_sorted()`
 - `.index_min()` & `.index_max()`
 - `ind2sub()`

B = sqrtmat(A)
sqrtmat(B, A)

- Complex square root of general square matrix A
- If A is not square sized, a *std::logic_error* exception is thrown
- If matrix A appears to be singular, an approximate square root is attempted; additionally, *sqrtmat(B,A)* returns a bool set to false
- **Caveats:**
 - the square root of a matrix is generally **not** the same as applying the *sqrt()* function to each element
 - if matrix A is symmetric positive definite, using *sqrtmat_sympd()* is faster

- Examples:

```
mat A(5, 5, fill::randu);
```

```
cx_mat B = sqrtmat(A);
```

- See also:
 - *sqrtmat_sympd()*
 - *powmat()*
 - *expmat()*
 - *logmat()*
 - *chol()*
 - *real()*
 - miscellaneous element-wise functions
 - square root of a matrix in Wikipedia

B = sqrtmat_sympd(A)
sqrtmat_sympd(B, A)

- Square root of symmetric/hermitian positive definite matrix A
- If A is not square sized, a `std::logic_error` exception is thrown
- If the square root cannot be found:
 - `B = sqrtmat_sympd(A)` resets B and throws a `std::runtime_error` exception
 - `sqrtmat_sympd(B,A)` resets B and returns a bool set to `false` (exception is not thrown)
- **Caveat:** the matrix square root operation is generally **not** the same as applying the `sqrt()` function to each element

- Examples:

```
mat A(5, 5, fill::randu);  
  
mat B = A*A.t();    // make symmetric matrix  
  
mat C = sqrtmat_sympd(B);
```

- See also:
 - `sqrtmat()`
 - `expmat_sym()`
 - `logmat_sympd()`
 - `chol()`
 - `.is_sympd()`
 - miscellaneous element-wise functions
 - square root of a matrix in Wikipedia
 - positive definite matrix in Wikipedia
 - positive definite matrix in MathWorld

sum(V)
sum(M)
sum(M, dim)
sum(Q)
sum(Q, dim)

- For vector V , return the sum of all elements
- For matrix M , return the sum of elements in each column ($dim = 0$), or each row ($dim = 1$)
- For cube Q , return the sums of elements along dimension dim , where $dim \in \{ 0, 1, 2 \}$; for example, $dim = 0$ indicates the sum of elements in each column within each slice
- The dim argument is optional; by default $dim = 0$ is used
- **Caveat:** to get a sum of all the elements regardless of the object type (ie. vector, or matrix, or cube), use `accu()` instead

- Examples:

```
colvec v(10, fill::randu);  
double x = sum(v);  
  
mat M(10, 10, fill::randu);  
  
rowvec a = sum(M);  
rowvec b = sum(M,0);  
colvec c = sum(M,1);  
  
double y = accu(M);    // find the overall sum regardless of object type
```

- See also:

- `accu()`
- `cumsum()`
- `trace()`
- `trapz()`
- `mean()`
- `as_scalar()`

uword index = **sub2ind(size(M), row, col)** (*M* is a matrix)
uvec indices = **sub2ind(size(M), matrix_of_subscripts)**

uword index = **sub2ind(size(Q), row, col, slice)** (*Q* is a cube)
uvec indices = **sub2ind(size(Q), matrix_of_subscripts)**

- Convert subscripts to a linear index
- The argument **size(X)** can be replaced with **size(n_rows, n_cols)** or **size(n_rows, n_cols, n_slices)**
- For the *matrix_of_subscripts* argument, the subscripts must be stored in each column of an $m \times n$ matrix of type **umat**; $m = 2$ for matrix subscripts, while $m = 3$ for cube subscripts
- If a subscript is out of range, a *std::logic_error* exception is thrown
- Examples:

```
mat M(4,5);  
cube Q(4,5,6);  
  
uword i = sub2ind( size(M), 2, 3 );  
uword j = sub2ind( size(Q), 2, 3, 4 );
```

- See also:
 - **size()**
 - **ind2sub()**
 - **element access**

symmatu(A)
symmatu(A, do_conj)

symmatl(A)
symmatl(A, do_conj)

- *symmatu(A)*: generate symmetric matrix from square matrix *A*, by reflecting the upper triangle to the lower triangle
- *symmatl(A)*: generate symmetric matrix from square matrix *A*, by reflecting the lower triangle to the upper triangle
- If *A* is a complex matrix, the reflection uses the complex conjugate of the elements; to disable the complex conjugate, set *do_conj* to *false*
- If *A* is non-square, a *std::logic_error* exception is thrown

- Examples:

```
mat A(5, 5, fill::randu);
```

```
mat B = symmatu(A);  
mat C = symmatl(A);
```

- See also:
 - [diagmat\(\)](#)
 - [trimatu\(\)](#) / [trimatl\(\)](#)
 - [.is_symmetric\(\)](#)
 - [.is_hermitian\(\)](#)
 - [Symmetric matrix in Wikipedia](#)

trace(X)

- Sum of the elements on the main diagonal of matrix X
- If X is an expression, the evaluation of the expression aims to calculate only the diagonal elements

- Examples:

```
mat A(5, 5, fill::randu);
```

```
double x = trace(A);
```

- See also:

- `accu()`
- `as_scalar()`
- `.diag()`
- `diagvec()`
- `diagmat()`
- `sum()`

trans(A)

strans(A)

- For real (non-complex) matrix:
 - *trans()* provides a transposed copy of the matrix
 - *strans()* is not applicable
- For complex matrix:
 - *trans()* provides a Hermitian (conjugate) transposed copy (ie. signs of imaginary components are flipped)
 - *strans()* provides a simple transposed copy (ie. signs of imaginary components are not flipped)

- Examples:

```
mat A(5, 10, fill::randu);
```

```
mat B = trans(A);
```

```
mat C = A.t();    // equivalent to trans(A), but more compact
```

- See also:
 - [.t\(\)](#)
 - [inplace_trans\(\)](#)
 - [reverse\(\)](#)
 - [transpose in Wikipedia](#)
 - [transpose in MathWorld](#)
 - [conjugate transpose in Wikipedia](#)
 - [conjugate transpose in MathWorld](#)

trapz(X, Y)
trapz(X, Y, dim)

trapz(Y)
trapz(Y, dim)

- Compute the trapezoidal integral of Y with respect to spacing in X , in each column ($dim = 0$) or each row ($dim = 1$) of Y
- X must be a vector; its length must equal either the number of rows in Y (when $dim = 0$), or the number of columns in Y (when $dim = 1$)
- If X is not specified, unit spacing is used
- The dim argument is optional; by default $dim = 0$

- Examples:

```
vec X = linspace<vec>(0, datum::pi, 1000);  
vec Y = sin(X);  
  
mat Z = trapz(X,Y);
```

- See also:
 - [diff\(\)](#)
 - [sum\(\)](#)
 - [linspace\(\)](#)
 - [numerical integration in Wikipedia](#)
 - [numerical integration in MathWorld](#)
 - [trapezoidal rule in Wikipedia](#)

trimatu(A)
trimatu(A, k)

trimatl(A)
trimatl(A, k)

- Create a new matrix by copying either the upper or lower triangular part from square matrix *A*, and setting the remaining elements to zero
 - *trimatu()* copies the upper triangular part
 - *trimatl()* copies the lower triangular part
- The argument *k* specifies the diagonal which inclusively delineates the boundary of the triangular part
 - for $k > 0$, the *k*-th super-diagonal is used (above main diagonal, towards top-right corner)
 - for $k < 0$, the *k*-th sub-diagonal is used (below main diagonal, towards bottom-left corner)
- The argument *k* is optional; by default the main diagonal is used ($k = 0$)
- If *A* is non-square, a *std::logic_error* exception is thrown

- Examples:

```
mat A(5, 5, fill::randu);

mat U = trimatu(A);
mat L = trimatl(A);

mat UU = trimatu(A, 1); // omit the main diagonal
mat LL = trimatl(A, -1); // omit the main diagonal
```

- See also:
 - [.is_trimatu\(\) / .is_trimatl\(\)](#)
 - [trimatu_ind\(\) / trimatl_ind\(\)](#)
 - [symmatu\(\) / symmatl\(\)](#)
 - [diagmat\(\)](#)
 - [nonzeros\(\)](#)
 - [Triangular matrix in MathWorld](#)
 - [Triangular matrix in Wikipedia](#)

trimatu_ind(size(A))
trimatu_ind(size(A), k)

trimatl_ind(size(A))
trimatl_ind(size(A), k)

- Return a column vector containing the indices of elements that form the upper or lower triangle part of matrix *A*
 - *trimatu_ind()* refers to the upper triangular part
 - *trimatl_ind()* refers to the lower triangular part
- The output vector must have the type **uvec** (ie. the indices are stored as unsigned integers of type **uword**)
- The argument *k* specifies the diagonal which inclusively delineates the boundary of the triangular part
 - for $k > 0$, the *k*-th super-diagonal is used (above main diagonal, towards top-right corner)
 - for $k < 0$, the *k*-th sub-diagonal is used (below main diagonal, towards bottom-left corner)
- The argument *k* is optional; by default the main diagonal is used ($k = 0$)
- The argument **size(A)** can be replaced with **size(n_rows, n_cols)**

- Examples:

```
mat A(5, 5, fill::randu);

uvec upper_indices = trimatu_ind( size(A) );
uvec lower_indices = trimatl_ind( size(A) );

// extract upper/lower triangle into vector
vec upper_part = A(upper_indices);
vec lower_part = A(lower_indices);

// obtain indices without the main diagonal
uvec alt_upper_indices = trimatu_ind( size(A), 1);
uvec alt_lower_indices = trimatl_ind( size(A), -1);
```

- See also:
 - **trimatu()** / **trimatl()**

- `find()`
- `submatrix views`

unique(A)

- Return the unique elements of A , sorted in ascending order
- If A is a vector, the output is also a vector with the same orientation (row or column) as A ; if A is a matrix, the output is always a column vector
- Examples:

```
mat X = { { 1, 2 }  
          { 2, 3 } };
```

```
mat Y = unique(X);
```

- See also:
 - `find()`
 - `find_unique()`
 - `sort()`
 - `shuffle()`
 - `nonzeros()`
 - `intersect()`

vecnorm(X)

vecnorm(X, p)

vecnorm(X, p, dim)

- Compute the p -norm of each column vector ($dim = 0$) or row vector ($dim = 1$) of matrix X
- p is an integer ≥ 1 , or `"-inf"` (minimum quasi-norm), or `"inf"` (maximum norm)
- The arguments p and dim are optional; by default $p = 2$ and $dim = 0$ are used
- **Caveat:** to compute the matrix norm, use `norm()` instead
- Examples:

```
mat X(4, 5, fill::randu);
```

```
rowvec r = vecnorm(X, 2);
```

```
colvec c = vecnorm(X, "inf", 1);
```

- See also:
 - `norm()`
 - `normalise()`
 - [vector norm in Wikipedia](#)
 - [vector norm in MathWorld](#)

vectorise(X)
vectorise(X, dim)
vectorise(Q)

- Generate a flattened version of matrix X or cube Q
- The argument dim is optional; by default $dim = 0$ is used
- For $dim = 0$, the elements are copied from X column-wise, resulting in a column vector; equivalent to concatenating all the columns of X
- For $dim = 1$, the elements are copied from X row-wise, resulting in a row vector; equivalent to concatenating all the rows of X
- **Caveats:**
 - column-wise vectorisation is faster than row-wise vectorisation
 - for sparse matrices, row-wise vectorisation is not recommended
- Examples:

```
mat X(4, 5, fill::randu);  
  
vec v = vectorise(X);
```
- See also:
 - `.as_col()` / `.as_row()`
 - `.col_as_mat()` / `.row_as_mat()`
 - `nonzeros()`
 - `reshape()`
 - `resize()`
 - `diagvec()`
 - `as_scalar()`

miscellaneous element-wise functions:

exp	log	square	floor	erf	tgamma	sign
exp2	log2	sqrt	ceil	erfc	lgamma	
exp10	log10		round			
expm1	log1p		trunc			
trunc_exp	trunc_log					

- Apply a function to each element
- Usage:
 - $B = fn(A)$, where $fn(A)$ is one of the functions below
 - A and B must have the same matrix type or cube type, such as *mat* or *cube*

<code>exp(A)</code>	base-e exponential: e^x
<code>exp2(A)</code>	base-2 exponential: 2^x
<code>exp10(A)</code>	base-10 exponential: 10^x
<code>expm1(A)</code>	compute $\exp(A) - 1$ accurately for values of A close to zero (only for <i>float</i> and <i>double</i> elements)
<code>trunc_exp(A)</code>	base-e exponential, truncated to avoid infinity (only for <i>float</i> and <i>double</i> elements)
<code>log(A)</code>	natural log: $\log_e x$
<code>log2(A)</code>	base-2 log: $\log_2 x$
<code>log10(A)</code>	base-10 log: $\log_{10} x$
<code>log1p(A)</code>	compute $\log(1+A)$ accurately for values of A close to zero (only for <i>float</i> and <i>double</i> elements)
<code>trunc_log(A)</code>	natural log, truncated to avoid \pm infinity (only for <i>float</i> and <i>double</i> elements)
<code>square(A)</code>	square: x^2
<code>sqrt(A)</code>	square root: \sqrt{x}
<code>floor(A)</code>	largest integral value that is not greater than the input value
<code>ceil(A)</code>	smallest integral value that is not less than the input value
<code>round(A)</code>	round to nearest integer, with halfway cases rounded away from zero
<code>trunc(A)</code>	round to nearest integer, towards zero
<code>erf(A)</code>	error function (only for <i>float</i> and <i>double</i> elements)

`erfc(A)` complementary error function (only for *float* and *double* elements)
`gamma(A)` gamma function (only for *float* and *double* elements)
`lgamma(A)` natural log of the absolute value of gamma function (only for *float* and *double* elements)
`sign(A)` signum function; for each element a in A , the corresponding element b in B is:

$$b = \begin{cases} -1 & \text{if } a < 0 \\ 0 & \text{if } a = 0 \\ +1 & \text{if } a > 0 \end{cases}$$
 if a is complex and non-zero, then $b = a / \text{abs}(a)$

- **Caveats:**

- all of the above functions are applied element-wise, where each element is treated independently
- the element-wise functions `exp()`, `log()` and `sqrt()` have the corresponding functions `expmat()`, `logmat()` and `sqrtmat()` which take into account matrix structure

- **Examples:**

```
mat A(5, 5, fill::randu);
mat B = exp(A);
```

- **See also:**

- `abs()`
- `pow()`
- `clamp()`
- `conj()`
- `imag()` / `real()`
- `.transform()` (apply user-defined function to each element)
- `expmat()`
- `logmat()`
- `sqrtmat()`
- trigonometric functions
- statistics functions
- miscellaneous constants

trigonometric element-wise functions (cos, sin, tan, ...)

- For single argument functions, $B = \text{trig_fn}(A)$, where *trig_fn* is applied to each element in *A*, with *trig_fn* as one of:
 - *cos*, *acos*, *cosh*, *acosh*
 - *sin*, *asin*, *sinh*, *asinh*
 - *tan*, *atan*, *tanh*, *atanh*
 - *sinc*, defined as $\text{sinc}(x) = \sin(\pi x) / (\pi x)$ for $x \neq 0$, and $\text{sinc}(x) = 1$ for $x = 0$
- For dual argument functions, apply the function to each tuple of two corresponding elements in *X* and *Y*:
 - $Z = \text{atan2}(Y, X)$
 - $Z = \text{hypot}(X, Y)$

- Examples:

```
mat X(5, 5, fill::randu);  
mat Y = cos(X);
```

- See also:
 - [miscellaneous element-wise functions](#)
 - [trigonometric functions in Wikipedia](#)
 - [atan2 function in Wikipedia](#)
 - [hypot function in Wikipedia](#)
 - [sinc function in Wikipedia](#)

Decompositions, Factorisations, Inverses and Equation Solvers (Dense Matrices)

R = chol(X) (form 1)

R = chol(X, layout) (form 2)

chol(R, X) (form 3)

chol(R, X, layout) (form 4)

chol(R, P, X, layout, "vector") (form 5)

chol(R, P, X, layout, "matrix") (form 6)

- Cholesky decomposition of symmetric/hermitian matrix X into triangular matrix R , with an optional permutation vector/matrix P
- By default, R is upper triangular
- The optional argument *layout* is either "upper" or "lower", which specifies whether R is upper or lower triangular
- Forms 1 to 4 require X to be positive definite
- Forms 5 to 6 require X to be positive semi-definite; these forms use pivoted decomposition and provide a permutation vector/matrix P with type `uvec` or `umat`
- The decomposition has the following form:
 - forms 1 and 3: $X = R.t() * R$
 - forms 2 and 4 with *layout* = "upper": $X = R.t() * R$
 - forms 2 and 4 with *layout* = "lower": $X = R * R.t()$
 - form 5 with *layout* = "upper": $X(P,P) = R.t() * R$, where $X(P,P)$ is a **non-contiguous view** of X
 - form 5 with *layout* = "lower": $X(P,P) = R * R.t()$, where $X(P,P)$ is a **non-contiguous view** of X
 - form 6 with *layout* = "upper": $X = P * R.t() * R * P.t()$
 - form 6 with *layout* = "lower": $X = P * R * R.t() * P.t()$
- If the decomposition fails:
 - the forms $R = chol(X)$ and $R = chol(X, layout)$ reset R and throw a `std::runtime_error` exception
 - the other forms reset R and P , and return a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 5, fill::randu);
mat X = A.t()*A;

mat R1 = chol(X);
mat R2 = chol(X, "lower");

mat R3;
bool ok = chol(R3, X);

mat R;
uvec P_vec;
umat P_mat;

chol(R, P_vec, X, "upper", "vector");
chol(R, P_mat, X, "lower", "matrix");
```

- See also:

- [sqrtmat\(\)](#)
- [lu\(\)](#)
- [qr\(\)](#)
- [.is_sympd\(\)](#)
- [Cholesky decomposition in MathWorld](#)
- [Cholesky decomposition in Wikipedia](#)
- [Definite matrix in Wikipedia](#)

vec eigval = eig_sym(X)

eig_sym(eigval, X)

eig_sym(eigval, eigvec, X)

eig_sym(eigval, eigvec, X, method)

- Eigen decomposition of **dense** symmetric/hermitian matrix X
- The eigenvalues and corresponding eigenvectors are stored in *eigval* and *eigvec*, respectively
- The eigenvalues are in ascending order
- The eigenvectors are stored as column vectors
- If X is not square sized, a *std::logic_error* exception is thrown
- The *method* argument is optional; *method* is either "dc" or "std"
 - "dc" indicates divide-and-conquer method (default setting)
 - "std" indicates standard method
 - the divide-and-conquer method provides slightly different results than the standard method, but is considerably faster for large matrices
- If the decomposition fails:
 - *eigval = eig_sym(X)* resets *eigval* and throws a *std::runtime_error* exception
 - *eig_sym(eigval,X)* resets *eigval* and returns a bool set to *false* (exception is not thrown)
 - *eig_sym(eigval,eigvec,X)* resets *eigval* & *eigvec* and returns a bool set to *false* (exception is not thrown)

- Examples:

```
// for matrices with real elements

mat A(50, 50, fill::randu);
mat B = A.t()*A; // generate a symmetric matrix

vec eigval;
mat eigvec;
```



```
eig_sym(eigval, eigvec, B);

// for matrices with complex elements

cx_mat C(50, 50, fill::randu);
cx_mat D = C.t()*C;

    vec eigval2;
cx_mat eigvec2;

eig_sym(eigval2, eigvec2, D);
```

- See also:
 - [eig_gen\(\)](#)
 - [eig_pair\(\)](#)
 - [svd\(\)](#)
 - [svd_econ\(\)](#)
 - [princomp\(\)](#)
 - [eigs_sym\(\)](#)
 - [.is_symmetric\(\)](#)
 - [.is_hermitian\(\)](#)
 - [eigen decomposition in MathWorld](#)
 - [eigenvalues & eigenvectors in Wikipedia](#)
 - [divide & conquer eigenvalue algorithm in Wikipedia](#)

```
cx_vec eigval = eig_gen( X )  
cx_vec eigval = eig_gen( X, bal )
```

```
eig_gen( eigval, X )  
eig_gen( eigval, X, bal )
```

```
eig_gen( eigval, eigvec, X )  
eig_gen( eigval, eigvec, X, bal )
```

```
eig_gen( eigval, leigvec, reigvec, X )  
eig_gen( eigval, leigvec, reigvec, X, bal )
```

- Eigen decomposition of **dense** general (non-symmetric/non-hermitian) square matrix X
- The eigenvalues and corresponding right eigenvectors are stored in *eigval* and *eigvec*, respectively
- If both left and right eigenvectors are requested they are stored in *leigvec* and *reigvec*, respectively
- The eigenvectors are stored as column vectors
- The *bal* argument is optional; *bal* is one of:
 - "balance" \mapsto diagonally scale and permute X to improve conditioning of the eigenvalues
 - "nobalance" \mapsto do not balance X ; this is the **default operation**
- If X is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails:
 - *eigval = eig_gen(X)* resets *eigval* and throws a *std::runtime_error* exception
 - *eig_gen(eigval,X)* resets *eigval* and returns a bool set to *false* (exception is not thrown)
 - *eig_gen(eigval,eigvec,X)* resets *eigval* & *eigvec* and returns a bool set to *false* (exception is not thrown)
 - *eig_gen(eigval,leigvec,reigvec,X)* resets *eigval*, *leigvec* & *reigvec* and returns a bool set to *false* (exception is not thrown)
- Examples:

```
mat A(10, 10, fill::randu);
```

```
cx_vec eigval;  
cx_mat eigvec;  
  
eig_gen(eigval, eigvec, A);  
eig_gen(eigval, eigvec, A, "balance");
```

- See also:

- [eig_pair\(\)](#)
- [eig_sym\(\)](#)
- [svd\(\)](#)
- [svd_econ\(\)](#)
- [schur\(\)](#)
- [eigs_gen\(\)](#)
- [eigen decomposition in MathWorld](#)
- [eigenvalues & eigenvectors in Wikipedia](#)

cx_vec eigval = eig_pair(A, B)

eig_pair(eigval, A, B)

eig_pair(eigval, eigvec, A, B)

eig_pair(eigval, leigvec, reigvec, A, B)

- Eigen decomposition for pair of general **dense** square matrices A and B of the same size, such that $A * eigvec = B * eigvec * diagmat(eigval)$
- The eigenvalues and corresponding right eigenvectors are stored in *eigval* and *eigvec*, respectively
- If both left and right eigenvectors are requested they are stored in *leigvec* and *reigvec*, respectively
- The eigenvectors are stored as column vectors
- If A or B is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails:
 - *eigval = eig_pair(A,B)* resets *eigval* and throws a *std::runtime_error* exception
 - *eig_pair(eigval,A,B)* resets *eigval* and returns a bool set to *false* (exception is not thrown)
 - *eig_pair(eigval,eigvec,A,B)* resets *eigval* & *eigvec* and returns a bool set to *false* (exception is not thrown)
 - *eig_pair(eigval,leigvec,reigvec,A,B)* resets *eigval*, *leigvec* & *reigvec* and returns a bool set to *false* (exception is not thrown)
- Examples:

```
mat A(10, 10, fill::randu);
mat B(10, 10, fill::randu);

cx_vec eigval;
cx_mat eigvec;

eig_pair(eigval, eigvec, A, B);
```

- See also:
 - `eig_gen()`
 - `eig_sym()`
 - `qz()`
 - eigen decomposition in MathWorld
 - eigenvalues & eigenvectors in Wikipedia

H = hess(X)

hess(H, X)

hess(U, H, X)

- Upper Hessenberg decomposition of square matrix X , such that $X = U*H*U.t()$
- U is a unitary matrix containing the Hessenberg vectors
- H is a square matrix known as the upper Hessenberg matrix, with elements below the first subdiagonal set to zero
- If X is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails:
 - $H = \text{hess}(X)$ resets H and throws a *std::runtime_error* exception
 - $\text{hess}(H,X)$ resets H and returns a bool set to *false* (exception is not thrown)
 - $\text{hess}(U,H,X)$ resets U & H and returns a bool set to *false* (exception is not thrown)
- **Caveat:** in general, upper Hessenberg decomposition is not unique

- Examples:

```
mat X(20,20, fill::randu);
```

```
mat U;
```

```
mat H;
```

```
hess(U, H, X);
```

- See also:
 - [qz\(\)](#)
 - [schur\(\)](#)
 - [Hessenberg decomposition in MathWorld](#)
 - [Hessenberg matrix in Wikipedia](#)

B = inv(A)
B = inv(A, settings)

inv(B, A)
inv(B, A, settings)

inv(B, rcond, A)

- Inverse of general square matrix *A*
- The *settings* argument is optional; it is one of the following:

<code>inv_opts::no_ugly</code>	do not provide inverses for poorly conditioned matrices (where $rcond < A.n_rows \cdot datum::eps$)
<code>inv_opts::allow_approx</code>	allow approximate inverses for rank deficient or poorly conditioned matrices

- The reciprocal condition number is optionally calculated and stored in *rcond*
 - *rcond* close to 1 suggests that *A* is well-conditioned
 - *rcond* close to 0 suggests that *A* is badly conditioned
- If *A* is not square sized, a *std::logic_error* exception is thrown
- If *A* appears to be singular:
 - *B = inv(A)* resets *B* and throws a *std::runtime_error* exception
 - *inv(B,A)* resets *B* and returns a bool set to *false* (exception is not thrown)
 - *inv(B,rcond,A)* resets *B*, sets *rcond* to zero, and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - if matrix *A* is known to be symmetric positive definite, using *inv_sympd()* is faster
 - if matrix *A* is known to be diagonal, use *inv(diagmat(A))*
 - if matrix *A* is known to be triangular, use *inv(trimatu(A))* or *inv(trimatl(A))*
 - to solve a system of linear equations, such as $Z = inv(X)*Y$, using *solve()* can be faster and/or more accurate
- Examples:

```
mat A(5, 5, fill::randu);

mat B = inv(A);

mat C;
bool success = inv(C, A);

mat D;
double rcond_val;
inv(D, rcond_val, A);

A.col(1).zeros();
mat E;
inv(E, A, inv_opts::allow_approx);
```

- See also:

- [.i\(\)](#)
- [inv_sympd\(\)](#)
- [rcond\(\)](#)
- [pinv\(\)](#)
- [solve\(\)](#)
- [diagmat\(\)](#)
- [trimatu\(\)](#) / [trimatl\(\)](#)
- [powmat\(\)](#)
- [matrix inverse](#) in MathWorld
- [invertible matrix](#) in Wikipedia

B = inv_sympd(A)
B = inv_sympd(A, settings)

inv_sympd(B, A)
inv_sympd(B, A, settings)

inv_sympd(B, rcond, A)

- Inverse of symmetric/hermitian positive definite matrix *A*
- The *settings* argument is optional; it is one of the following:

<code>inv_opts::no_ugly</code>	do not provide inverses for poorly conditioned matrices (where <i>rcond</i> < <i>A.n_rows</i> · <i>datum::eps</i>)
<code>inv_opts::allow_approx</code>	allow approximate inverses for rank deficient or poorly conditioned symmetric matrices

- The reciprocal condition number is optionally calculated and stored in *rcond*
 - *rcond* close to 1 suggests that *A* is well-conditioned
 - *rcond* close to 0 suggests that *A* is badly conditioned
- If *A* is not square sized, a *std::logic_error* exception is thrown
- If *A* appears to be singular or not positive definite:
 - *B = inv_sympd(A)* resets *B* and throws a *std::runtime_error* exception
 - *inv_sympd(B,A)* resets *B* and returns a bool set to *false* (exception is not thrown)
 - *inv_sympd(B,rcond,A)* resets *B*, sets *rcond* to zero, and returns a bool set to *false* (exception is not thrown)
- **Caveat:** to solve a system of linear equations, such as $Z = \text{inv}(X) * Y$, using `solve()` can be faster and/or more accurate
- Examples:

```
mat A(5, 5, fill::randu);  
mat B = A.t() * A;
```

```
mat C = inv_sympd(B);

mat D;
bool success = inv_sympd(D, B);

mat E;
double rcond_val;
inv_sympd(E, rcond_val, B);

B.col(1).zeros();
B.row(1).zeros();
mat F;
inv_sympd(F, B, inv_opts::allow_approx);
```

- See also:

- [inv\(\)](#)
- [rcond\(\)](#)
- [pinv\(\)](#)
- [solve\(\)](#)
- [eig_sym\(\)](#)
- [.is_sympd\(\)](#)
- [matrix inverse in MathWorld](#)
- [invertible matrix in Wikipedia](#)
- [positive definite matrix in MathWorld](#)
- [positive definite matrix in Wikipedia](#)

lu(L, U, P, X)

lu(L, U, X)

- Lower-upper decomposition (with partial pivoting) of matrix X
- The first form provides a lower-triangular matrix L , an upper-triangular matrix U , and a permutation matrix P , such that $P.t()*L*U = X$
- The second form provides permuted L and U , such that $L*U = X$; note that in this case L is generally not lower-triangular
- If the decomposition fails:
 - `lu(L,U,P,X)` resets L , U , P and returns a bool set to *false* (exception is not thrown)
 - `lu(L,U,X)` resets L , U and returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 5, fill::randu);
```

```
mat L, U, P;
```

```
lu(L, U, P, A);
```

```
mat B = P.t()*L*U;
```

- See also:
 - [chol\(\)](#)
 - [LU decomposition in Wikipedia](#)
 - [LU decomposition in MathWorld](#)

B = null(A)
B = null(A, tolerance)

null(B, A)
null(B, A, tolerance)

- Find the orthonormal basis of the null space of matrix A
- The dimension of the range space is the number of singular values of A not greater than *tolerance*
- The *tolerance* argument is optional; by default $tolerance = max_rc \cdot max_sv \cdot epsilon$, where:
 - $max_rc = \max(A.n_rows, A.n_cols)$
 - max_sv = maximum singular value of A (ie. spectral norm)
 - $epsilon$ = difference between 1 and the least value greater than 1 that is representable
- The computation is based on singular value decomposition; if the decomposition fails:
 - $B = null(A)$ resets B and throws a `std::runtime_error` exception
 - $null(B,A)$ resets B and returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 6, fill::randu);  
  
A.row(0).zeros();  
A.col(0).zeros();  
  
mat B = null(A);
```

- See also:
 - [orth\(\)](#)
 - [qr\(\)](#)
 - [svd\(\)](#)
 - [rank\(\)](#)
 - [datum::eps](#)
 - [Orthonormal basis in Wikipedia](#)

B = orth(A)
B = orth(A, tolerance)

orth(B, A)
orth(B, A, tolerance)

- Find the orthonormal basis of the range space of matrix A , so that $B.t()*B \approx \text{eye}(r,r)$, where $r = \text{rank}(A)$
- The dimension of the range space is the number of singular values of A greater than *tolerance*
- The *tolerance* argument is optional; by default $\text{tolerance} = \text{max_rc} \cdot \text{max_sv} \cdot \text{epsilon}$, where:
 - $\text{max_rc} = \max(A.\text{n_rows}, A.\text{n_cols})$
 - max_sv = maximum singular value of A (ie. spectral norm)
 - epsilon = difference between 1 and the least value greater than 1 that is representable
- The computation is based on singular value decomposition; if the decomposition fails:
 - $B = \text{orth}(A)$ resets B and throws a `std::runtime_error` exception
 - $\text{orth}(B,A)$ resets B and returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 6, fill::randu);
```

```
mat B = orth(A);
```

- See also:
 - `null()`
 - `qr()`
 - `svd()`
 - `rank()`
 - `datum::eps`
 - [Orthonormal basis in Wikipedia](#)

B = pinv(A)
B = pinv(A, tolerance)
B = pinv(A, tolerance, method)

pinv(B, A)
pinv(B, A, tolerance)
pinv(B, A, tolerance, method)

- Moore-Penrose pseudo-inverse (generalised inverse) of matrix *A*
- The computation is based on singular value decomposition
- The *tolerance* argument is optional; by default *tolerance* = *max_rc* · *max_sv* · *epsilon*, where:
 - *max_rc* = max(A.n_rows, A.n_cols)
 - *max_sv* = maximum singular value of *A* (ie. spectral norm)
 - *epsilon* = difference between 1 and the least value greater than 1 that is representable
- Any singular values less than *tolerance* are treated as zero
- The *method* argument is optional; *method* is either "dc" or "std"
 - "dc" indicates divide-and-conquer method (default setting)
 - "std" indicates standard method
 - the divide-and-conquer method provides slightly different results than the standard method, but is considerably faster for large matrices
- If the decomposition fails:
 - *B = pinv(A)* resets *B* and throws a *std::runtime_error* exception
 - *pinv(B,A)* resets *B* and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - to find approximate solutions to under-/over-determined or rank deficient systems of linear equations, *solve()* can be considerably faster and/or more accurate
 - if the given matrix *A* is square-sized and only occasionally rank deficient, using *inv()* or *inv_sympd()* with the *inv_opts::allow_approx* option is faster
- Examples:

```
mat A(4, 5, fill::randu);  
  
mat B = pinv(A);          // use default tolerance  
  
mat C = pinv(A, 0.01);    // set tolerance to 0.01
```

- See also:

- [.i\(\)](#)
- [inv\(\)](#)
- [solve\(\)](#)
- [svd\(\)](#)
- [datum::eps](#)
- [Pseudoinverse in MathWorld](#)
- [Moore-Penrose Matrix Inverse in MathWorld](#)
- [Moore-Penrose inverse in Wikipedia](#)

qr(Q, R, X) (form 1)

qr(Q, R, P, X, "vector") (form 2)

qr(Q, R, P, X, "matrix") (form 3)

- Decomposition of X into an orthogonal matrix Q and a right triangular matrix R , with an optional permutation matrix/vector P
 - form 1: decomposition has the form $Q^*R = X$
 - form 2: P is permutation vector with type **uvec**; decomposition has the form $Q^*R = X.cols(P)$
 - form 3: P is permutation matrix with type **umat**; decomposition has the form $Q^*R = X^*P$
- If P is specified, a column pivoting decomposition is used; the diagonal entries of R are ordered from largest to smallest magnitude
- If the decomposition fails, Q , R and P are reset and the function returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat X(5, 5, fill::randu);
```

```
mat Q;  
mat R;
```

```
qr(Q, R, X);
```

```
uvec P_vec;  
umat P_mat;
```

```
qr(Q, R, P_vec, X, "vector");  
qr(Q, R, P_mat, X, "matrix");
```

- See also:

- **qr_econ()**
- **chol()**
- **orth()**
- **orthogonal matrix in Wikipedia**
- **QR decomposition in Wikipedia**
- **QR decomposition in MathWorld**

qr_econ(Q, R, X)

- Economical decomposition of X (with size $m \times n$) into an orthogonal matrix Q and a right triangular matrix R , such that $Q \cdot R = X$
- If $m > n$, only the first n rows of R and the first n columns of Q are calculated (ie. the zero rows of R and the corresponding columns of Q are omitted)
- If the decomposition fails, Q and R are reset and the function returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat X(6, 5, fill::randu);
```

```
mat Q;  
mat R;
```

```
qr_econ(Q, R, X);
```

- See also:

- [qr\(\)](#)
- [orthogonal matrix in Wikipedia](#)
- [QR decomposition in Wikipedia](#)
- [QR decomposition in Octave](#)
- [QR decomposition in MathWorld](#)

qz(AA, BB, Q, Z, A, B)
qz(AA, BB, Q, Z, A, B, select)

- Generalised Schur decomposition for pair of general square matrices A and B of the same size, such that $A = Q.t()*AA*Z.t()$ and $B = Q.t()*BB*Z.t()$
- The *select* argument is optional and specifies the ordering of the top left of the Schur form; it is one of the following:
 - "none" no ordering (default operation)
 - "lhp" left-half-plane: eigenvalues with real part < 0
 - "rhp" right-half-plane: eigenvalues with real part > 0
 - "iuc" inside-unit-circle: eigenvalues with absolute value < 1
 - "ouc" outside-unit-circle: eigenvalues with absolute value > 1
- The left and right Schur vectors are stored in Q and Z , respectively
- In the complex-valued problem, the generalised eigenvalues are found in $diagvec(AA) / diagvec(BB)$
- If A or B is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails, AA , BB , Q and Z are reset, and the function returns a bool set to *false* (exception is not thrown)
- Examples:

```
mat A(10, 10, fill::randu);
mat B(10, 10, fill::randu);

mat AA;
mat BB;
mat Q;
mat Z;

qz(AA, BB, Q, Z, A, B);
```

- See also:
 - [hess\(\)](#)

- `schur()`
- `eig_pair()`
- generalised Schur decomposition in Wikipedia

S = schur(X)

schur(S, X)

schur(U, S, X)

- Schur decomposition of square matrix X , such that $X = U*S*U.t()$
- U is a unitary matrix containing the Schur vectors
- S is an upper triangular matrix, called the Schur form of X
- If X is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails:
 - $S = \text{schur}(X)$ resets S and throws a *std::runtime_error* exception
 - $\text{schur}(S,X)$ resets S and returns a bool set to *false* (exception is not thrown)
 - $\text{schur}(U,S,X)$ resets U & S and returns a bool set to *false* (exception is not thrown)
- **Caveat:** in general, Schur decomposition is not unique

- Examples:

```
mat X(20,20, fill::randu);
```

```
mat U;  
mat S;
```

```
schur(U, S, X);
```

- See also:
 - [hess\(\)](#)
 - [qz\(\)](#)
 - [eig_gen\(\)](#)
 - [Schur decomposition in MathWorld](#)
 - [Schur decomposition in Wikipedia](#)

X = solve(A, B)
X = solve(A, B, settings)

solve(X, A, B)
solve(X, A, B, settings)

- Solve a **dense** system of linear equations, $A*X = B$, where X is unknown; similar functionality to the `\` operator in Matlab/Octave, ie. $X = A \setminus B$
- A can be square sized (critically determined system), or non-square (under/over-determined system); A can be rank deficient
- B can be a vector or matrix
- The number of rows in A and B must be the same
- By default, matrix A is analysed to automatically determine whether it is a general matrix, band matrix, diagonal matrix, or symmetric/hermitian positive definite (SPD) matrix; based on the detected matrix structure, a specialised solver is used for faster execution; if no solution is found, an approximate solver is automatically used as a fallback; see the [associated paper](#) for more details
- If A is known to be a triangular matrix, the solution can be computed faster by explicitly indicating that A is triangular through `trimatu()` or `trimatl()`; see examples below
- The *settings* argument is optional; it is one of the following, or a combination thereof:

<code>solve_opts::fast</code>	fast mode: disable determining solution quality via <code>rcond</code> , disable iterative refinement, disable equilibration
<code>solve_opts::refine</code>	apply iterative refinement to improve solution quality (matrix A must be square)
<code>solve_opts::equilibrate</code>	equilibrate the system before solving (matrix A must be square)
<code>solve_opts::likely_sympd</code>	indicate that matrix A is likely symmetric/hermitian positive definite
<code>solve_opts::allow_ugly</code>	keep solutions of systems that are singular to working precision
<code>solve_opts::no_approx</code>	do not find approximate solutions for rank deficient systems
<code>solve_opts::no_band</code>	do not use specialised solver for band matrices or diagonal matrices
<code>solve_opts::no_trimat</code>	do not use specialised solver for triangular matrices

`solve_opts::no_sympd` do not use specialised solver for symmetric/hermitian positive definite matrices
`solve_opts::force_approx` skip the standard solver and directly use of the approximate solver

the above settings can be combined using the `+` operator; for example: `solve_opts::fast + solve_opts::no_approx`

- If a rank deficient system is detected and the `solve_opts::no_approx` option is **not** enabled, a warning is emitted and an approximate solution is attempted;
since Armadillo 10.4, this warning can be disabled by setting `ARMA_WARN_LEVEL` to 1 before including the `armadillo` header:

```
#define ARMA_WARN_LEVEL 1
#include <armadillo>
```

- **Caveats:**

- using `solve_opts::fast` will speed up finding the solution, but for poorly conditioned systems the solution may have lower quality
- not all SPD matrices are automatically detected; to skip the analysis step and directly indicate that matrix A is likely SPD, use `solve_opts::likely_sympd`
- using `solve_opts::force_approx` is only advised if the system is known to be rank deficient; the approximate solver is considerably slower

- If no solution is found:
 - `X = solve(A,B)` resets X and throws a `std::runtime_error` exception
 - `solve(X,A,B)` resets X and returns a bool set to `false` (exception is not thrown)

- Implementation details are available in the following paper:
Conrad Sanderson and Ryan Curtin.
[An Adaptive Solver for Systems of Linear Equations.](#)
International Conference on Signal Processing and Communication Systems, 2020.

- Examples:

```
mat A(5, 5, fill::randu);
vec b(5, fill::randu);
mat B(5, 5, fill::randu);
```

```
vec x1 = solve(A, b);
```

```
vec x2;
```

```
bool status = solve(x2, A, b);  
  
mat X1 = solve(A, B);  
  
mat X2 = solve(A, B, solve_opts::fast); // enable fast mode  
  
mat X3 = solve(trimatu(A), B); // indicate that A is triangular
```

- See also:

- [inv\(\)](#)
- [pinv\(\)](#)
- [rcond\(\)](#)
- [roots\(\)](#)
- [sylv\(\)](#)
- [spsolve\(\)](#) - solve sparse system of linear equations
- [linear system of equations in MathWorld](#)
- [system of linear equations in Wikipedia](#)
- [band matrix in Wikipedia](#)
- [definiteness of a matrix in Wikipedia](#)
- [positive definite matrix in MathWorld](#)
- [iterative refinement](#)
- [ensmallen](#) - C++ library for solving numerical optimisation problems

vec s = svd(X)

svd(vec s, X)

svd(mat U, vec s, mat V, mat X)

svd(mat U, vec s, mat V, mat X, method)

svd(cx_mat U, vec s, cx_mat V, cx_mat X)

svd(cx_mat U, vec s, cx_mat V, cx_mat X, method)

- Singular value decomposition of **dense** matrix X
- If X is square, it can be reconstructed using $X = U*diagmat(s)*V.t()$
- The singular values are in descending order
- The *method* argument is optional; *method* is either "dc" or "std"
 - "dc" indicates divide-and-conquer method (default setting)
 - "std" indicates standard method
 - the divide-and-conquer method provides slightly different results than the standard method, but is considerably faster for large matrices
- If the decomposition fails, the output objects are reset and:
 - $s = svd(X)$ resets s and throws a `std::runtime_error` exception
 - $svd(s,X)$ resets s and returns a bool set to *false* (exception is not thrown)
 - $svd(U,s,V,X)$ resets U , s , V and returns a bool set to *false* (exception is not thrown)
- Examples:

```
mat X(5, 5, fill::randu);
```

```
mat U;
```

```
vec s;
```

```
mat V;
```

```
svd(U,s,V,X);
```

- See also:
 - `svd_econ()`
 - `eig_gen()`
 - `eig_sym()`
 - `princomp()`
 - `pinv()`
 - `svds()`
 - [singular value decomposition in Wikipedia](#)
 - [singular value decomposition in MathWorld](#)

```
svd_econ( mat U, vec s, mat V, mat X )  
svd_econ( mat U, vec s, mat V, mat X, mode )  
svd_econ( mat U, vec s, mat V, mat X, mode, method )
```

```
svd_econ( cx_mat U, vec s, cx_mat V, cx_mat X )  
svd_econ( cx_mat U, vec s, cx_mat V, cx_mat X, mode )  
svd_econ( cx_mat U, vec s, cx_mat V, cx_mat X, mode, method )
```

- Economical singular value decomposition of **dense** matrix X
- The singular values are in descending order
- The *mode* argument is optional; *mode* is one of:
 - "both" = compute both left and right singular vectors (default operation)
 - "left" = compute only left singular vectors
 - "right" = compute only right singular vectors
- The *method* argument is optional; *method* is either "dc" or "std"
 - "dc" indicates divide-and-conquer method (default setting)
 - "std" indicates standard method
 - the divide-and-conquer method provides slightly different results than the standard method, but is considerably faster for large matrices
- If the decomposition fails, U , s , V are reset and a bool set to *false* is returned (exception is not thrown)
- Examples:

```
mat X(4, 5, fill::randu);  
  
mat U;  
vec s;  
mat V;  
  
svd_econ(U, s, V, X);
```

- See also:
 - `svd()`

- `eig_gen()`
- `eig_sym()`
- `princomp()`
- `svds()`
- singular value decomposition in Wikipedia
- singular value decomposition in MathWorld

`X = syl(A, B, C)`
`syl(X, A, B, C)`

- Solve the Sylvester equation, ie. $AX + XB + C = 0$, where X is unknown
- Matrices A , B and C must be square sized
- If no solution is found:
 - `syl(A,B,C)` resets X and throws a `std::runtime_error` exception
 - `syl(X,A,B,C)` resets X and returns a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 5, fill::randu);  
mat B(5, 5, fill::randu);  
mat C(5, 5, fill::randu);
```

```
mat X1 = syl(A, B, C);
```

```
mat X2;  
syl(X2, A, B, C);
```

- See also:
 - [solve\(\)](#)
 - [Sylvester equation in Wikipedia](#)

Decompositions, Factorisations and Equation Solvers (Sparse Matrices)

```

vec eigval = eigs_sym( X, k )
vec eigval = eigs_sym( X, k, form )
vec eigval = eigs_sym( X, k, form, opts )
vec eigval = eigs_sym( X, k, sigma )
vec eigval = eigs_sym( X, k, sigma, opts )

```

```

eigs_sym( eigval, X, k )
eigs_sym( eigval, X, k, form )
eigs_sym( eigval, X, k, form, opts )
eigs_sym( eigval, X, k, sigma )
eigs_sym( eigval, X, k, sigma, opts )

```

```

eigs_sym( eigval, eigvec, X, k )
eigs_sym( eigval, eigvec, X, k, form )
eigs_sym( eigval, eigvec, X, k, form, opts )
eigs_sym( eigval, eigvec, X, k, sigma )
eigs_sym( eigval, eigvec, X, k, sigma, opts )

```

- Obtain a limited number of eigenvalues and eigenvectors of **sparse** symmetric real matrix X
- k specifies the number of eigenvalues and eigenvectors
- The argument *form* is optional; *form* is one of:
 - "lm" = obtain eigenvalues with largest magnitude (default operation)
 - "sm" = obtain eigenvalues with smallest magnitude (see the caveats below)
 - "la" = obtain eigenvalues with largest algebraic value
 - "sa" = obtain eigenvalues with smallest algebraic value
- The argument *sigma* is optional; if *sigma* is given, eigenvalues closest to *sigma* are found via shift-invert mode
NOTE: to use *sigma*, both `ARMA_USE_ARPACK` and `ARMA_USE_SUPERLU` must be enabled in `config.hpp`
- The *opts* argument is optional; *opts* is an instance of the *eigs_opts* structure:

```

struct eigs_opts
{
    double      tol;      // default: 0

```

```

unsigned int maxiter; // default: 1000
unsigned int subdim; // default: max(2*k+1, 20)
};

```

- *tol* specifies the tolerance for convergence
 - *maxiter* specifies the maximum number of Arnoldi iterations
 - *subdim* specifies the dimension of the Krylov subspace, with the constraint $k < \text{subdim} \leq X.\text{n_rows}$; recommended value is $\text{subdim} \geq 2*k$
- The eigenvalues and corresponding eigenvectors are stored in *eigval* and *eigvec*, respectively
 - If *X* is not square sized, a *std::logic_error* exception is thrown
 - If the decomposition fails:
 - *eigval = eigs_sym(X,k)* resets *eigval* and throws a *std::runtime_error* exception
 - *eigs_sym(eigval,X,k)* resets *eigval* and returns a bool set to *false* (exception is not thrown)
 - *eigs_sym(eigval,eigvec,X,k)* resets *eigval* & *eigvec* and returns a bool set to *false* (exception is not thrown)
 - **Caveats:**
 - the number of obtained eigenvalues/eigenvectors may be lower than requested, depending on the given data
 - if the decomposition fails, try first increasing *opts.subdim* (Krylov subspace dimension), and, as secondary options, try increasing *opts.maxiter* (maximum number of iterations), and/or *opts.tol* (tolerance for convergence), and/or *k* (number of eigenvalues)
 - for an alternative to the "sm" form, use the shift-invert mode with *sigma* set to 0.0
 - **NOTE:** the implementation in Armadillo 12.6 is considerably faster than earlier versions; further speedups can be obtained by enabling OpenMP in your compiler (eg. *-fopenmp* in GCC and clang)
 - Examples:

```

// generate sparse matrix
sp_mat A = sprandu<sp_mat>(1000, 1000, 0.1);
sp_mat B = A.t()*A;

vec eigval;
mat eigvec;

```



```
eigs_sym(eigval, eigvec, B, 5); // find 5 eigenvalues/eigenvectors

eigs_opts opts;
opts.maxiter = 10000;          // increase max iterations to 10000

eigs_sym(eigval, eigvec, B, 5, "lm", opts);
```

- See also:

- [eigs_gen\(\)](#)
- [eig_sym\(\)](#)
- [svds\(\)](#)
- [.is_symmetric\(\)](#)
- [shift-invert mode in ARPACK](#)
- [eigen decomposition in MathWorld](#)
- [eigenvalues & eigenvectors in Wikipedia](#)

```
cx_vec eigval = eigs_gen( X, k )
cx_vec eigval = eigs_gen( X, k, form )
cx_vec eigval = eigs_gen( X, k, sigma )
cx_vec eigval = eigs_gen( X, k, form, opts )
cx_vec eigval = eigs_gen( X, k, sigma, opts )
```

```
eigs_gen( eigval, X, k )
eigs_gen( eigval, X, k, form )
eigs_gen( eigval, X, k, sigma )
eigs_gen( eigval, X, k, form, opts )
eigs_gen( eigval, X, k, sigma, opts )
```

```
eigs_gen( eigval, eigvec, X, k )
eigs_gen( eigval, eigvec, X, k, form )
eigs_gen( eigval, eigvec, X, k, sigma )
eigs_gen( eigval, eigvec, X, k, form, opts )
eigs_gen( eigval, eigvec, X, k, sigma, opts )
```

- Obtain a limited number of eigenvalues and eigenvectors of **sparse** general (non-symmetric/non-hermitian) square matrix X
- k specifies the number of eigenvalues and eigenvectors
- The argument *form* is optional; *form* is one of:
 - "lm" = obtain eigenvalues with largest magnitude (default operation)
 - "sm" = obtain eigenvalues with smallest magnitude (see the caveats below)
 - "lr" = obtain eigenvalues with largest real part
 - "sr" = obtain eigenvalues with smallest real part
 - "li" = obtain eigenvalues with largest imaginary part
 - "si" = obtain eigenvalues with smallest imaginary part
- The argument *sigma* is optional; if *sigma* is given, eigenvalues closest to *sigma* are found via shift-invert mode
NOTE: to use *sigma*, both `ARMA_USE_ARPACK` and `ARMA_USE_SUPERLU` must be enabled in `config.hpp`

- The *opts* argument is optional; *opts* is an instance of the *eigs_opts* structure:

```
struct eigs_opts
{
    double      tol;      // default: 0
    unsigned int maxiter; // default: 1000
    unsigned int subdim;  // default: max(2*k+1, 20)
};
```

- *tol* specifies the tolerance for convergence
 - *maxiter* specifies the maximum number of Arnoldi iterations
 - *subdim* specifies the dimension of the Krylov subspace, with the constraint $k + 2 < \text{subdim} \leq X.\text{n_rows}$; recommended value is $\text{subdim} \geq 2*k + 1$
- The eigenvalues and corresponding eigenvectors are stored in *eigval* and *eigvec*, respectively
- If *X* is not square sized, a *std::logic_error* exception is thrown
- If the decomposition fails:
 - *eigval = eigs_gen(X,k)* resets *eigval* and throws a *std::runtime_error* exception
 - *eigs_gen(eigval,X,k)* resets *eigval* and returns a bool set to *false* (exception is not thrown)
 - *eigs_gen(eigval,eigvec,X,k)* resets *eigval* & *eigvec* and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - the number of obtained eigenvalues/eigenvectors may be lower than requested, depending on the given data
 - if the decomposition fails, try first increasing *opts.subdim* (Krylov subspace dimension) and, as secondary options, try increasing *opts.maxiter* (maximum number of iterations), and/or *opts.tol* (tolerance for convergence), and/or *k* (number of eigenvalues)
 - for an alternative to the "sm" form, use the shift-invert mode with *sigma* set to 0.0
- **NOTE:** the implementation in Armadillo 12.6 is considerably faster than earlier versions; further speedups can be obtained by enabling OpenMP in your compiler (eg. *-fopenmp* in GCC and clang)
- Examples:

```
// generate sparse matrix
sp_mat A = sprandu<sp_mat>(1000, 1000, 0.1);
```

```
cx_vec eigval;  
cx_mat eigvec;  
  
eigs_gen(eigval, eigvec, A, 5); // find 5 eigenvalues/eigenvectors  
  
eigs_opts opts;  
opts.maxiter = 10000;           // increase max iterations to 10000  
  
eigs_gen(eigval, eigvec, A, 5, "lm", opts);
```

- See also:

- [eigs_sym\(\)](#)
- [eig_gen\(\)](#)
- [svds\(\)](#)
- [shift-invert mode in ARPACK](#)
- [eigen decomposition in MathWorld](#)
- [eigenvalues & eigenvectors in Wikipedia](#)

```
vec s = svds( X, k )  
vec s = svds( X, k, tol )
```

```
svds( vec s, X, k )  
svds( vec s, X, k, tol )
```

```
svds( mat U, vec s, mat V, sp_mat X, k )  
svds( mat U, vec s, mat V, sp_mat X, k, tol )
```

```
svds( cx_mat U, vec s, cx_mat V, sp_cx_mat X, k )  
svds( cx_mat U, vec s, cx_mat V, sp_cx_mat X, k, tol )
```

- Obtain a limited number of singular values and singular vectors (truncated SVD) of **sparse** matrix X
- The singular values and vectors are calculated via sparse eigen decomposition of:
$$\begin{bmatrix} \text{zeros}(X.n_rows, X.n_rows) & X \\ X.t() & \text{zeros}(X.n_cols, X.n_cols) \end{bmatrix}$$
- k specifies the number of singular values and singular vectors
- The singular values are in descending order
- The argument *tol* is optional; it specifies the tolerance for convergence; it is passed as $(tol \div \sqrt{2})$ to `eigs_sym()`
- If the decomposition fails, the output objects are reset and:
 - $s = \text{svds}(X, k)$ resets s and throws a `std::runtime_error` exception
 - $\text{svds}(s, X, k)$ resets s and returns a bool set to *false* (exception is not thrown)
 - $\text{svds}(U, s, V, X, k)$ resets U , s , V and returns a bool set to *false* (exception is not thrown)
- **Caveats:**
 - `svds()` is intended only for finding a few singular values from a large sparse matrix; to find all singular values, use `svd()` instead
 - depending on the given matrix, `svds()` may find fewer singular values than specified
- **NOTE:** the implementation in Armadillo 12.6 is considerably faster than earlier versions; further speedups

can be obtained by enabling OpenMP in your compiler (eg. *-fopenmp* in GCC and clang)

- Examples:

```
sp_mat X = sprandu<sp_mat>(100, 200, 0.1);  
  
mat U;  
vec s;  
mat V;  
  
svds(U, s, V, X, 10);
```

- See also:

- [eigs_gen\(\)](#)
- [eigs_sym\(\)](#)
- [svd\(\)](#)
- [singular value decomposition in Wikipedia](#)
- [singular value decomposition in MathWorld](#)

X = spsolve(A, B)
X = spsolve(A, B, solver)
X = spsolve(A, B, solver, opts)

spsolve(X, A, B)
spsolve(X, A, B, solver)
spsolve(X, A, B, solver, opts)

- Solve a **sparse** system of linear equations, $A \cdot X = B$, where A is a sparse matrix, B is a dense matrix or vector, and X is unknown
- The number of rows in A and B must be the same
- If no solution is found:
 - $X = \text{spsolve}(A, B)$ resets X and throws a `std::runtime_error` exception
 - $\text{spsolve}(X, A, B)$ resets X and returns a bool set to `false` (no exception is thrown)
- The *solver* argument is optional; *solver* is either "superlu" or "lapack"; by default "superlu" is used
 - for "superlu", `ARMA_USE_SUPERLU` must be enabled in `config.hpp`
 - for "lapack", sparse matrix A is converted to a dense matrix before using the LAPACK solver; this considerably increases memory usage
- **Notes:**
 - the SuperLU solver is mainly useful for very large and/or very sparse matrices
 - to reuse the SuperLU factorisation of A for finding solutions where B is iteratively changed, see the `spsolve_factoriser` class
 - if there is sufficient amount of memory to store a dense version of matrix A , the LAPACK solver can be faster
- The *opts* argument is optional and applicable to the SuperLU solver; *opts* is an instance of the `superlu_opts` structure:

```
struct superlu_opts
{
    bool        allow_ugly;    // default: false
    bool        equilibrate;  // default: false
    bool        symmetric;    // default: false
    double      pivot_thresh; // default: 1.0
```

```

permutation_type permutation; // default: superlu_opts::COLAMD
refine_type          refine;   // default: superlu_opts::REF_NONE
};

```

- *allow_ugly* is either *true* or *false*; indicates whether to keep solutions of systems singular to working precision
- *equilibrate* is either *true* or *false*; indicates whether to equilibrate the system (scale the rows and columns of *A* to have unit norm)
- *symmetric* is either *true* or *false*; indicates whether to use SuperLU symmetric mode, which gives preference to diagonal pivots
- *pivot_threshold* is in the range [0.0, 1.0], used for determining whether a diagonal entry is an acceptable pivot (details in SuperLU documentation)
- *permutation* specifies the type of column permutation; it is one of:

<code>superlu_opts::NATURAL</code>	natural ordering
<code>superlu_opts::MMD_ATA</code>	minimum degree ordering on structure of $A.t() * A$
<code>superlu_opts::MMD_AT_PLUS_A</code>	minimum degree ordering on structure of $A.t() + A$
<code>superlu_opts::COLAMD</code>	approximate minimum degree column ordering
- *refine* specifies the type of iterative refinement; it is one of:

<code>superlu_opts::REF_NONE</code>	no refinement
<code>superlu_opts::REF_SINGLE</code>	iterative refinement in single precision
<code>superlu_opts::REF_DOUBLE</code>	iterative refinement in double precision
<code>superlu_opts::REF_EXTRA</code>	iterative refinement in extra precision

- Examples:

```

sp_mat A = sprandu<sp_mat>(1000, 1000, 0.1);

vec b(1000, fill::randu);
mat B(1000, 5, fill::randu);

vec x = spsolve(A, b); // solve one system
mat X = spsolve(A, B); // solve several systems

bool status = spsolve(x, A, b); // use default solver
if(status == false) { cout << "no solution" << endl; }

```



```
spsolve(x, A, b, "lapack" ); // use LAPACK solver
spsolve(x, A, b, "superlu"); // use SuperLU solver

superlu_opts opts;

opts.allow_ugly = true;
opts.equilibrate = true;

spsolve(x, A, b, "superlu", opts);
```

- See also:
 - [spsolve_factoriser](#)
 - [solve\(\)](#) - solve dense system of linear equations
 - [SuperLU home page](#)
 - [linear system of equations in MathWorld](#)
 - [system of linear equations in Wikipedia](#)

spsolve_factoriser

- Class for factorisation of **sparse** matrix A for solving systems of linear equations in the form $A*X = B$
- Allows the SuperLU factorisation of A to be reused for finding solutions in cases where B is iteratively changed
- For an instance of *spsolve_factoriser* named as SF , the member functions are:

SF.factorise(A)

SF.factorise(A, opts)

- factorise square-sized sparse matrix A
- optional settings are given in the *opts* argument as per the `spsolve()` function
- if the factorisation fails, a bool set to *false* is returned

SF.solve(X, B)

- using the given dense matrix B and the computed factorisation, store in X the solution to $A*X = B$
- if computing the solution fails, X is reset and a bool set to *false* is returned

SF.rcond()

- return the 1-norm estimate of the reciprocal condition number computed during the factorisation
 - values close to 1 suggest that the factorised matrix is well-conditioned
 - values close to 0 suggest that the factorised matrix is badly conditioned

SF.reset()

- reset the instance and release all memory related to the stored factorisation; this is automatically done when the instance goes out of scope

- **Notes:**

- if the factorisation of A does not need to be reused, use `spsolve()` instead
- this class internally uses the SuperLU solver; `ARMA_USE_SUPERLU` must be enabled in `config.hpp`

- **Examples:**

```
sp_mat A = sprandu<sp_mat>(1000, 1000, 0.1);  
  
spsolve_factoriser SF;
```

```
bool status = SF.factorise(A);

if(status == false) { cout << "factorisation failed" << endl; }

double rcond_value = SF.rcond();

vec B1(1000, fill::randu);
vec B2(1000, fill::randu);

vec X1;
vec X2;

bool solution1_ok = SF.solve(X1,B1);
bool solution2_ok = SF.solve(X2,B2);

if(solution1_ok == false) { cout << "couldn't find X1" << endl; }
if(solution2_ok == false) { cout << "couldn't find X2" << endl; }
```

- See also:
 - [spsolve\(\)](#)
 - [linear system of equations in MathWorld](#)
 - [system of linear equations in Wikipedia](#)

Signal & Image Processing

conv(A, B)
conv(A, B, shape)

- 1D convolution of vectors A and B
- The orientation of the result vector is the same as the orientation of A (ie. either column or row vector)
- The *shape* argument is optional; it is one of:
 - "full" = return the full convolution (**default setting**), with the size equal to $A.n_elem + B.n_elem - 1$
 - "same" = return the central part of the convolution, with the same size as vector A

- The convolution operation is also equivalent to FIR filtering

- Examples:

```
vec A(256, fill::randu);  
  
vec B(16, fill::randu);  
  
vec C = conv(A, B);  
  
vec D = conv(A, B, "same");
```

- See also:

- [conv2\(\)](#)
- [fft\(\)](#)
- [cor\(\)](#)
- [interp1\(\)](#)
- [Convolution in MathWorld](#)
- [Convolution in Wikipedia](#)
- [FIR filter in Wikipedia](#)

conv2(A, B)
conv2(A, B, shape)

- 2D convolution of matrices A and B
- The *shape* argument is optional; it is one of:
 - "full" = return the full convolution (**default setting**), with the size equal to $size(A) + size(B) - 1$
 - "same" = return the central part of the convolution, with the same size as matrix A
- The implementation of 2D convolution in this version is preliminary

- Examples:

```
mat A(256, 256, fill::randu);
```

```
mat B(16, 16, fill::randu);
```

```
mat C = conv2(A, B);
```

```
mat D = conv2(A, B, "same");
```

- See also:
 - [conv\(\)](#)
 - [fft2\(\)](#)
 - [interp2\(\)](#)
 - [Convolution in MathWorld](#)
 - [Convolution in Wikipedia](#)
 - [Kernel \(image processing\) in Wikipedia](#)

```
cx_mat Y = fft( X )  
cx_mat Y = fft( X, n )
```

```
cx_mat Z = ifft( cx_mat Y )  
cx_mat Z = ifft( cx_mat Y, n )
```

- *fft()*: fast Fourier transform of a vector or matrix (real or complex)
- *ifft()*: inverse fast Fourier transform of a vector or matrix (complex only)
- If given a matrix, the transform is done on each column vector of the matrix
- The optional *n* argument specifies the transform length:
 - if *n* is larger than the length of the input vector, a zero-padded version of the vector is used
 - if *n* is smaller than the length of the input vector, only the first *n* elements of the vector are used
- If *n* is not specified, the transform length is the same as the length of the input vector
- **Caveat:** the transform is fastest when the transform length is a power of 2, eg. 64, 128, 256, 512, 1024, ...
- By default, an internal FFT algorithm based on **KISS FFT** is used
- Since Armadillo 12.0, the **FFTW3 library** can be optionally used for faster execution, as follows:
 - **ARMA_USE_FFTW3** must be defined before including the armadillo header:

```
#define ARMA_USE_FFTW3  
#include <armadillo>
```
 - you will also need to link with the FFTW3 runtime library (eg. `-lfftw3`)
- Examples:

```
vec X(100, fill::randu);  
  
cx_vec Y = fft(X, 128);
```
- See also:
 - **fft2()**

- `conv()`
- `real()`
- fast Fourier transform in MathWorld
- fast Fourier transform in Wikipedia


```
cx_mat Y = fft2( X )  
cx_mat Y = fft2( X, n_rows, n_cols )
```

```
cx_mat Z = ifft2( cx_mat Y )  
cx_mat Z = ifft2( cx_mat Y, n_rows, n_cols )
```

- *fft2()*: 2D fast Fourier transform of a matrix (real or complex)
- *ifft2()*: 2D inverse fast Fourier transform of a matrix (complex only)
- The optional arguments *n_rows* and *n_cols* specify the size of the transform; a truncated and/or zero-padded version of the input matrix is used
- **Caveat:** the transform is fastest when both *n_rows* and *n_cols* are a power of 2, eg. 64, 128, 256, 512, 1024, ...
- The implementation of the 2D transform in this version is preliminary
- Examples:

```
mat A(100, 100, fill::randu);  
  
cx_mat B = fft2(A);  
cx_mat C = fft2(A, 128, 128);
```

- See also:
 - [fft\(\)](#)
 - [conv2\(\)](#)
 - [real\(\)](#)
 - [fast Fourier transform in MathWorld](#)
 - [fast Fourier transform in Wikipedia](#)

interp1(X, Y, XI, YI)
interp1(X, Y, XI, YI, method)
interp1(X, Y, XI, YI, method, extrapolation_value)

- 1D data interpolation
- Given a 1D function specified in vectors *X* and *Y* (where *X* specifies locations and *Y* specifies the corresponding values), generate vector *YI* which contains interpolated values at locations *XI*
- The *method* argument is optional; it is one of:
 - "nearest" = interpolate using single nearest neighbour
 - "linear" = linear interpolation between two nearest neighbours (**default setting**)
 - "*nearest" = as per "nearest", but faster by assuming that *X* and *XI* are monotonically increasing
 - "*linear" = as per "linear", but faster by assuming that *X* and *XI* are monotonically increasing
- If a location in *XI* is outside the domain of *X*, the corresponding value in *YI* is set to *extrapolation_value*
- The *extrapolation_value* argument is optional; by default it is `datum::nan` (not-a-number)
- Examples:

```
vec x = linspace<vec>(0, 3, 20);  
vec y = square(x);  
  
vec xx = linspace<vec>(0, 3, 100);  
  
vec yy;  
  
interp1(x, y, xx, yy); // use linear interpolation by default  
  
interp1(x, y, xx, yy, "*linear"); // faster than "linear"  
  
interp1(x, y, xx, yy, "nearest");
```

- See also:
 - `interp2()`
 - `polyval()`

- `linspace()`
- `regspace()`
- `conv()`
- interpolation in Wikipedia

interp2(X, Y, Z, XI, YI, ZI)

interp2(X, Y, Z, XI, YI, ZI, method)

interp2(X, Y, Z, XI, YI, ZI, method, extrapolation_value)

- 2D data interpolation
- Given a 2D function specified by matrix Z with coordinates given by vectors X and Y , generate matrix ZI which contains interpolated values at the coordinates given by vectors XI and YI
- The vector pairs (X, Y) and (XI, YI) define 2D coordinates in a grid; for example, X defines the horizontal coordinates and Y defines the corresponding vertical coordinates, so that $(X(m), Y(n))$ is the 2D coordinate of element $Z(n,m)$
- The length of vector X must be equal to the number of columns in matrix Z
- The length of vector Y must be equal to the number of rows in matrix Z
- Vectors X, Y, XI, YI must contain monotonically increasing values (eg. 0.1, 0.2, 0.3, ...)
- The *method* argument is optional; it is one of:
 - "nearest" = interpolate using nearest neighbours
 - "linear" = linear interpolation between nearest neighbours (**default setting**)
- If a coordinate in the 2D grid specified by (XI, YI) is outside the domain of the 2D grid specified by (X, Y) , the corresponding value in ZI is set to *extrapolation_value*
- The *extrapolation_value* argument is optional; by default it is `datum::nan` (not-a-number)
- Examples:

```
mat Z;
```

```
Z.load("input_image.pgm", pgm_binary); // load an image in pgm format
```

```
vec X = regspace(1, Z.n_cols); // X = horizontal spacing
```

```
vec Y = regspace(1, Z.n_rows); // Y = vertical spacing
```

```
vec XI = regspace(X.min(), 1.0/2.0, X.max()); // magnify by approx 2
vec YI = regspace(Y.min(), 1.0/3.0, Y.max()); // magnify by approx 3

mat ZI;

interp2(X, Y, Z, XI, YI, ZI); // use linear interpolation by default

ZI.save("output_image.pgm", pgm_binary);
```

- See also:

- [interp1\(\)](#)
- [regspace\(\)](#)
- [conv2\(\)](#)
- [reshape\(\)](#)
- [bilinear interpolation in Wikipedia](#)

P = polyfit(X, Y, N)
polyfit(P, X, Y, N)

- Given a 1D function specified in vectors X and Y (where X holds independent values and Y holds the corresponding dependent values), model the function as a polynomial of order N and store the polynomial coefficients in column vector P
- The given function is modelled as:
$$y = p_0x^N + p_1x^{N-1} + p_2x^{N-2} + \dots + p_{N-1}x^1 + p_N$$
where p_i is the i -th polynomial coefficient; the coefficients are selected to minimise the overall error of the fit (least squares)
- The column vector P has $N+1$ coefficients
- N must be smaller than the number of elements in X
- If the polynomial coefficients cannot be found:
 - $P = \text{polyfit}(X, Y, N)$ resets P and throws a *std::runtime_error* exception
 - $\text{polyfit}(P, X, Y, N)$ resets P and returns a bool set to *false* (exception is not thrown)

- Examples:

```
vec x = linspace<vec>(0,4*datum::pi,100);  
vec y = cos(x);  
  
vec p = polyfit(x,y,10);
```

- See also:

- [polyval\(\)](#)
- [roots\(\)](#)
- [interp1\(\)](#)
- [polynomial in Wikipedia](#)
- [least squares in Wikipedia](#)
- [curve fitting in Wikipedia](#)
- [least squares fitting in MathWorld](#)

Y = polyval(P, X)

- Given vector P of polynomial coefficients and vector X containing the independent values of a 1D function, generate vector Y which contains the corresponding dependent values

- For each x value in vector X , the corresponding y value in vector Y is generated using:

$$y = p_0x^N + p_1x^{N-1} + p_2x^{N-2} + \dots + p_{N-1}x^1 + p_N$$

where p_i is the i -th polynomial coefficient in vector P

- P must contain polynomial coefficients in descending powers (eg. generated by the [polyfit\(\)](#) function)

- Examples:

```
vec x1 = linspace<vec>(0,4*datum::pi,100);  
vec y1 = cos(x1);  
vec p1 = polyfit(x1,y1,10);  
  
vec y2 = polyval(p1,x1);
```

- See also:

- [polyfit\(\)](#)
- [roots\(\)](#)
- [interp1\(\)](#)
- [polynomial in Wikipedia](#)

Statistics & Clustering

mean, median, stddev, var, range

mean(V)	}	mean (average value)
mean(M)		
mean(M, dim)		
mean(Q)		
mean(Q, dim)		
median(V)	}	median
median(M)		
median(M, dim)		
stddev(V)	}	standard deviation
stddev(V, norm_type)		
stddev(M)		
stddev(M, norm_type)		
stddev(M, norm_type, dim)		
var(V)	}	variance
var(V, norm_type)		
var(M)		
var(M, norm_type)		
var(M, norm_type, dim)		
range(V)	}	range (difference between max and min)
range(M)		
range(M, dim)		

- For vector V , return the statistic calculated using all the elements of the vector
- For matrix M , find the statistic for each column ($dim = 0$), or each row ($dim = 1$)
- For cube Q , find the statistics of elements along dimension dim , where $dim \in \{ 0, 1, 2 \}$
- The dim argument is optional; by default $dim = 0$ is used

- The *norm_type* argument is optional; by default *norm_type* = 0 is used
- For the *var()* and *stddev()* functions:
 - the default *norm_type* = 0 performs normalisation using $N-1$ (where N is the number of samples), providing the best unbiased estimator
 - using *norm_type* = 1 performs normalisation using N , which provides the second moment around the mean
- **Caveat:** to obtain statistics for integer matrices/vectors (eg. *umat*, *imat*, *uvec*, *ivec*), convert to a matrix/vector with floating point values (eg. *mat*, *vec*) using the *conv_to()* function

- Examples:

```
mat A(5, 5, fill::randu);

mat B    = mean(A);
mat C    = var(A);
double m = mean(mean(A));

vec v(5, fill::randu);
double x = var(v);
```

- See also:

- *cov()*
- *cor()*
- *diff()*
- *hist()*
- *histc()*
- *quantile()*
- *min()* & *max()*
- *running_stat* - class for running statistics of scalars
- *running_stat_vec* - class for running statistics of vectors
- *gmm_diag* / *gmm_full* - model and evaluate data using Gaussian Mixture Models (GMMs)
- *kmeans()*

cov(X, Y)
cov(X, Y, norm_type)

cov(X)
cov(X, norm_type)

- For two matrix arguments X and Y , if each row of X and Y is an observation and each column is a variable, the (i,j) -th entry of $\text{cov}(X,Y)$ is the covariance between the i -th variable in X and the j -th variable in Y
- For vector arguments, the type of vector is ignored and each element in the vector is treated as an observation
- For matrices, X and Y must have the same dimensions
- For vectors, X and Y must have the same number of elements
- $\text{cov}(X)$ is equivalent to $\text{cov}(X, X)$
- The *norm_type* argument is optional; by default *norm_type* = 0 is used
- the *norm_type* argument controls the type of normalisation used, with N denoting the number of observations:
 - for *norm_type* = 0, normalisation is done using $N-1$, providing the best unbiased estimation of the covariance matrix (if the observations are from a normal distribution)
 - for *norm_type* = 1, normalisation is done using N , which provides the second moment matrix of the observations about their mean
- Examples:

```
mat X(4, 5, fill::randu);  
mat Y(4, 5, fill::randu);  
  
mat C = cov(X,Y);  
mat D = cov(X,Y, 1);
```
- See also:
 - `cor()`

- statistics functions
- running_stat_vec
- Covariance in MathWorld

cor(X, Y)
cor(X, Y, norm_type)

cor(X)
cor(X, norm_type)

- For two matrix arguments X and Y , if each row of X and Y is an observation and each column is a variable, the (i,j) -th entry of $cor(X,Y)$ is the correlation coefficient between the i -th variable in X and the j -th variable in Y
- For vector arguments, the type of vector is ignored and each element in the vector is treated as an observation
- For matrices, X and Y must have the same dimensions
- For vectors, X and Y must have the same number of elements
- $cor(X)$ is equivalent to $cor(X, X)$
- The *norm_type* argument is optional; by default *norm_type* = 0 is used
- the *norm_type* argument controls the type of normalisation used, with N denoting the number of observations:
 - for *norm_type* = 0, normalisation is done using $N-1$
 - for *norm_type* = 1, normalisation is done using N

- Examples:

```
mat X(4, 5, fill::randu);  
mat Y(4, 5, fill::randu);
```

```
mat R = cor(X,Y);  
mat S = cor(X,Y, 1);
```

- See also:
 - `cov()`
 - `conv()`

- statistics functions
- running_stat_vec
- Correlation in MathWorld
- Autocorrelation in MathWorld

hist(V)
hist(V, n_bins)
hist(V, centers)

hist(X, centers)
hist(X, centers, dim)

- For vector *V*, produce an unsigned vector of the same orientation as *V* (ie. either *uvec* or *urowvec*) that represents a histogram of counts
- For matrix *X*, produce a *umat* matrix containing either column histogram counts (for *dim* = 0, default), or row histogram counts (for *dim* = 1)
- The bin centers can be automatically determined from the data, with the number of bins specified via *n_bins* (default is 10); the range of the bins is determined by the range of the data
- The bin centers can also be explicitly specified via the *centers* vector; the vector must contain monotonically increasing values (eg. 0.1, 0.2, 0.3, ...)

- Examples:

```
vec v(1000, fill::randn); // Gaussian distribution
```

```
uvec h1 = hist(v, 11);
```

```
uvec h2 = hist(v, linspace<vec>(-2,2,11));
```

- See also:
 - *histc()*
 - *quantile()*
 - *statistics functions*
 - *conv_to()*

histc(V, edges)
histc(X, edges)
histc(X, edges, dim)

- For vector *V*, produce an unsigned vector of the same orientation as *V* (ie. either *uvec* or *urowvec*) that contains the counts of the number of values that fall between the elements in the *edges* vector
- For matrix *X*, produce a *umat* matrix containing either column histogram counts (for *dim* = 0, default), or row histogram counts (for *dim* = 1)
- The *edges* vector must contain monotonically increasing values (eg. 0.1, 0.2, 0.3, ...)

- Examples:

```
vec v(1000, fill::randn); // Gaussian distribution
uvec h = histc(v, linspace<vec>(-2,2,11));
```

- See also:
 - *hist()*
 - *statistics functions*
 - *conv_to()*

quantile(V, P)
quantile(X, P)
quantile(X, P, dim)

- For a dataset stored in vector V or matrix X , calculate the quantiles corresponding to the cumulative probability values in the given P vector
- For vector V , produce a vector with the same orientation as V and the same length as P
- For matrix X , produce a matrix with the quantiles for each column vector ($dim = 0$) or each row vector ($dim = 1$)
- The dim argument is optional; by default $dim = 0$
- The P vector must contain values in the $[0,1]$ interval (eg. 0.00, 0.25, 0.50, 0.75, 1.00)
- The algorithm for calculating the quantiles is based on *Definition 5* in:
Rob J. Hyndman and Yanan Fan. Sample Quantiles in Statistical Packages. The American Statistician, 50(4), 361-365, 1996. DOI: [10.2307/2684934](https://doi.org/10.2307/2684934)

- Examples:

```
vec V(1000, fill::randn); // Gaussian distribution

vec P = { 0.25, 0.50, 0.75 };

vec Q = quantile(V, P);
```

- See also:
 - [hist\(\)](#)
 - [median\(\)](#)
 - [normcdf\(\)](#)
 - [quantile in Wikipedia](#)

```
mat coeff = princomp( mat X )  
cx_mat coeff = princomp( cx_mat X )
```

```
princomp( mat coeff, mat X )  
princomp( cx_mat coeff, cx_mat X )
```

```
princomp( mat coeff, mat score, mat X )  
princomp( cx_mat coeff, cx_mat score, cx_mat X )
```

```
princomp( mat coeff, mat score, vec latent, mat X )  
princomp( cx_mat coeff, cx_mat score, vec latent, cx_mat X )
```

```
princomp( mat coeff, mat score, vec latent, vec tsquared, mat X )  
princomp( cx_mat coeff, cx_mat score, vec latent, cx_vec tsquared, cx_mat X )
```

- Principal component analysis of matrix X
- Each row of X is an observation and each column is a variable
- output objects:
 - *coeff*: principal component coefficients
 - *score*: projected data
 - *latent*: eigenvalues of the covariance matrix of X
 - *tsquared*: Hotelling's statistic for each sample
- The computation is based on singular value decomposition
- If the decomposition fails:
 - *coeff = princomp(X)* resets *coeff* and throws a *std::runtime_error* exception
 - remaining forms of *princomp()* reset all output matrices and return a bool set to *false* (exception is not thrown)

- Examples:

```
mat A(5, 4, fill::randu);  
  
mat coeff;
```

```
mat score;  
vec latent;  
vec tsquared;  
  
princomp(coeff, score, latent, tsquared, A);
```

- See also:
 - [eig_sym\(\)](#)
 - [svd\(\)](#)
 - [svd_econ\(\)](#)
 - [principal components analysis in Wikipedia](#)

normpdf(X)

normpdf(X, M, S)

- For each scalar x in X , compute its probability density function according to a Gaussian (normal) distribution using the corresponding mean value m in M and the corresponding standard deviation value s in S :

$$y = \frac{1}{s \sqrt{2\pi}} \exp \left[-0.5 \frac{(x-m)^2}{s^2} \right]$$

- X can be a scalar, vector, or matrix
- M and S can jointly be either scalars, vectors, or matrices
- If M and S are omitted, their values are assumed to be 0 and 1, respectively
- **Caveat:** to reduce the incidence of numerical underflows, consider using [log_normpdf\(\)](#)
- Examples:

```
vec X(10, fill::randu);
vec M(10, fill::randu);
vec S(10, fill::randu);

vec P1 = normpdf(X);
vec P2 = normpdf(X, M, S);
vec P3 = normpdf(1.23, M, S);
vec P4 = normpdf(X, 4.56, 7.89);
double P5 = normpdf(1.23, 4.56, 7.89);
```

- See also:
 - [log_normpdf\(\)](#)
 - [normcdf\(\)](#)
 - [randn\(\)](#)
 - [gmm_diag](#) / [gmm_full](#) - model and evaluate data using Gaussian Mixture Models (GMMs)
 - [normal distribution in Wikipedia](#)

log_normpdf(X)

log_normpdf(X, M, S)

- For each scalar x in X , compute the logarithm version of probability density function according to a Gaussian (normal) distribution using the corresponding mean value m in M and the corresponding standard deviation value s in S :

$$y = \log \left[\frac{1}{s \sqrt{2\pi}} \exp \left[-0.5 \frac{(x-m)^2}{s^2} \right] \right]$$
$$= -\log(s \sqrt{2\pi}) + \left[-0.5 \frac{(x-m)^2}{s^2} \right]$$

- X can be a scalar, vector, or matrix
- M and S can jointly be either scalars, vectors, or matrices
- If M and S are omitted, their values are assumed to be 0 and 1, respectively
- Examples:

```
vec X(10, fill::randu);  
vec M(10, fill::randu);  
vec S(10, fill::randu);  
  
vec P1 = log_normpdf(X);  
vec P2 = log_normpdf(X, M, S);  
vec P3 = log_normpdf(1.23, M, S);  
vec P4 = log_normpdf(X, 4.56, 7.89);  
double P5 = log_normpdf(1.23, 4.56, 7.89);
```

- See also:
 - [normpdf\(\)](#)
 - [gmm_diag / gmm_full](#) - model and evaluate data using Gaussian Mixture Models (GMMs)
 - [normal distribution in Wikipedia](#)

normcdf(X)
normcdf(X, M, S)

- For each scalar x in X , compute its cumulative distribution function according to a Gaussian (normal) distribution using the corresponding mean value m in M and the corresponding standard deviation value s in S
- X can be a scalar, vector, or matrix
- M and S can jointly be either scalars, vectors, or matrices
- If M and S are omitted, their values are assumed to be 0 and 1, respectively

- Examples:

```
vec X(10, fill::randu);  
vec M(10, fill::randu);  
vec S(10, fill::randu);  
  
vec P1 = normcdf(X);  
vec P2 = normcdf(X, M, S);  
vec P3 = normcdf(1.23, M, S);  
vec P4 = normcdf(X, 4.56, 7.89);  
double P5 = normcdf(1.23, 4.56, 7.89);
```

- See also:
 - [normpdf\(\)](#)
 - [quantile\(\)](#)
 - [randn\(\)](#)
 - [normal distribution in Wikipedia](#)
 - [cumulative distribution function in Wikipedia](#)

`X = mvnrnd(M, C)`
`X = mvnrnd(M, C, N)`

`mvnrnd(X, M, C)`
`mvnrnd(X, M, C, N)`

- Generate a matrix with random column vectors from a multivariate Gaussian (normal) distribution with parameters M and C :
 - M is the mean; must be a column vector
 - C is the covariance matrix; must be symmetric positive semi-definite (preferably positive definite)
- N is the number of column vectors to generate; if N is omitted, it is assumed to be 1
- **Caveat:** repeated generation of one vector (or a small number of vectors) using the same M and C parameters can be inefficient;
for repeated generation consider using the `generate()` function in the `gmm_diag` and `gmm_full` classes
- If generating the random vectors fails:
 - `X = mvnrnd(M, C)` and `X = mvnrnd(M, C, N)` reset X and throw a `std::runtime_error` exception
 - `mvnrnd(X, M, C)` and `mvnrnd(X, M, C, N)` reset X and return a bool set to `false` (exception is not thrown)

- Examples:

```
vec M(5, fill::randu);

mat B(5, 5, fill::randu);
mat C = B.t() * B;

mat X = mvnrnd(M, C, 100);
```

- See also:
 - `randn()`
 - `chi2rnd()`
 - `wishrnd()`
 - `cov()`
 - `.is_sympd()`
 - `gmm_diag` / `gmm_full` - model and evaluate data using Gaussian Mixture Models (GMMs)

- [multivariate normal distribution in Wikipedia](#)

chi2rnd(DF)
chi2rnd(DF_scalar)
chi2rnd(DF_scalar, n_elem)
chi2rnd(DF_scalar, n_rows, n_cols)
chi2rnd(DF_scalar, size(X))

- Generate a random scalar, vector or matrix with elements sampled from a chi-squared distribution with the degrees of freedom specified by parameter *DF* or *DF_scalar*
- *DF* is a vector or matrix, while *DF_scalar* is a scalar
- Each value in *DF* and *DF_scalar* must be greater than zero
- For the *chi2rnd(DF)* form, the output vector/matrix has the same size and type as *DF*; each element in *DF* specifies a separate degree of freedom
- Usage:
 - *vector_type* v = chi2rnd(DF), where the type of DF is a real *vector_type*
 - *matrix_type* X = chi2rnd(DF), where the type of DF is a real *matrix_type*
 - *scalar_type* s = chi2rnd<*scalar_type*>(DF_scalar), where *scalar_type* is either *float* or *double*
 - *vector_type* v = chi2rnd<*vector_type*>(DF_scalar, n_elem)
 - *matrix_type* X = chi2rnd<*matrix_type*>(DF_scalar, n_rows, n_cols)
 - *matrix_type* Y = chi2rnd<*matrix_type*>(DF_scalar, size(X))

- Examples:

```
mat X = chi2rnd(2, 5, 6);  
  
mat A = randi<mat>(5, 6, distr_param(1, 10));  
mat B = chi2rnd(A);
```

- See also:

- [randg\(\)](#)
- [randn\(\)](#)
- [mvnrnd\(\)](#)
- [wishrnd\(\)](#)

- `size()`
- chi-squared distribution in Wikipedia

W = wishrnd(S, df)
W = wishrnd(S, df, D)

wishrnd(W, S, df)
wishrnd(W, S, df, D)

- Generate a random matrix sampled from the Wishart distribution with parameters S and df :
 - S is a symmetric positive definite matrix (eg. a covariance matrix)
 - df is a scalar specifying the degrees of freedom; it can be a non-integer value
- D is an optional argument; it specifies the Cholesky decomposition of S ; if D is provided, S is ignored; using D is more efficient if `wishrnd()` needs to be used many times for the same S matrix
- If generating the random matrix fails:
 - `W = wishrnd(S, df)` and `W = wishrnd(S, df, D)` reset W and throw a `std::runtime_error` exception
 - `wishrnd(W, S, df)` and `wishrnd(W, S, df, D)` reset W and return a bool set to `false` (exception is not thrown)
- **Caveat:** for the inverse (inverted) Wishart distribution, use `iwishrnd()`

- Examples:

```
mat X(5, 5, fill::randu);  
  
mat S = X.t() * X;  
  
mat W = wishrnd(S, 6.7);
```

- See also:
 - `iwishrnd()`
 - `chi2rnd()`
 - `mvnrnd()`
 - `chol()`
 - `cov()`
 - `randn()`
 - `.is_sympd()`
 - [Wishart distribution in Wikipedia](#)

W = iwishrnd(T, df)
W = iwishrnd(T, df, Dinv)

iwishrnd(W, T, df)
iwishrnd(W, T, df, Dinv)

- Generate a random matrix sampled from the **inverse** Wishart distribution with parameters T and df :
 - T is a symmetric positive definite matrix
 - df is a scalar specifying the degrees of freedom; it can be a non-integer value
- $Dinv$ is an optional argument; it specifies the Cholesky decomposition of the inverse of T ; if $Dinv$ is provided, T is ignored
using $Dinv$ is more efficient if $iwishrnd()$ needs to be used many times for the same T matrix
- If generating the random matrix fails:
 - $W = iwishrnd(T, df)$ and $W = iwishrnd(T, df, Dinv)$ reset W and throw a `std::runtime_error` exception
 - $iwishrnd(W, T, df)$ and $iwishrnd(W, T, df, Dinv)$ reset W and return a bool set to `false` (exception is not thrown)

- **Caveat:** for the plain (non-inverse) Wishart distribution, use `wishrnd()`

- Examples:

```
mat X(5, 5, fill::randu);  
  
mat T = X.t() * X;  
  
mat W = iwishrnd(T, 6.7);
```

- See also:
 - `wishrnd()`
 - `chol()`
 - `inv_sympd()`
 - `.is_sympd()`
 - [inverse Wishart distribution in Wikipedia](#)

running_stat<type>

- Class for running statistics (online statistics) of scalars (one dimensional process/signal)
- Useful if the storage of all samples (scalars) is impractical, or if the number of samples is not known in advance
- *type* is one of: *float*, *double*, *cx_float*, *cx_double*
- For an instance of *running_stat* named as *X*, the member functions are:

X(scalar)	update the statistics using the given scalar
X.min()	current minimum value
X.max()	current maximum value
X.range()	current range
X.mean()	current mean or average value
X.var() and X.var(norm_type)	current variance
X.stddev() and X.stddev(norm_type)	current standard deviation
X.reset()	reset all statistics and set the number of samples to zero
X.count()	current number of samples

- The *norm_type* argument is optional; by default *norm_type = 0* is used
- For the *.var()* and *.stddev()* functions, the default *norm_type = 0* performs normalisation using *N-1* (where *N* is the number of samples so far), providing the best unbiased estimator; using *norm_type = 1* causes normalisation to be done using *N*, which provides the second moment around the mean
- The return type of *.count()* depends on the underlying form of *type*: it is either *float* or *double*
- Examples:

```
running_stat<double> stats;  
  
for(uword i=0; i<10000; ++i)  
{  
    double sample = randn();
```

```
    stats(sample);  
}  
  
cout << "mean = " << stats.mean() << endl;  
cout << "var  = " << stats.var()  << endl;  
cout << "min  = " << stats.min()  << endl;  
cout << "max  = " << stats.max()  << endl;
```

- See also:
 - [running_stat_vec](#) (running statistics of vectors)
 - [statistics functions](#)
 - [gmm_diag / gmm_full](#) - model and evaluate data using Gaussian Mixture Models (GMMs)

running_stat_vec<vec_type>
running_stat_vec<vec_type>(calc_cov)

- Class for running statistics (online statistics) of vectors (multi-dimensional process/signal)
- Useful if the storage of all samples (vectors) is impractical, or if the number of samples is not known in advance
- This class is similar to *running_stat*, with the difference that vectors are processed instead of scalars
- *vec_type* is the vector type of the samples; for example: *vec*, *cx_vec*, *rowvec*, ...
- For an instance of *running_stat_vec* named as *X*, the member functions are:

X(vector)	update the statistics using the given vector
X.min()	vector of current minimum values
X.max()	vector of current maximum values
X.range()	vector of current ranges
X.mean()	vector of current means
X.var() and X.var(norm_type)	vector of current variances
X.stddev() and X.stddev(norm_type)	vector of current standard deviations
X.cov() and X.cov(norm_type)	matrix of current covariances; valid if <i>calc_cov=true</i> during construction of <i>running_stat_vec</i>
X.reset()	reset all statistics and set the number of samples to zero
X.count()	current number of samples

- The *calc_cov* argument is optional; by default *calc_cov=false*, indicating that the covariance matrix will not be calculated; to enable the covariance matrix, use *calc_cov=true* during construction; for example:
running_stat_vec<vec> X(true);
- The *norm_type* argument is optional; by default *norm_type = 0* is used
- For the *.var()* and *.stddev()* functions, the default *norm_type = 0* performs normalisation using *N-1* (where *N*

is the number of samples so far), providing the best unbiased estimator; using *norm_type* = 1 causes normalisation to be done using *N*, which provides the second moment around the mean

- The return type of *.count()* depends on the underlying form of *vec_type*: it is either *float* or *double*
- Examples:

```
running_stat_vec<vec> stats;

vec sample;

for(uword i=0; i<10000; ++i)
{
    sample = randu<vec>(5);
    stats(sample);
}

cout << "mean = " << endl << stats.mean() << endl;
cout << "var  = " << endl << stats.var()  << endl;
cout << "max  = " << endl << stats.max()  << endl;

//
//

running_stat_vec<rowvec> more_stats(true);

for(uword i=0; i<20; ++i)
{
    sample = randu<rowvec>(3);

    sample(1) -= sample(0);
    sample(2) += sample(1);

    more_stats(sample);
}

cout << "covariance matrix = " << endl;
cout << more_stats.cov() << endl;

rowvec sd = more_stats.stddev();

cout << "correlations = " << endl;
cout << more_stats.cov() / (sd.t() * sd);
```

- See also:
 - `running_stat` (running statistics of scalars)
 - `statistics functions`
 - `cov()`
 - `cor()`
 - `gmm_diag` / `gmm_full` - model and evaluate data using Gaussian Mixture Models (GMMs)

kmeans(means, data, k, seed_mode, n_iter, print_mode)

- Cluster given data into k disjoint sets
- The *means* parameter is the output matrix for storing the resulting centroids of the sets, with each centroid stored as a column vector
- The *data* parameter is the input data matrix, with each sample stored as a column vector
- The k parameter indicates the number of centroids; the number of samples in the *data* matrix should be much larger than k
- The *seed_mode* parameter specifies how the initial centroids are seeded; it is one of:
 - `keep_existing` use the centroids specified in the *means* matrix as the starting point
 - `static_subset` use a subset of the data vectors (repeatable)
 - `random_subset` use a subset of the data vectors (random)
 - `static_spread` use a maximally spread subset of data vectors (repeatable)
 - `random_spread` use a maximally spread subset of data vectors (random start)

caveat: seeding the initial centroids with `static_spread` and `random_spread` can be much more time consuming than with `static_subset` and `random_subset`

- The *n_iter* parameter specifies the number of clustering iterations; this is data dependent, but about 10 is typically sufficient
- The *print_mode* parameter is either *true* or *false*, indicating whether progress is printed during clustering
- If the clustering fails, the *means* matrix is reset and a bool set to *false* is returned
- The clustering will run faster on multi-core machines when OpenMP is enabled in your compiler (eg. *-fopenmp* in GCC and clang)
- Examples:

```
uword d = 5;      // dimensionality
uword N = 10000;  // number of vectors
```

```
mat data(d, N, fill::randu);

mat means;

bool status = kmeans(means, data, 2, random_subset, 10, true);

if(status == false)
{
    cout << "clustering failed" << endl;
}

means.print("means:");
```

- See also:
 - [gmm_diag / gmm_full](#) - model and evaluate data using Gaussian Mixture Models (GMMs)
 - [statistics functions](#)
 - [running_stat_vec](#)
 - [k-means clustering in Wikipedia](#)
 - [k-means clustering in MathWorld](#)
 - [OpenMP in Wikipedia](#)

`gmm_diag` **`gmm_full`**

- Classes for multivariate data modelling and evaluation via **Gaussian Mixture Models** (GMMs)
- The *gmm_diag* class is tailored for **diagonal covariance matrices** (ie. in each covariance matrix, all entries outside the main diagonal are assumed to be zero)
- The *gmm_full* class is tailored for **full covariance matrices**
- The *gmm_diag* class is typically much faster to train and use than the *gmm_full* class, at the potential cost of some reduction in modelling accuracy
- The *gmm_diag* and *gmm_full* classes include dedicated optimisation algorithms for learning (training) the model parameters from data:
 - k-means clustering, for quick initial estimates
 - Expectation-Maximisation (EM), for maximum-likelihood estimates

The optimisation algorithms are multi-threaded and can run much quicker on multi-core machines when OpenMP is enabled in your compiler (eg. *-fopenmp* in GCC and clang)

- The classes can also be used for probabilistic clustering and vector quantisation (VQ)
- Data is modelled as:

$$p(x) = \sum_{g=0}^{n_gaus-1} h_g \text{ N}(x | m_g, C_g)$$

where:

- *n_gaus* is the number of Gaussians; $n_gaus \geq 1$
- $\text{N}(x | m_g, C_g)$ represents a Gaussian (normal) distribution
- each Gaussian *g* has the following parameters:
 - h_g is the heft (weight), with constraints $h_g \geq 0$ and $\sum h_g = 1$
 - m_g is the mean vector (centroid) with dimensionality *n_dims*
 - C_g is the covariance matrix (either diagonal or full)

- Mathematical implementation details are available in the following paper:
 Conrad Sanderson and Ryan Curtin.
An Open Source C++ Implementation of Multi-Threaded Gaussian Mixture Models, k-Means and Expectation Maximisation.
 International Conference on Signal Processing and Communication Systems, 2017.

- For an instance of *gmm_diag* or *gmm_full* named as *M*, the member functions and variables are:

M.log_p(V)	return a scalar representing the log-likelihood of vector <i>V</i> (of type <i>vec</i>)
M.log_p(V, g)	return a scalar representing the log-likelihood of vector <i>V</i> (of type <i>vec</i>), according to Gaussian with index <i>g</i>
M.log_p(X)	return a row vector (of type <i>rowvec</i>) containing log-likelihoods of each column vector in matrix <i>X</i> (of type <i>mat</i>)
M.log_p(X, g)	return a row vector (of type <i>rowvec</i>) containing log-likelihoods of each column vector in matrix <i>X</i> (of type <i>mat</i>), according to Gaussian with index <i>g</i>
M.sum_log_p(X)	return a scalar representing the sum of log-likelihoods for all column vectors in matrix <i>X</i> (of type <i>mat</i>)
M.sum_log_p(X, g)	return a scalar representing the sum of log-likelihoods for all column vectors in matrix <i>X</i> (of type <i>mat</i>), according to Gaussian with index <i>g</i>
M.avg_log_p(X)	return a scalar representing the average log-likelihood of all column vectors in matrix <i>X</i> (of type <i>mat</i>)
M.avg_log_p(X, g)	return a scalar representing the average log-likelihood of all column vectors in matrix <i>X</i> (of type <i>mat</i>), according to Gaussian with index <i>g</i>
M.assign(V, dist_mode)	return the index of the closest mean (or Gaussian) to vector <i>V</i> (of type <i>vec</i>);

parameter *dist_mode* is one of:

eucl_dist Euclidean distance (takes only means into account)

prob_dist probabilistic "distance", defined as the inverse likelihood
(takes into account means, covariances and hefts)

M.assign(X, dist_mode)

return a row vector (of type *urowvec*) containing the indices of the closest means (or Gaussians) to each column vector in matrix *X* (of type *mat*);

parameter *dist_mode* is either *eucl_dist* or *prob_dist* (as per the **.assign()** function above)

M.raw_hist(X, dist_mode)

return a row vector (of type *urowvec*) representing the raw histogram of counts; each entry is the number of counts corresponding to a Gaussian; each count is the number times the corresponding Gaussian was the closest to each column vector in matrix *X*; parameter *dist_mode* is either *eucl_dist* or *prob_dist* (as per the **.assign()** function above)

M.norm_hist(X, dist_mode)

similar to the **.raw_hist()** function above; return a row vector (of type *rowvec*) containing normalised counts; the vector sums to one; parameter *dist_mode* is either *eucl_dist* or *prob_dist* (as per the **.assign()** function above)

M.generate()

return a column vector (of type *vec*) representing a random sample generated according to the model's parameters

M.generate(N)

return a matrix (of type *mat*) containing *N* column vectors, with each vector representing a random sample generated according to the model's parameters

M.save(filename)

save the model to a file and return a *bool* indicating either success (*true*) or failure (*false*)

M.load(filename)

load the model from a file and return a *bool* indicating either success (*true*) or failure (*false*)

M.n_gaus()

return the number of means/Gaussians in the model

M.n_dims()	return the dimensionality of the means/Gaussians in the model
M.reset(n_dims, n_gaus)	set the model to have dimensionality <i>n_dims</i> , with <i>n_gaus</i> number of Gaussians; all the means are set to zero, all covariance matrix representations are equivalent to the identity matrix, and all the hefts (weights) are set to be uniform
M.hefts	read-only row vector (of type <i>rowvec</i>) containing the hefts (weights)
M.means	read-only matrix (of type <i>mat</i>) containing the means (centroids), stored as column vectors
M.dcovs [only in <i>gmm_diag</i>]	read-only matrix (of type <i>mat</i>) containing the representation of diagonal covariance matrices, with the set of diagonal covariances for each Gaussian stored as a column vector; applicable only to the <i>gmm_diag</i> class
M.fcovs [only in <i>gmm_full</i>]	read-only <i>cube</i> containing the full covariance matrices, with each covariance matrix stored as a slice within the cube; applicable only to the <i>gmm_full</i> class
M.set_hefts(V)	set the hefts (weights) of the model to be as specified in row vector <i>V</i> (of type <i>rowvec</i>); the number of hefts must match the existing model
M.set_means(X)	set the means to be as specified in matrix <i>X</i> (of type <i>mat</i>); the number of means and their dimensionality must match the existing model
M.set_dcovs(X) [only in <i>gmm_diag</i>]	set the diagonal covariances matrices to be as specified in matrix <i>X</i> (of type <i>mat</i>), with the set of diagonal covariances for each Gaussian stored as a column vector; the number of covariance matrices and their dimensionality must match the existing model; applicable only to the <i>gmm_diag</i> class

M.set_fcovs(X) <i>[only in gmm_full]</i>	set the full covariances matrices to be as specified in cube <i>X</i> , with each covariance matrix stored as a slice within the cube; the number of covariance matrices and their dimensionality must match the existing model; applicable only to the <i>gmm_full</i> class
M.set_params(means, covs, hefts)	set all the parameters at the same time; the type and layout of the parameters is as per the .set_hefts() , .set_means() , .set_dcovs() and .set_fcovs() functions above; the number of Gaussians and dimensionality can be different from the existing model
M.learn(data, n_gaus, dist_mode, seed_mode, km_iter, em_iter, var_floor, print_mode)	learn the model parameters via multi-threaded k-means and/or EM algorithms; return a <code>bool</code> value, with <i>true</i> indicating success, and <i>false</i> indicating failure; the parameters have the following meanings:
<i>data</i>	matrix (of type <i>mat</i>) containing training samples; each sample is stored as a column vector
<i>n_gaus</i>	set the number of Gaussians to <i>n_gaus</i> ; to help convergence, it is recommended that the given <i>data</i> matrix (above) contains at least 10 samples for each Gaussian
<i>dist_mode</i>	specifies the distance used during the seeding of initial means and k-means clustering: <div> <div>eucl_dist</div> <div>Euclidean distance</div> <div>maha_dist</div> <div>Mahalanobis distance, which uses a global diagonal covariance matrix estimated from the training samples; this is recommended for probabilistic applications</div> </div>
<i>seed_mode</i>	specifies how the initial means are seeded prior to running k-means and/or EM algorithms:

<code>keep_existing</code>	keep the existing model (do not modify the means, covariances and hefts)
<code>static_subset</code>	a subset of the training samples (repeatable)
<code>random_subset</code>	a subset of the training samples (random)
<code>static_spread</code>	a maximally spread subset of training samples (repeatable)
<code>random_spread</code>	a maximally spread subset of training samples (random start)

caveat: seeding the initial means with `static_spread` and `random_spread` can be much more time consuming than with `static_subset` and `random_subset`

<code>km_iter</code>	the number of iterations of the k-means algorithm; this is data dependent, but typically 10 iterations are sufficient
<code>em_iter</code>	the number of iterations of the EM algorithm; this is data dependent, but typically 5 to 10 iterations are sufficient
<code>var_floor</code>	the variance floor (smallest allowed value) for the diagonal covariances; setting this to a small non-zero value can help with convergence and/or better quality parameter estimates
<code>print_mode</code>	either <code>true</code> or <code>false</code> ; enable or disable printing of progress during the k-means and EM algorithms

- Examples:

```
// create synthetic data with 2 Gaussians

uword d = 5;      // dimensionality
uword N = 10000;  // number of vectors

mat data(d, N, fill::zeros);
```

```

vec mean0 = linspace<vec>(1,d,d);
vec mean1 = mean0 + 2;

uword i = 0;

while(i < N)
{
    if(i < N) { data.col(i) = mean0 + randn<vec>(d); ++i; }
    if(i < N) { data.col(i) = mean0 + randn<vec>(d); ++i; }
    if(i < N) { data.col(i) = mean1 + randn<vec>(d); ++i; }
}

// model the data as a diagonal GMM with 2 Gaussians

gmm_diag model;

bool status = model.learn(data, 2, maha_dist, random_subset, 10, 5, 1e-10, true);

if(status == false)
{
    cout << "learning failed" << endl;
}

model.means.print("means:");

double scalar_likelihood = model.log_p( data.col(0) );
rowvec set_likelihood = model.log_p( data.cols(0,9) );

double overall_likelihood = model.avg_log_p(data);

uword gaus_id = model.assign( data.col(0), eucl_dist );
urowvec gaus_ids = model.assign( data.cols(0,9), prob_dist );

urowvec hist1 = model.raw_hist (data, prob_dist);
rowvec hist2 = model.norm_hist(data, eucl_dist);

model.save("my_model.gmm");

```

- See also:
 - [normpdf\(\)](#)
 - [mvnrnd\(\)](#)
 - [statistics functions](#)

- `running_stat_vec`
- `kmeans()`
- covariance matrix in Wikipedia
- covariance matrix in MathWorld
- Mahalanobis distance in Wikipedia
- multivariate normal distribution in Wikipedia
- mixture model in Wikipedia
- k-means clustering in Wikipedia
- k-means clustering in MathWorld
- Expectation-Maximisation algorithm in Wikipedia
- maximum likelihood in MathWorld
- vector quantisation in Wikipedia
- OpenMP in Wikipedia

Miscellaneous

constants (pi, inf, eps, ...)

datum::pi	π , the ratio of any circle's circumference to its diameter
datum::tau	τ , the ratio of any circle's circumference to its radius; equivalent to 2π
datum::inf	∞ , infinity
datum::nan	"not a number" (NaN); caveat: NaN is not equal to anything, even itself
datum::eps	machine epsilon; approximately 2.2204e-16; difference between 1 and the next representable value
datum::e	base of the natural logarithm
datum::sqrt2	square root of 2
datum::log_min	log of minimum non-zero value (type and machine dependent)
datum::log_max	log of maximum value (type and machine dependent)
datum::euler	Euler's constant, aka Euler-Mascheroni constant
datum::gratio	golden ratio
datum::m_u	atomic mass constant (in kg)
datum::N_A	Avogadro constant
datum::k	Boltzmann constant (in joules per kelvin)
datum::k_evk	Boltzmann constant (in eV/K)
datum::a_0	Bohr radius (in meters)
datum::mu_B	Bohr magneton
datum::Z_0	characteristic impedance of vacuum (in ohms)
datum::G_0	conductance quantum (in siemens)
datum::k_e	Coulomb's constant (in meters per farad)
datum::eps_0	electric constant (in farads per meter)

<code>datum::m_e</code>	electron mass (in kg)
<code>datum::eV</code>	electron volt (in joules)
<code>datum::ec</code>	elementary charge (in coulombs)
<code>datum::F</code>	Faraday constant (in coulombs)
<code>datum::alpha</code>	fine-structure constant
<code>datum::alpha_inv</code>	inverse fine-structure constant
<code>datum::K_J</code>	Josephson constant
<code>datum::mu_0</code>	magnetic constant (in henries per meter)
<code>datum::phi_0</code>	magnetic flux quantum (in webers)
<code>datum::R</code>	molar gas constant (in joules per mole kelvin)
<code>datum::G</code>	Newtonian constant of gravitation (in newton square meters per kilogram squared)
<code>datum::h</code>	Planck constant (in joule seconds)
<code>datum::h_bar</code>	Planck constant over 2 pi, aka reduced Planck constant (in joule seconds)
<code>datum::m_p</code>	proton mass (in kg)
<code>datum::R_inf</code>	Rydberg constant (in reciprocal meters)
<code>datum::c_0</code>	speed of light in vacuum (in meters per second)
<code>datum::sigma</code>	Stefan-Boltzmann constant
<code>datum::R_k</code>	von Klitzing constant (in ohms)
<code>datum::b</code>	Wien wavelength displacement law constant

- The constants are stored in the *Datum<type>* class, where *type* is either *float* or *double*; for convenience, *Datum<double>* is typedefed as *datum*, and *Datum<float>* is typedefed as *fdatum*
- **Caveat:** `datum::nan` is not equal to anything, even itself; to check whether a scalar *x* is finite, use

`std::isfinite(x)`

- The physical constants were mainly taken from [NIST 2018 CODATA values](#), and some from [WolframAlpha](#) (as of 2009-06-23)

- Examples:

```
cout << "speed of light = " << datum::c_0 << endl;
```

```
cout << "log_max for floats = ";  
cout << fdatum::log_max << endl;
```

```
cout << "log_max for doubles = ";  
cout << datum::log_max << endl;
```

- See also:

- [.is_finite\(\)](#)
- [.fill\(\)](#)
- [NaN](#) in Wikipedia
- [physical constant](#) in Wikipedia
- [replacement of \$2\pi\$ with \$\tau\$](#) in Wikipedia
- [The Tau Manifesto](#) by Michael Hartl
- [std::numeric_limits](#) in [cplusplus.com](#)
- [std::numeric_limits](#) in [cppreference.com](#)

wall_clock

- Simple timer class for measuring the number of elapsed seconds
- An instance of the class has two member functions:

`.tic()` start the timer

`.toc()` return the number of seconds since the last call to `.tic()`

- Examples:

```
wall_clock timer;  
  
timer.tic();  
  
// ... do something ...  
  
double n = timer.toc();  
  
cout << "number of seconds: " << n << endl;
```

- See also:
 - [elapsed real time](#) in Wikipedia

RNG seed setting

- There are two functions to change the seed used by the random number generator (RNG):

`arma_rng::set_seed(value)` set the RNG seed to the specified *value*

`arma_rng::set_seed_random()` set the RNG seed to a value drawn from `std::random_device` (if the reported entropy is not zero), or `/dev/urandom` (on Linux and macOS), or based on the current time (on systems without `/dev/urandom`)

- Since Armadillo 12.6.2, the default RNG on all systems is the 64-bit version of the **Mersenne Twister** (aka MT19937-64)
- **Caveat:** when using a multi-threading framework (such as OpenMP) and the underlying system supports the `thread_local` storage specifier, the above functions change the seed only within the thread they are running on
- To change the seeds on all OpenMP threads to the same value, adapt the following code:

```
#pragma omp parallel
{
    arma_rng::set_seed(123);
}
```

- To change the seeds on all OpenMP threads to unique values, adapt the following code:

```
std::atomic<std::size_t> counter(0);

#pragma omp parallel
{
    arma_rng::set_seed(123 + counter++);
}
```

- See also:
 - `randu()`, `randn()`, `randg()`, `randi()`, `sprandu()`, `sprandn()`, `shuffle()`, `randperm()`
 - **Pseudo-random number generator** in Wikipedia
 - **Mersenne Twister** in Wikipedia
 - `/dev/random` in Wikipedia

- `std::mersenne_twister_engine` in `cppreference.com`
- `std::random_device` in `cppreference.com`

output streams

- The default stream for printing matrices and cubes is `std::cout`
the stream can be changed via the `ARMA_COUT_STREAM` define; see [config.hpp](#)
- The default stream for printing warnings and errors is `std::cerr`
the stream can be changed via the `ARMA_CERR_STREAM` define; see [config.hpp](#)
- The degree of printed warnings is controlled by the `ARMA_WARN_LEVEL` define; see [config.hpp](#)
- Example of changing the warning level:

```
#define ARMA_WARN_LEVEL 1  
#include <armadillo>
```

- See also:
 - [config.hpp](#)
 - [.print\(\)](#)
 - [std::cout](#)
 - [std::cerr](#)
 - [std::ostream](#)

uword, sword

- *uword* is a typedef for an unsigned integer type; it is used for matrix **indices** as well as all internal counters and loops
- *sword* is a typedef for a signed integer type
- The minimum width of both *uword* and *sword* is either 32 or 64 bits:
 - the default width is 32 bits on 32-bit platforms
 - the default width is 64 bits on 64-bit platforms
 - the default width is 32 bits when using Armadillo in the R environment (via RcppArmadillo) on either 32-bit or 64-bit platforms
- The width can also be forcefully set to 64 bits by enabling **ARMA_64BIT_WORD** via editing *include/armadillo_bits/config.hpp*
- See also:
 - **C++ variable types**
 - **explanation of typedef**
 - **imat & umat** matrix types
 - **ivec & uvec** vector types

cx_double, cx_float

- Convenience short forms (typedefs) for complex element types:

`cx_double` equivalent to `std::complex<double>`

`cx_float` equivalent to `std::complex<float>`

- Example:

```
cx_mat X(5, 5, fill::randu);
```

```
X(1,2) = cx_double(3.4, 5.6);
```

```
cx_double val = X(2,3);
```

- See also:
 - [complex numbers in the standard C++ library](#)
 - [explanation of *typedef*](#)
 - `cx_mat` matrix type
 - `cx_vec` vector type

Examples of Matlab/Octave syntax and conceptually corresponding Armadillo syntax

Matlab/Octave	Armadillo	Notes
A(1, 1)	A(0, 0)	indexing in Armadillo starts at 0
A(k, k)	A(k-1, k-1)	
size(A,1)	A.n_rows	read only
size(A,2)	A.n_cols	
size(Q,3)	Q.n_slices	Q is a cube (3D array)
numel(A)	A.n_elem	
A(:, k)	A.col(k)	this is a conceptual example only; exact conversion from Matlab/Octave to Armadillo syntax will require taking into account that indexing starts at 0
A(k, :)	A.row(k)	
A(:, p:q)	A.cols(p, q)	
A(p:q, :)	A.rows(p, q)	
A(p:q, r:s)	A(span(p,q), span(r,s))	A(span(first_row, last_row), span(first_col, last_col))
Q(:, :, k)	Q.slice(k)	Q is a cube (3D array)
Q(:, :, t:u)	Q.slices(t, u)	
Q(p:q, r:s, t:u)	Q(span(p,q), span(r,s), span(t,u))	
A'	A.t() or trans(A)	matrix transpose / Hermitian transpose (for complex matrices, the conjugate of each element is taken)
A = zeros(size(A))	A.zeros()	
A = ones(size(A))	A.ones()	

<code>A = zeros(k)</code>	<code>A = zeros<mat>(k,k)</code>
<code>A = ones(k)</code>	<code>A = ones<mat>(k,k)</code>

<code>C = complex(A,B)</code>	<code>cx_mat C = cx_mat(A,B)</code>
-------------------------------	-------------------------------------

<code>A .* B</code>	<code>A % B</code>
<code>A ./ B</code>	<code>A / B</code>
<code>A \ B</code>	<code>solve(A,B)</code>
<code>A = A + 1;</code>	<code>A++</code>
<code>A = A - 1;</code>	<code>A--</code>

element-wise multiplication

element-wise division

conceptually similar to `inv(A)*B`, but more efficient

<code>A = [1 2; 3 4;]</code>	<code>A = { { 1, 2 }, { 3, 4 } };</code>
--------------------------------	--

element initialisation

<code>X = A(:)</code>	<code>X = vectorise(A)</code>
<code>X = [A B]</code>	<code>X = join_horiz(A,B)</code>
<code>X = [A; B]</code>	<code>X = join_vert(A,B)</code>

<code>A</code>	<code>cout << A << endl;</code>
	or
	<code>A.print("A =");</code>

<code>save -ascii 'A.txt' A</code>	<code>A.save("A.txt", raw_ascii);</code>
------------------------------------	--

Matlab/Octave matrices saved as ascii are readable by Armadillo (and vice-versa)

<code>load -ascii 'A.txt'</code>	<code>A.load("A.txt", raw_ascii);</code>
----------------------------------	--

<code>A = randn(2,3);</code>	<code>mat A = randn(2,3);</code>
<code>B = randn(4,5);</code>	<code>mat B = randn(4,5);</code>
<code>F = { A; B }</code>	<code>field<mat> F(2,1);</code>
	<code>F(0,0) = A;</code>
	<code>F(1,0) = B;</code>

fields store arbitrary objects, such as matrices

example program

```
#include <iostream>
#include <armadillo>

using namespace std;
using namespace arma;

int main()
{
    mat A(4, 5, fill::randu);
    mat B(4, 5, fill::randu);

    cout << A*B.t() << endl;

    return 0;
}
```

- If the above program is stored as *example.cpp*, under Linux and macOS it can be compiled using:
g++ example.cpp -o example -std=c++11 -O2 -larmadillo
- Armadillo extensively uses template meta-programming, so it's recommended to enable optimisation when compiling programs (eg. use the -O2 or -O3 options for GCC or clang)
- See the [Questions](#) page for more info on compiling and linking
- See also the example program that comes with the Armadillo archive

config.hpp

- Armadillo can be configured via editing the file *include/armadillo_bits/config.hpp*
- Specific functionality can be enabled or disabled by uncommenting or commenting out a particular *#define*, listed below.
- Some options can also be specified by explicitly defining them **before** including the armadillo header.

ARMA_DONT_USE_WRAPPER	Disable going through the run-time Armadillo wrapper library (<i>libarmadillo.so</i>) when calling LAPACK, BLAS, ARPACK, SuperLU and HDF5 functions. You will need to directly link with BLAS, LAPACK, etc (eg. <i>-lblas -llapack</i>)
ARMA_USE_LAPACK	Enable use of LAPACK, or a high-speed replacement for LAPACK (eg. OpenBLAS, Intel MKL, or the Accelerate framework). Armadillo requires LAPACK for functions such as <i>svd()</i> , <i>inv()</i> , <i>eig_sym()</i> , <i>solve()</i> , etc.
ARMA_DONT_USE_LAPACK	Disable use of LAPACK; overrides <i>ARMA_USE_LAPACK</i>
ARMA_USE_BLAS	Enable use of BLAS, or a high-speed replacement for BLAS (eg. OpenBLAS, Intel MKL, or the Accelerate framework). BLAS is used for <i>matrix multiplication</i> . Without BLAS, Armadillo will use a built-in matrix multiplication routine, which might be slower for large matrices.
ARMA_DONT_USE_BLAS	Disable use of BLAS; overrides <i>ARMA_USE_BLAS</i>
ARMA_USE_NEWARP	Enable use of NEWARP (built-in alternative to ARPACK). This is used for the eigen decomposition of real (non-complex) sparse matrices, ie. <i>eigs_gen()</i> , <i>eigs_sym()</i> and <i>svds()</i> . Requires <i>ARMA_USE_LAPACK</i> to be enabled. If use of both NEWARP and ARPACK is enabled, NEWARP will be preferred.

ARMA_DONT_USE_NEWARP	Disable use of NEWARP (built-in alternative to ARPACK); overrides <i>ARMA_USE_NEWARP</i>
ARMA_USE_ARPACK	Enable use of ARPACK, or a high-speed replacement for ARPACK. Armadillo requires ARPACK for the eigen decomposition of complex sparse matrices, ie. <i>eigs_gen()</i> , <i>eigs_sym()</i> and <i>svds()</i> . If use of NEWARP is disabled, ARPACK will also be used for the eigen decomposition of real sparse matrices.
ARMA_DONT_USE_ARPACK	Disable use of ARPACK; overrides <i>ARMA_USE_ARPACK</i>
ARMA_USE_SUPERLU	Enable use of SuperLU, which is used by <i>spsolve()</i> for finding the solutions of sparse systems, as well as <i>eigs_sym()</i> and <i>eigs_gen()</i> in shift-invert mode. You will need to link with the superlu library, for example <i>-lsuperlu</i>
ARMA_DONT_USE_SUPERLU	Disable use of SuperLU; overrides <i>ARMA_USE_SUPERLU</i>
ARMA_USE_HDF5	Enable the ability to <i>save and load</i> matrices stored in the HDF5 format; the <i>hdf5.h</i> header file must be available on your system and you will need to link with the hdf5 library (eg. <i>-lhdf5</i>)
ARMA_DONT_USE_HDF5	Disable the use of the HDF5 library; overrides <i>ARMA_USE_HDF5</i>
ARMA_USE_FFTW3	Enable use of the FFTW3 library by <i>fft()</i> and <i>ifft()</i> ; you will need to link with the FFTW3 library (eg. <i>-lfftw3</i>)
ARMA_DONT_USE_FFTW3	Disable the use of the FFTW3 library; overrides <i>ARMA_USE_FFTW3</i>
ARMA_DONT_USE_STD_MUTEX	Disable use of <i>std::mutex</i> ; applicable if your compiler and/or environment doesn't support <i>std::mutex</i>

ARMA_DONT_OPTIMISE_BAND	Disable automatically optimised handling of band matrices by <code>solve()</code> and <code>chol()</code>
ARMA_DONT_OPTIMISE_SYMPD	Disable automatically optimised handling of symmetric/hermitian positive definite matrices by <code>solve()</code> , <code>inv()</code> , <code>pinv()</code> , <code>expmat()</code> , <code>logmat()</code> , <code>sqrmat()</code> , <code>powmat()</code> , <code>rcond()</code>
ARMA_USE_OPENMP	Use OpenMP for parallelisation of computationally expensive element-wise operations (such as <code>exp()</code> , <code>log()</code> , <code>cos()</code> , etc). Automatically enabled when using a compiler which has OpenMP 3.1+ active (eg. the <code>-fopenmp</code> option for gcc and clang).
ARMA_DONT_USE_OPENMP	Disable use of OpenMP for parallelisation of element-wise operations; overrides <code>ARMA_USE_OPENMP</code>
ARMA_OPENMP_THRESHOLD	The minimum number of elements in a matrix to enable OpenMP based parallelisation of computationally expensive element-wise functions; default value is 320
ARMA_OPENMP_THREADS	The maximum number of threads for OpenMP based parallelisation of computationally expensive element-wise functions; default value is 8
ARMA_BLAS_CAPITALS	Use capitalised (uppercase) BLAS and LAPACK function names (eg. DGEMM vs dgemm)
ARMA_BLAS_UNDERSCORE	Append an underscore to BLAS and LAPACK function names (eg. dgemm_ vs dgemm). Enabled by default.
ARMA_BLAS_LONG	Use "long" instead of "int" when calling BLAS and LAPACK functions
ARMA_BLAS_LONG_LONG	Use "long long" instead of "int" when calling BLAS and LAPACK functions

ARMA_USE_FORTRAN_HIDDEN_ARGS	Use so-called "hidden arguments" when calling BLAS and LAPACK functions. Enabled by default. See Fortran argument passing conventions for more details.
ARMA_DONT_USE_FORTRAN_HIDDEN_ARGS	Disable use of so-called "hidden arguments" when calling BLAS and LAPACK functions. May be necessary when using Armadillo in conjunction with broken MKL headers (eg. if you have <code>#include "mkl_lapack.h"</code> in your code).
ARMA_USE_TBB_ALLOC	Use Intel TBB <i>scalable_malloc()</i> and <i>scalable_free()</i> instead of standard <i>malloc()</i> and <i>free()</i> for managing matrix memory
ARMA_USE_MKL_ALLOC	Use Intel MKL <i>mkl_malloc()</i> and <i>mkl_free()</i> instead of standard <i>malloc()</i> and <i>free()</i> for managing matrix memory
ARMA_USE_MKL_TYPES	Use Intel MKL types for complex numbers. You will need to include appropriate MKL headers before the Armadillo header. You may also need to enable one or more of the following options: ARMA_BLAS_LONG, ARMA_BLAS_LONG_LONG, ARMA_DONT_USE_FORTRAN_HIDDEN_ARGS
ARMA_64BIT_WORD	Use 64 bit integers. Automatically enabled when using a 64-bit platform, except when using Armadillo in the R environment (via RcppArmadillo). Useful if matrices/vectors capable of holding more than 4 billion elements are required. This can also be enabled by adding <code>#define ARMA_64BIT_WORD</code> before each instance of <code>#include <armadillo></code>
ARMA_NO_DEBUG	Disable all run-time checks, including size conformance and bounds checks . NOT RECOMMENDED. DO NOT USE UNLESS YOU KNOW WHAT YOU ARE DOING AND ARE WILLING TO RISK THE DOWNSIDES. Keeping run-time checks enabled during development and deployment greatly aids in finding mistakes in your code.
ARMA_EXTRA_DEBUG	Print out the trace of internal functions used for evaluating expressions. Not recommended for normal use. This is mainly useful for debugging the

library.

ARMA_MAT_PREALLOC

The number of pre-allocated elements used by matrices and vectors. Must be always enabled and set to an integer that is at least 1. By default set to 16. If you mainly use lots of very small vectors (eg. ≤ 4 elements), change the number to the size of your vectors.

ARMA_COUT_STREAM

The default stream used for printing matrices and cubes by `.print()`. Must be always enabled. By default defined to `std::cout`

ARMA_CERR_STREAM

The default stream used for printing warnings and errors. Must be always enabled. By default defined to `std::cerr`

ARMA_WARN_LEVEL

The level of warning messages printed to `ARMA_CERR_STREAM`. Must be an integer ≥ 0 . By default defined to 2.

0 = no warnings; generally not recommended

1 = only critical warnings about arguments and/or data which are likely to lead to incorrect results

2 = as per level 1, and warnings about poorly conditioned systems (low `rcond`) detected by `solve()`, `spsolve()`, etc

3 = as per level 2, and warnings about failed decompositions, failed saving / loading, etc

Example usage:

```
#define ARMA_WARN_LEVEL 1
#include <armadillo>
```

- See also:
 - [element access](#)
 - [element initialisation](#)
 - [uword/sword](#)
 - [output streams](#)

History of API Additions and Changes

- API Stability and Version Policy:
 - Each release of Armadillo has its public API (functions, classes, constants) described in the accompanying API documentation specific to that release.
 - Each release of Armadillo has its full version specified as $A.B.C$, where A is a major version number, B is a minor version number, and C is a patch level. The version specification has explicit meaning (similar to **Semantic Versioning**) as follows:
 - Within a major version (eg. 10), each minor version (eg. 10.2) has a public API that strongly strives to be **backwards compatible** (at the source level) with the public API of preceding minor versions. For example, user code written for version 10.0 should work with version 10.1, 10.2, etc. However, subsequent minor versions may have more features (API additions and extensions) than preceding minor versions. As such, user code *specifically* written for version 10.2 may not work with 10.1.
 - An increase in the patch level, while the major and minor versions are retained, indicates modifications to the code and/or documentation which aim to fix bugs without altering the public API.
 - We don't like changes to existing public API and strongly prefer not to break any user software. However, to allow evolution, the public API may be altered in future major versions while remaining backwards compatible in as many cases as possible (eg. major version 11 may have slightly different public API than major version 10).
 - **Caveat:** the above policy applies only to the public API described in the documentation. Any functionality within Armadillo which is not explicitly described in the public API documentation is considered as internal implementation detail, and may be changed or removed without notice.
- List of additions and changes for each version:
 - Version 12.6:
 - faster multiplication of dense vectors by **sparse matrices** (and vice versa)

- faster `eigs_sym()`, `eigs_gen()`, `svds()`
 - faster `conv()` and `conv2()` when using OpenMP
 - added `diags()` and `spdiags()` for generating band matrices from set of vectors
- Version 12.4:
 - added `norm2est()` for finding fast estimates of matrix 2-norm (spectral norm)
 - added `vecnorm()` for obtaining the vector norm of each row or column of a matrix
- Version 12.2:
 - more efficient use of FFTW3 by `fft()` and `ifft()`
 - faster in-place element-wise multiplication of `sparse matrices` by dense matrices
 - added `spsolve_factoriser` class to allow reuse of sparse matrix factorisation for solving systems of linear equations
- Version 12.0:
 - faster `fft()` and `ifft()` via optional use of FFTW3
 - faster `min()` and `max()`
 - faster `index_min()` and `index_max()`
 - added `.col_as_mat()` and `.row_as_mat()` which return matrix representation of cube column and cube row
 - added `csv_opts::strict` option to `loading CSV files` to interpret missing values as NaN
 - added `check_for_zeros` option to form 4 of `sparse matrix batch constructors`
 - `inv()` and `inv_sympd()` with options `inv_opts::no_ugly` OR `inv_opts::allow_approx` now use a scaled threshold similar to `pinv()`
 - `set_cout_stream()` and `set_cerr_stream()` are now no-ops; instead use the options `ARMA_WARN_LEVEL`, or `ARMA_COUT_STREAM`, or `ARMA_CERR_STREAM`
- Version 11.4:
 - extended `pow()` with various forms of element-wise power operations
 - added `find_nan()` to find indices of NaN elements
 - faster handling of compound expressions by `sum()`
- Version 11.2:
 - extended `randu()` and `randn()` to allow specification of distribution parameters
 - added `inv_opts::no_ugly` option to `inv()` and `inv_sympd()` to disallow inverses of poorly conditioned matrices

- more efficient handling of rank-deficient matrices via `inv_opts::allow_approx` option in `inv()` and `inv_sympd()`
 - faster handling of sparse submatrix column views by `norm()`, `accu()`, `nonzeros()`
 - faster handling of symmetric and diagonal matrices by `cond()`
 - better detection of rank deficient matrices by `solve()`
- Version 11.0:
 - added variants of `inv()` and `inv_sympd()` that provide *rcond* (reciprocal condition number)
 - added `inv_opts::allow_approx` option to `inv()` and `inv_sympd()` to allow approximate inverses of poorly conditioned matrices
 - stricter handling of singular matrices by `inv()` and `inv_sympd()`
 - stricter handling of non-sympd matrices by `inv_sympd()`
 - stricter handling of non-finite matrices by `pinv()`
 - more robust handling of rank deficient matrices by `solve()`
 - faster handling of diagonal matrices by `rcond()`
 - changed `eigs_sym()` and `eigs_gen()` to use higher quality RNG
 - `quantile()` and `median()` will now throw an exception if given matrices/vectors have NaN elements
 - Version 10.8:
 - faster handling of symmetric matrices by `pinv()` and `rank()`
 - faster handling of diagonal matrices by `inv_sympd()`, `pinv()`, `rank()`
 - expanded `norm()` to handle integer vectors and matrices
 - added `datum::tau` to replace 2π
 - Version 10.7:
 - faster handling of **submatrix views** accessed by `X.cols(first_col, last_col)`
 - faster handling of element-wise `min()` and `max()` in compound expressions
 - expanded `solve()` with `solve_opts::force_approx` option to force use of the approximate solver
 - Version 10.6:
 - expanded `chol()` to optionally use pivoted decomposition
 - expanded vector, matrix and cube constructors to allow element initialisation via `fill::value(scalar)`, eg. `mat X(4,5,fill::value(123))`
 - faster **loading** of CSV files when using OpenMP
 - added `csv_opts::semicolon` option to allow **saving / loading** of CSV files with semicolon (;) instead of comma (,) as the separator

- Version 10.5:
 - added `.clamp()` member function
 - expanded the standalone `clamp()` function to handle complex values
 - more efficient use of OpenMP
 - vector, matrix and cube constructors now initialise elements to zero by default; element initialisation can be disabled via the `fill::none` specifier, eg. `mat X(4,5,fill::none)`
- Version 10.4:
 - faster handling of triangular matrices by `log_det()`
 - added `log_det_sympd()` for log determinant of symmetric positive matrices
 - added `ARMA_WARN_LEVEL` `configuration option`, to control the degree of emitted warning messages
 - reduced the default degree of warning messages, so that failed decompositions, failed saving / loading, etc, no longer emit warnings
- Version 10.3:
 - faster handling of symmetric positive definite matrices by `pinv()`
 - expanded `.save()` / `.load()` for dense matrices to handle `coord_ascii` format
 - for out of bounds access, `element accessors` now throw the more nuanced `std::out_of_range` exception, instead of only `std::logic_error`
 - improved quality of random numbers
- Version 10.2:
 - faster handling of `subcubes`
 - added `tgamma()`
 - added `.brief_print()` for abridged printing of matrices & cubes
 - expanded sparse matrix forms of `trimatu()` and `trimatl()` to allow specifying the diagonal delimiter
 - expanded `eigs_sym()` and `eigs_gen()` with optional shift-invert mode
- Version 10.1:
 - C++11 is now the minimum required C++ standard
 - faster handling of compound expressions by `trimatu()` and `trimatl()`
 - faster sparse matrix addition, subtraction and element-wise multiplication
 - expanded sparse `submatrix views` to handle the non-contiguous form of `X.cols(vector_of_column_indices)`
 - expanded `eigs_sym()` and `eigs_gen()` with optional fine-grained parameters

- Version 9.900:
 - faster `solve()` for under/over-determined systems
 - faster `eig_gen()` and `eig_pair()` for large matrices
 - faster handling of matrix multiplication expressions by `diagvec()` and `diagmat()`
 - faster handling of relational expressions by `accu()`
 - faster handling of sympd matrices by `expmat()`, `logmat()`, `sqrmat()`
 - faster access to columns in sparse submatrix views
 - added `quantile()`
 - added `powmat()`
 - added `trimatu_ind()` and `trimatl_ind()`
 - added `log_normpdf()`
 - added `.is_zero()`
 - added `ARMA_DONT_USE_CXX11_MUTEX` configuration option to disable use of `std::mutex`
 - expanded `eig_gen()` and `eig_pair()` to optionally provide left and right eigenvectors
 - expanded `qr()` to optionally use pivoted decomposition
 - expanded `.save()` and `.load()` to handle CSV files with headers via `csv_name(filename,header)` specification
 - more consistent detection of sparse vector expressions
 - updated `physical constants` to NIST 2018 CODATA values
 - workaround for `save/load` issues with HDF5 v1.12
- Version 9.800:
 - faster `solve()` in default operation; iterative refinement is no longer applied by default; use `solve_opts::refine` to explicitly enable refinement
 - faster `expmat()`
 - faster handling of triangular matrices by `rcond()`
 - added `.front()` and `.back()`
 - added `.is_trimatu()` and `.is_trimatl()`
 - added `.is_diagmat()`
- Version 9.700:
 - faster handling of cubes by `vectorise()`
 - faster handling of sparse matrices by `nonzeros()`
 - faster row-wise `index_min()` and `index_max()`
 - expanded `join_rows()` and `join_cols()` to handle joining up to 4 matrices

- expanded `.save()` and `.load()` to allow storing sparse matrices in CSV format
- added `randperm()` to generate a vector with random permutation of a sequence of integers
- Version 9.600:
 - faster handling of sparse `submatrices`
 - faster handling of sparse `diagonal views`
 - faster handling of sparse matrices by `symmatu()` and `symmatl()`
 - faster handling of sparse matrices by `join_cols()`
 - expanded `clamp()` to handle sparse matrices
 - added `.clean()` to replace elements below a threshold with zeros
- Version 9.500:
 - expanded `solve()` with `solve_opts::likely_sympd` to indicate that the given matrix is likely positive definite
 - more robust automatic detection of positive definite matrices by `solve()` and `inv()`
 - faster handling of sparse submatrices
 - expanded `eigs_sym()` to print a warning if the given matrix is not symmetric
 - extended LAPACK function prototypes to follow Fortran `passing conventions` for so-called "hidden arguments", in order to address [GCC Bug 90329](#);
to use previous LAPACK function prototypes without the "hidden arguments", `#define ARMA_DONT_USE_FORTRAN_HIDDEN_ARGS` before `#include <armadillo>`
- Version 9.400:
 - faster `cov()` and `cor()`
 - added `.as_col()` and `.as_row()`
 - expanded `.shed_rows()` / `.shed_cols()` / `.shed_slices()` to remove rows/columns/slices specified in a vector
 - expanded `vectorise()` to handle sparse matrices
 - expanded element-wise versions of `max()` and `min()` to handle sparse matrices
 - optimised handling of sparse matrix expressions: `sparse % (sparse +- scalar)` and `sparse / (sparse +- scalar)`
 - expanded `eig_sym()`, `chol()`, `expmat_sym()`, `logmat_sympd()`, `sqrtmat_sympd()`, `inv_sympd()` to print a warning if the given matrix is not symmetric
 - more consistent detection of vector expressions
- Version 9.300:

- faster handling of compound complex matrix expressions by `trace()`
 - more efficient handling of element access for inplace modifications in `sparse matrices`
 - added `.is_sympd()` to check whether a matrix is symmetric/hermitian positive definite
 - added `interp2()` for 2D data interpolation
 - added `expm1()` and `log1p()`
 - expanded `.is_sorted()` with options "strictascend" and "strictdescend"
 - expanded `eig_gen()` to optionally perform balancing prior to decomposition
- Version 9.200:
 - faster handling of symmetric positive definite matrices by `rcond()`
 - faster transpose of matrices with size $\geq 512 \times 512$
 - faster handling of compound sparse matrix expressions by `accu()`, `diagmat()`, `trace()`
 - faster handling of sparse matrices by `join_rows()`
 - added `sinc()`
 - expanded `sign()` to handle scalar arguments
 - expanded `operators` (*, %, +, -) to handle sparse matrices with differing element types (eg. multiplication of complex matrix by real matrix)
 - expanded `conv_to()` to allow conversion between sparse matrices with differing element types
 - expanded `solve()` to optionally allow keeping solutions of systems singular to working precision
 - Version 9.100:
 - faster handling of symmetric/hermitian positive definite matrices by `solve()`
 - faster handling of `inv_sympd()` in compound expressions
 - added `.is_symmetric()`
 - added `.is_hermitian()`
 - expanded `spsolve()` to optionally allow keeping solutions of systems singular to working precision
 - new `configuration` options ARMA_OPTIMISE_SOLVE_BAND and ARMA_OPTIMISE_SOLVE_SYMPD
 - smarter use of the element cache in sparse matrices
 - Version 8.600:
 - added `hess()` for Hessenberg decomposition
 - added `.row()`, `.rows()`, `.col()`, `.cols()` to `subcube views`
 - expanded `.shed_rows()` and `.shed_cols()` to handle cubes
 - expanded `.insert_rows()` and `.insert_cols()` to handle cubes
 - expanded `subcube views` to allow non-contiguous access to slices
 - improved tuning of `sparse matrix` element access operators

- faster handling of tridiagonal matrices by `solve()`
- faster multiplication of matrices with differing element types when using OpenMP
- Version 8.500:
 - faster handling of sparse matrices by `kron()` and `repmat()`
 - faster `transpose` of sparse matrices
 - faster `element access` in sparse matrices
 - faster `row iterators` for sparse matrices
 - faster handling of compound expressions by `trace()`
 - more efficient handling of aliasing in `submatrix views`
 - expanded `normalise()` to handle sparse matrices
 - expanded `.transform()` and `.for_each()` to handle sparse matrices
 - added `reverse()` for reversing order of elements
 - added `repelem()` for replicating elements
 - added `roots()` for finding the roots of a polynomial
- Version 8.400:
 - faster handling of sparse matrices by `repmat()`
 - faster `loading` of CSV files
 - expanded `kron()` to handle sparse matrices
 - expanded `index_min()` and `index_max()` to handle cubes
 - expanded `randi()`, `randu()`, `randn()`, `randg()` to output single scalars
 - added `submatrix & subcube iterators`
 - added `normcdf()`
 - added `mvnrnd()`
 - added `chi2rnd()`
 - added `wishrnd()` and `iwishrnd()`
- Version 8.300:
 - faster handling of band matrices by `solve()`
 - faster handling of band matrices by `chol()`
 - faster `randg()` when using OpenMP
 - added `normpdf()`
 - expanded `.save()` to allow appending new datasets to existing HDF5 files
- Version 8.200:

- added `intersect()` for finding common elements in two vectors/matrices
- expanded `affmul()` to handle non-square matrices
- Version 8.100:
 - faster incremental construction of `sparse matrices` via element access operators
 - faster `diagonal views` in `sparse matrices`
 - expanded `SpMat` to save/load sparse matrices in `coord format`
 - expanded `.save()/load()` to allow specification of datasets within HDF5 files
 - added `affmul()` to simplify application of affine transformations
 - warnings and errors are now printed by default to the `std::cerr` stream
 - added `set_cerr_stream()` and `get_cerr_stream()` to replace `set_stream_err1()`, `set_stream_err2()`, `get_stream_err1()`, `get_stream_err2()`
 - new `configuration` options `ARMA_COUT_STREAM` and `ARMA_CERR_STREAM`
- Version 7.960:
 - faster `randn()` when using OpenMP
 - faster `gmm_diag` class, for Gaussian mixture models with diagonal covariance matrices
 - added `.sum_log_p()` to the `gmm_diag` class
 - added `gmm_full` class, for Gaussian mixture models with full covariance matrices
 - expanded `.each_slice()` to optionally use OpenMP for multi-threaded execution
- Version 7.950:
 - expanded `accu()` and `sum()` to use OpenMP for processing expressions with computationally expensive element-wise functions
 - expanded `trimatu()` and `trimatl()` to allow specification of the diagonal which delineates the boundary of the triangular part
- Version 7.900:
 - expanded `clamp()` to handle cubes
 - computationally expensive element-wise functions (such as `exp()`, `log()`, `cos()`, etc) can now be automatically sped up via `OpenMP`; this requires a C++11/C++14 compiler with OpenMP 3.1+ support
 - for GCC and clang compilers use the following options to enable both C++11 and OpenMP:
`-std=c++11 -fopenmp`
 - **Caveat:** when using GCC, use of `-march=native` in conjunction with `-fopenmp` may lead to speed regressions on recent processors

- Version 7.800:
 - changed license to the permissive [Apache License 2.0](#); see the [Questions page](#) for more info
- Version 7.700:
 - added [polyfit\(\)](#) and [polyval\(\)](#)
 - added second form of [log_det\(\)](#) to directly return the result as a complex number
 - added [range\(\)](#) to statistics functions
 - expanded [trimatu\(\)/trimatl\(\)](#) and [symmatu\(\)/symmatl\(\)](#) to handle sparse matrices
- Version 7.600:
 - more accurate [eigs_sym\(\)](#) and [eigs_gen\(\)](#)
 - expanded [floor\(\)](#), [ceil\(\)](#), [round\(\)](#), [trunc\(\)](#), [sign\(\)](#) to handle sparse matrices
 - added [arg\(\)](#), [atan2\(\)](#), [hypot\(\)](#)
- Version 7.500:
 - expanded [qz\(\)](#) to optionally specify ordering of the Schur form
 - expanded [.each_slice\(\)](#) to support matrix multiplication
- Version 7.400:
 - added [expmat_sym\(\)](#), [logmat_sympd\(\)](#), [sqrtmat_sympd\(\)](#)
 - added [.replace\(\)](#)
- Version 7.300:
 - added [index_min\(\)](#) and [index_max\(\)](#) standalone functions
 - expanded [.subvec\(\)](#) to accept [size\(\)](#) arguments
 - more robust handling of non-square matrices by [lu\(\)](#)
- Version 7.200:
 - added [.index_min\(\)](#) and [.index_max\(\)](#) member functions
 - expanded [ind2sub\(\)](#) to handle vectors of indices
 - expanded [sub2ind\(\)](#) to handle matrix of subscripts
 - expanded [expmat\(\)](#), [logmat\(\)](#) and [sqrtmat\(\)](#) to optionally return a bool indicating success
 - faster handling of compound expressions by [vectorise\(\)](#)
- Version 7.100:

- added `erf()`, `erfc()`, `lgamma()`
 - added `.head_slices()` and `.tail_slices()` to subcube views
 - `spsolve()` now requires SuperLU 5.2
 - `eigs_sym()`, `eigs_gen()` and `svds()` now use a built-in reimplementaion of ARPACK for real (non-complex) matrices; contributed by Yixuan Qiu
- Version 6.700:
 - added `trapz()` for numerical integration
 - added `logmat()` for calculating the matrix logarithm
 - added `regspace()` for generating vectors with regularly spaced elements
 - added `logspace()` for generating vectors with logarithmically spaced elements
 - added `approx_equal()` for determining approximate equality
- Version 6.600:
 - expanded `sum()`, `mean()`, `min()`, `max()` to handle cubes
 - expanded `Cube` class to handle arbitrarily sized empty cubes (eg. 0x5x2)
 - added `shift()` for circular shifts of elements
 - added `sqrmat()` for finding the square root of a matrix
- Version 6.500:
 - added `conv2()` for 2D convolution
 - added stand-alone `kmeans()` function for clustering data
 - added `trunc()`
 - extended `conv()` to optionally provide central convolution
 - faster handling of multiply-and-accumulate by `accu()` when using Intel MKL, ATLAS or OpenBLAS
- Version 6.400:
 - expanded `each_col()`, `each_row()` and `each_slice()` to handle C++11 lambda functions
 - added `ind2sub()` and `sub2ind()`
- Version 6.300:
 - expanded `solve()` to find approximate solutions for rank-deficient systems
 - faster handling of `non-contiguous submatrix views` in compound expressions
 - added `.for_each()` to `Mat`, `Row`, `Col`, `Cube` and `field` classes
 - added `rcond()` for estimating the reciprocal condition number

- Version 6.200:
 - expanded `diagmat()` to handle non-square matrices and arbitrary diagonals
 - expanded `trace()` to handle non-square matrices
- Version 6.100:
 - faster `norm()` and `normalise()` when using Intel MKL, ATLAS or OpenBLAS
 - added Schur decomposition: `schur()`
 - stricter handling of matrix objects by `hist()` and `histc()`
 - `advanced constructors` for using auxiliary memory by Mat, Col, Row and Cube now have the default of `strict = false`
 - `Cube` class now delays allocation of `.slice()` related structures until needed
 - expanded `join_slices()` to handle joining cubes with matrices
- Version 5.600:
 - added `.each_slice()` for matrix operations applied to each slice of a cube
 - expanded `.each_col()` and `.each_row()` to handle out-of-place operations
- Version 5.500:
 - expanded object constructors and generators to handle `size()` based specification of dimensions
- Version 5.400:
 - added `find_unique()` for finding indices of unique values
 - added `diff()` for calculating differences between consecutive elements
 - added `cumprod()` for calculating cumulative product
 - added `null()` for finding the orthonormal basis of null space
 - expanded `interp1()` to handle repeated locations
 - expanded `unique()` to handle complex numbers
 - faster `flipud()`
 - faster row-wise `cumsum()`
- Version 5.300:
 - added generalised Schur decomposition: `qz()`
 - added `.has_inf()` and `.has_nan()`
 - expanded `interp1()` to handle out-of-domain locations
 - expanded `sparse matrix` class with `.set_imag()` and `.set_real()`
 - expanded `imag()`, `real()` and `conj()` to handle sparse matrices

- expanded `diagmat()`, `reshape()` and `resize()` to handle sparse matrices
 - faster sparse `sum()`
 - faster row-wise `sum()`, `mean()`, `min()`, `max()`
 - updated `physical constants` to NIST 2014 CODATA values
- Version 5.200:
 - added `orth()` for finding the orthonormal basis of the range space of a matrix
 - expanded `element initialisation` to handle nested initialiser lists (C++11)
- Version 5.100:
 - added `interp1()` for 1D interpolation
 - added `.is_sorted()` for checking whether a vector or matrix has sorted elements
 - updated `physical constants` to NIST 2010 CODATA values
- Version 5.000:
 - added `spsolve()` for solving sparse systems of linear equations
 - added `svds()` for singular value decomposition of sparse matrices
 - added `nonzeros()` for extracting non-zero values from matrices
 - added handling of `diagonal views` by sparse matrices
 - expanded `repmat()` to handle sparse matrices
 - expanded `join_rows()` and `join_cols()` to handle sparse matrices
 - `sort_index()` and `stable_sort_index()` have been placed in the delayed operations framework for increased efficiency
 - use of `64 bit integers` is automatically enabled when using a C++11 compiler
- Version 4.650:
 - added `randg()` for generating random values from gamma distributions (C++11 only)
 - added `.head_rows()` and `.tail_rows()` to `submatrix views`
 - added `.head_cols()` and `.tail_cols()` to `submatrix views`
 - expanded `eigs_sym()` to optionally calculate eigenvalues with smallest/largest algebraic values
- Version 4.600:
 - added `.head()` and `.tail()` to `submatrix views`
 - faster matrix transposes within compound expressions
 - faster in-place matrix multiplication
 - faster `accu()` and `norm()` when compiling with `-O3 -ffast-math -march=native` (gcc and clang)

- Version 4.550:
 - added matrix exponential function: `expmat()`
 - faster `.log_p()` and `.avg_log_p()` functions in the `gmm_diag` class when compiling with OpenMP enabled
 - faster handling of in-place addition/subtraction of expressions with an outer product
- Version 4.500:
 - faster handling of complex vectors by `norm()`
 - expanded `chol()` to optionally specify output matrix as upper or lower triangular
 - better handling of non-finite values when `saving` matrices as text files
- Version 4.450:
 - faster handling of matrix transposes within compound expressions
 - expanded `symmatu()/symmatl()` to optionally disable taking the complex conjugate of elements
 - expanded `sort_index()` to handle complex vectors
 - expanded the `gmm_diag` class with functions to generate random samples
- Version 4.400:
 - faster handling of subvectors by `dot()`
 - faster handling of aliasing by `submatrix` views
 - added `clamp()` for clamping values to be between lower and upper limits
 - added `gmm_diag` class for statistical modelling of data using Gaussian Mixture Models
 - expanded `batch insertion constructors` for sparse matrices to add values at repeated locations
- Version 4.320:
 - expanded `eigs_sym()` and `eigs_gen()` to use an optional tolerance parameter
 - expanded `eig_sym()` to automatically fall back to standard decomposition method if divide-and-conquer fails
 - cmake-based installer enables use of C++11 random number generator when using gcc 4.8.3+ in C++11 mode
- Version 4.300:
 - added `find_finite()` and `find_nonfinite()`
 - expressions $X = \text{inv}(A) * B * C$ and $X = A.i() * B * C$ are automatically converted to $X = \text{solve}(A, B * C)$

- Version 4.200:
 - faster transpose of sparse matrices
 - more efficient handling of aliasing during matrix multiplication
 - faster `inverse` of matrices marked as diagonal
- Version 4.100:
 - added `normalise()` for normalising vectors to unit p -norm
 - extended the `field class` to handle 3D layout
 - extended `eigs_sym()` and `eigs_gen()` to obtain eigenvalues of various forms (eg. largest or smallest magnitude)
 - automatic SIMD vectorisation of elementary expressions (eg. matrix addition) when using Clang 3.4+ with -O3 optimisation
 - faster handling of sparse submatrix views
- Version 4.000:
 - added eigen decompositions of sparse matrices: `eigs_sym()` and `eigs_gen()`
 - added eigen decomposition for pair of matrices: `eig_pair()`
 - added simpler forms of `eig_gen()`
 - added condition number of matrices: `cond()`
 - expanded `find()` to handle cubes
 - expanded `subcube views` to access elements specified in a vector
 - template argument for `running_stat_vec` expanded to accept vector types
 - more robust fast `inverse` of 4x4 matrices
 - faster divide-and-conquer decompositions are now used by default for `eig_sym()`, `pinv()`, `princomp()`, `rank()`, `svd()`, `svd_econ()`
 - the form `inv(sympd(X))` no longer assumes that X is positive definite; use `inv_sympd()` instead
- Version 3.930:
 - added `size()` based specifications of `submatrix view` sizes
 - added element-wise variants of `min()` and `max()`
 - added divide-and-conquer variant of `svd_econ()`
 - added divide-and-conquer variant of `pinv()`
 - added `randi()` for generating matrices with random integer values
 - added `inplace_trans()` for memory efficient in-place transposes
 - added more intuitive specification of sort direction in `sort()` and `sort_index()`
 - added more intuitive specification of method in `det()`, `.i()`, `inv()` and `solve()`

- more precise timer for the `wall_clock` class when using C++11
- Version 3.920:
 - faster `.zeros()`
 - faster `round()`, `exp2()` and `log2()` when using C++11
 - added signum function: `sign()`
 - added move constructors when using C++11
 - added 2D fast Fourier transform: `fft2()`
 - added `.tube()` for easier extraction of vectors and subcubes from cubes
 - added specification of a fill type during construction of `Mat`, `Col`, `Row` and `Cube` classes, eg. `mat X(4, 5, fill::zeros)`
- Version 3.910:
 - faster multiplication of a matrix with a transpose of itself, ie. $X * X.t()$ and $X.t() * X$
 - added `vectorise()` for reshaping matrices into vectors
 - added `all()` and `any()` for indicating presence of elements satisfying a relational condition
- Version 3.900:
 - added automatic SSE2 vectorisation of elementary expressions (eg. matrix addition) when using GCC 4.7+ with -O3 optimisation
 - faster `median()`
 - faster handling of compound expressions with transposes of `submatrix` rows
 - faster handling of compound expressions with transposes of complex vectors
 - added support for `saving & loading` of `cubes` in HDF5 format
- Version 3.820:
 - faster `as_scalar()` for compound expressions
 - faster transpose of small vectors
 - faster matrix-vector product for small vectors
 - faster multiplication of small `fixed size matrices`
- Version 3.810:
 - added fast Fourier transform: `fft()`
 - added handling of `.imbue()` and `.transform()` by submatrices and subcubes
 - added `batch insertion constructors` for sparse matrices

- Version 3.800:
 - added `.imbue()` for filling a matrix/cube with values provided by a functor or lambda expression
 - added `.swap()` for swapping contents with another matrix
 - added `.transform()` for transforming a matrix/cube using a functor or lambda expression
 - added `round()` for rounding matrix elements towards nearest integer
 - faster `find()`
 - changed license to the [Mozilla Public License 2.0](#)
- Version 3.6:
 - faster handling of compound expressions with submatrices and subcubes
 - faster `trace()`
 - added support for loading matrices as text files with *NaN* and *Inf* elements
 - added `stable_sort_index()`, which preserves the relative order of elements with equivalent values
 - added handling of [sparse matrices](#) by `mean()`, `var()`, `norm()`, `abs()`, `square()`, `sqrt()`
 - added saving and loading of sparse matrices in *arma_binary* format
- Version 3.4:
 - added economical QR decomposition: `qr_econ()`
 - added `.each_col()` & `.each_row()` for vector operations applied to each column or row of a matrix
 - added preliminary support for [sparse matrices](#)
 - added ability to [save and load](#) matrices in HDF5 format
 - faster [singular value decomposition](#) via optional use of divide-and-conquer algorithm
 - faster `.randn()`
 - faster `dot()` and `cdot()` for complex numbers
- Version 3.2:
 - added `unique()`, for finding unique elements of a matrix
 - added `.eval()`, for forcing the evaluation of delayed expressions
 - faster [eigen decomposition](#) via optional use of divide-and-conquer algorithm
 - faster `transpose` of vectors and compound expressions
 - faster handling of [diagonal views](#)
 - faster handling of tiny [fixed size](#) vectors (≤ 4 elements)
- Version 3.0:
 - added shorthand for inverse: `.i()`
 - added `datum` class

- added `hist()` and `histc()`
- added non-contiguous `submatrix views`
- faster handling of `submatrix views` with a single row or column
- faster element access in `fixed size matrices`
- faster `repmat()`
- expressions `X=inv(A)*B` and `X=A.i()*B` are automatically converted to `X=solve(A,B)`
- better detection of vector expressions by `sum()`, `cumsum()`, `prod()`, `min()`, `max()`, `mean()`, `median()`, `stddev()`, `var()`
- faster generation of random numbers (eg. `randu()` and `randn()`), via an algorithm that produces slightly different numbers than in 2.x
- support for tying writable auxiliary (external) memory to fixed size matrices has been removed; instead, you can use standard matrices with `writable auxiliary memory`, or initialise fixed size matrices by `copying the memory`; using auxiliary memory with standard matrices is unaffected
- `.print_trans()` and `.raw_print_trans()` have been removed; instead, you can chain `.t()` and `.print()` to achieve a similar result: `X.t().print()`

○ Version 2.4:

- added shorter forms of transposes: `.t()` and `.st()`
- added `.resize()` and `resize()`
- added optional use of 64 bit indices (allowing matrices to have more than 4 billion elements), enabled via `ARMA_64BIT_WORD` in `include/armadillo_bits/config.hpp`
- added experimental support for C++11 initialiser lists, enabled via `ARMA_USE_CXX11` in `include/armadillo_bits/config.hpp`
- refactored code to eliminate warnings when using the Clang C++ compiler
- `umat`, `uvec`, `.min()` and `.max()` have been changed to use the `uword` type instead of the `u32` type; by default the `uword` and `u32` types are equivalent (ie. unsigned integer type with a minimum width 32 bits); however, when the use of 64 bit indices is enabled via `ARMA_64BIT_WORD` in `include/armadillo_bits/config.hpp`, the `uword` type then has a minimum width of 64 bits

○ Version 2.2:

- added `svd_econ()`
- added `circ_toeplitz()`
- added `.is_colvec()` and `.is_rowvec()`

○ Version 2.0:

- `det()`, `inv()` and `solve()` can be forced to use more precise algorithms for tiny matrices ($\leq 4 \times 4$)

- added `syl()`, for solving Sylvester's equation
- added `strans()`, for transposing a complex matrix without taking the complex conjugate
- added `symmatu()` and `symmatl()`
- added submatrices of `submatrices`
- faster `inverse` of symmetric positive definite matrices
- faster element access for `fixed size` matrices
- faster multiplication of tiny matrices (eg. 4x4)
- faster compound expressions containing `submatrices`
- added handling of arbitrarily sized empty matrices (eg. 5x0)
- added `.count()` member function in `running_stat` and `running_stat_vec`
- added `loading & saving` of matrices as CSV text files
- `trans()` now takes the complex conjugate when transposing a complex matrix
- forms of `chol()`, `eig_sym()`, `eig_gen()`, `inv()`, `lu()`, `pinv()`, `princomp()`, `qr()`, `solve()`, `svd()`, `syl()` that do not return a bool indicating success now throw `std::runtime_error` exceptions when failures are detected
- `princomp_cov()` has been removed; `eig_sym()` in conjunction with `cov()` can be used instead
- `.is_vec()` now outputs `true` for empty vectors (eg. 0x1)
- `set_log_stream()` & `get_log_stream()` have been replaced by `set_stream_err1()` & `get_stream_err1()`

◦ Version 1.2:

- added `.min()` & `.max()` member functions of `Mat` and `Cube`
- added `floor()` and `ceil()`
- added representation of "not a number": `math::nan()`
- added representation of infinity: `math::inf()`
- `.in_range()` expanded to use **`span()`** arguments
- `fixed size` matrices and vectors can use auxiliary (external) memory
- `submatrices` and `subfields` can be accessed via **`X(span(a,b), span(c,d))`**
- `subcubes` can be accessed via **`X(span(a,b), span(c,d), span(e,f))`**
- the two argument version of **`span`** can be replaced by **`span::all`** or **`span()`**, to indicate an entire range
- for cubes, the two argument version of **`span`** can be replaced by a single argument version, **`span(a)`**, to indicate a single column, row or slice
- arbitrary "flat" subcubes can be interpreted as matrices; for example:

```
cube Q = randu<cube>(5,3,4);
mat A = Q( span(1), span(1,2), span::all );
// A has a size of 2x4
```

```
vec v = ones<vec>(4);
Q( span(1), span(1), span::all ) = v;
```

- added interpretation of matrices as triangular through `trimatu()` / `trimatl()`
 - added explicit handling of triangular matrices by `solve()` and `inv()`
 - extended syntax for `submatrices`, including access to elements whose indices are specified in a vector
 - added ability to change the stream used for `logging` of errors and warnings
 - added ability to `save/load matrices` in raw binary format
 - added cumulative sum function: `cumsum()`
- Changed in 1.0 (compared to earlier 0.x development versions):
- the 3 argument version of `lu()`, eg. `lu(L,U,X)`, provides L and U which should be the same as produced by Octave 3.2 (this was not the case in versions prior to 0.9.90)
 - `rand()` has been replaced by `randu()`; this has been done to avoid confusion with `std::rand()`, which generates random numbers in a different interval
 - In versions earlier than 0.9.0, some multiplication operations directly converted result matrices with a size of 1x1 into scalars. This is no longer the case. If you know the result of an expression will be a 1x1 matrix and wish to treat it as a pure scalar, use the `as_scalar()` wrapping function
 - Almost all functions have been placed in the delayed operations framework (for speed purposes). This may affect code which assumed that the output of some functions was a pure matrix. The solution is easy, as explained below.

In general, Armadillo queues operations before executing them. As such, the direct output of an operation or function cannot be assumed to be a directly accessible matrix. The queued operations are executed when the output needs to be stored in a matrix, eg. `mat B = trans(A)` or `mat B(trans(A))`. If you need to force the execution of the delayed operations, place the operation or function inside the corresponding Mat constructor. For example, if your code assumed that the output of some functions was a pure matrix, eg. `chol(m).diag()`, change the code to `mat(chol(m)).diag()`. Similarly, if you need to pass the result of an operation such as `A+B` to one of your own functions, use `my_function(mat(A+B))`.

